MCS-068
Predictive Data Analysis

ignou
THE PEOPLE'S UNIVERSITY
Indira Gandhi National Open University
School of Computer and Information Sciences (SOCIS)

MCS-068PREDICTIVE DATA ANALYSIS

BLOCK 1: DATA ANALYSIS – AN INTRODUCTION

BLOCK 2: PREDICTIVE DATA ANALYSIS – I

BLOCK 3: PREDICTIVE DATA ANALYSIS – II

BLOCK 4: R - PROGRAMMING FOR DATA ANALYSIS

COURSE PREPARATION TEAM

Content Writer (Unit-1,2,3,4,6,7)
Dr. VenuGopal, General Manager,
Sify Digital Services Ltd., Noida, Uttar Pradesh
New Delhi

Content Writer (Unit-5)
(Adopted from MCS-226: Block-3 unit 9)
Dr. Akshay Kumar, Associate Professor
SOCIS, IGNOU
Delhi

Content Writer (Unit-8,9,10,11)
Prof. Durgansh Sharma
School of Business and Management
Christ University-Delhi (NCR Campus)

Content Writer (Unit-12)
Dr. Jaidev Sharma, Assistant Professor
School of Engineering and Sciences
Sharda University, Greater Noida, U.P.

Content Writer (Unit-13,14,15,16)
(Adopted from MCS226:Block-4 Unit-13,14,15,16)
Dr. Priya Gupta
Jawahar Lal Nehru University (JNU), New Delhi

Content Editor *(Unit: 1,2,3,4,6,7)*
Prof. Vikas Rao Vadi
Director, Don Bosco Institute of Technology,

Content Editor *(Unit-5,13,14,15,16)*
Prof. T.V. Vijay Kumar
Jawahar Lal Nehru University (JNU), New

Content Editor *(Unit-8,9,10,11,12)*
Prof. Manish Trivedi
Director, School of Sciences, IGNOU

Language Editor
Prof. Parmod Kumar
School of Humanities, IGNOU

Course Coordinator
Dr. Sudhansh Sharma, Associate Professor,
School of Computer and Information Sciences, IGNOU

PRINT PRODUCTION

Sh. Sanjay Aggarwal
Assistant Registrar, MPDD, IGNOU, New Delhi

,2026

ISBN – xx-xxxx-xxx-x

Further information on the Indira Gandhi National Open University courses may be obtained from the University's office at Maidan Garhi, New Delhi 110068.

Printed and published on behalf of the Indira Gandhi National Open University, New Delhi by MPDD, IGNOU.

COURSE INTRODUCTION

The course *Predictive Data Analysis* is designed to develop a strong understanding of data-driven decision-making through statistical techniques, machine learning models, and programming tools. In the modern data-centric environment, predictive analytics plays an important role in forecasting trends, identifying hidden patterns, and supporting strategic decisions across industries. The course enables learners to analyze data, build predictive models, and interpret analytical results for both academic and practical applications.

The course is organized into four blocks comprising sixteen units, ensuring a gradual progression from foundational concepts to advanced predictive techniques.

Block-1: Data Analysis – An Introduction introduces learners to descriptive, diagnostic, predictive, and prescriptive analytics, building the conceptual foundation for data interpretation and decision-making.

Block-2: Predictive Data Analysis-I focuses on analytical techniques such as similarity measures, distance metrics, outlier detection, and introductory predictive models, helping learners bridge the gap between basic data understanding and advanced modeling.

Block-3: Predictive Data Analysis-II forms the core of the course and covers advanced predictive models including regression, time series forecasting methods such as ARIMA and LSTM, supervised learning techniques like KNN, Logistic Regression, Decision Trees, Random Forest, Naïve Bayes, and Support Vector Machines, along with unsupervised learning methods such as K-Means, Hierarchical Clustering, and Association Rules. The block also introduces important performance evaluation metrics including accuracy, precision, recall, F1-score, ROC curve, and PR curve.

Block-4: R Programming for Data Analysis emphasizes practical implementation using R programming, covering data handling, visualization, statistical analysis, and predictive modeling to provide hands-on analytical experience.

The course is highly relevant to both industry and academia. In industry, predictive analytics is widely applied in finance, healthcare, marketing, e-commerce, and manufacturing for forecasting, risk analysis, recommendation systems, and customer segmentation. In academia, it supports research in data science, artificial intelligence, and management analytics by enabling data-driven insights and innovation.

Overall, the course provides a balanced combination of theoretical knowledge and practical implementation, helping learners develop analytical thinking, predictive modeling skills, and real-world problem-solving abilities using data analytics techniques.

BLOCK 1: DATA ANALYSIS – AN INTRODUCTION

COURSE PREPARATION TEAM

Content Writer (Unit-1, 2, 3 & 4)
Dr. VenuGopal
General Manager,
Sify Digital Services Ltd., Noida, Uttar Pradesh

Content Editor
Prof. Vikas Rao Vadi,
Professor & Director,
Don Bosco Institute of Technlogy, New Delhi

Language Editor
Prof. Parmod Kumar,
School of Humanities, IGNOU

Course Coordinator
Dr. Sudhansh Sharma, Associate Professor,
School of Computer and Information Sciences(SOCIS), IGNOU

PRINT PRODUCTION

Sh. Sanjay Aggarwal
Assistant Registrar, MPDD, IGNOU, New Delhi

,2026

ISBN – xx-xxxx-xxx-x

Further information on the Indira Gandhi National Open University courses may be obtained from the University's office at Maidan Garhi, New Delhi 110068.

Printed and published on behalf of the Indira Gandhi National Open University, New Delhi by MPDD, IGNOU.

BLOCK INTRODUCTION (BLOCK 1)

Block-1, titled *Data Analysis – An Introduction*, forms the foundation of the course *Predictive Data Analysis* by introducing learners to the key stages and concepts involved in analyzing data. The block is divided into four units that progress systematically from understanding data to making informed decisions. It emphasizes both theoretical understanding and practical application, helping learners develop a strong analytical mindset.

The block begins with **Unit-1: Descriptive Data Analysis**, which focuses on collecting, organizing, summarizing, and visualizing data. It includes topics such as data types, frequency distribution, graphical representation, and statistical measures like mean and dispersion. This unit helps learners transform raw data into meaningful information for analysis.

Unit-2: Diagnostic Data Analysis explains how to identify the reasons behind trends and patterns in data. It introduces techniques such as correlation analysis and data mining to understand relationships between variables and identify causes of specific outcomes. This unit enables learners to move from understanding "what happened" to analyzing "why it happened."

Unit-3: Predictive Data Analysis introduces forecasting techniques used to predict future outcomes. It covers methods such as regression analysis, time series analysis, classification, and clustering. By using historical data, learners gain the ability to anticipate trends and support future planning and decision-making.

Finally, **Unit-4: Prescriptive Data Analysis** focuses on recommending suitable actions based on analytical insights. It includes optimization techniques, simulation, game theory, and decision-analysis methods. This unit connects data analysis directly with strategic decision-making by answering "what should be done."

The importance of this block is significant in both industry and academia. In industry, descriptive, diagnostic, predictive, and prescriptive analytics are widely used in finance, healthcare, marketing, operations, and business analytics for performance evaluation, forecasting, and decision optimization. In academia, the block provides the conceptual base for advanced studies in data science, machine learning, and analytics research.

Overall, Block-1 provides a complete introduction to the data analysis lifecycle. It equips learners with the essential skills required to understand data, identify patterns, make predictions, and support effective decision-making in real-world applications.

UNIT 1 DESCRIPTIVE DATA ANALYSIS

Structure

1.1 INTRODUCTION

Descriptive analysis is a foundational aspect of data analysis that involves summarizing and organizing data to uncover meaningful patterns, trends, and insights. Before conducting any advanced statistical modelling or inferential analysis, it is essential to understand the basic characteristics of the data at hand. This unit introduces the key concepts and techniques used to perform descriptive analysis, providing a solid groundwork for deeper statistical exploration.

The unit begins by explaining the process of collecting data, which is the first critical step in any data analysis task. It discusses various kinds of data—including qualitative and quantitative types—and how these classifications influence the choice of analytical methods.

Once data has been collected, organizing it into a frequency distribution helps to identify how values are distributed across different categories or intervals. To enhance interpretability, the unit explores graphical representations of frequency distributions, such as histograms, bar charts, and pie charts, which visually depict the spread and concentration of data values.

Following this, the unit covers techniques for summarization of data, highlighting how large and complex datasets can be distilled into more manageable forms. This leads to a discussion on measures of central tendency, including the mean, median, and mode, which describe the central or typical value in a dataset. Complementing these are the measures of dispersion or variability, such as range, variance, and standard deviation, which provide insights into the spread or consistency of the data.

Through these topics, the unit equips readers with essential tools for understanding and describing data in a clear, concise, and informative manner. These skills are not only crucial for statistical analysis but also for effective communication of findings in academic, professional, and real-world contexts.

1.2 EXPECTED LEARNING OUTCOMES

After completing this unit, you should be able to:

- Understand the process of collecting data and identify appropriate sources and methods for different types of analysis.
- Distinguish between different kinds of data such as qualitative and quantitative, and understand their relevance in statistical analysis.
- Organize raw data into a frequency distribution to facilitate easier interpretation and analysis.
- Create and interpret various graphical representations of frequency distributions, including bar charts, histograms, and pie charts.
- Summarize datasets effectively using basic descriptive statistics.
- Calculate and interpret measures of central tendency such as mean, median, and mode to describe the central value of a dataset.
- Calculate and interpret measures of dispersion or variability, such as range, variance, and standard deviation, to understand the spread of data.
- Apply descriptive analysis techniques to extract meaningful insights and support data-driven decision-making.

1.3 DESCRIPTIVE DATA ANALYSIS

Descriptive Data Analysis is the process of examining, summarizing, and visualizing data to understand its main characteristics before applying any formal statistical modeling or hypothesis testing. It helps in making sense of raw data by organizing it into meaningful patterns and trends. This method is widely used in research, business, healthcare, education, and many other fields for preliminary data interpretation. It forms the backbone of any data-driven inquiry. By systematically collecting, organizing, visualizing, and summarizing data, it allows analysts to extract meaningful insights that inform decision-making and further statistical procedures. Whether dealing with student scores, sales figures, or survey responses, the principles covered in this unit—ranging from frequency distributions to measures of central tendency and variability—equip learners with essential tools for effective data analysis.

Data Collection is the very first step in the process of analysis, thus before any analysis can begin, data must be collected systematically. Data collection involves gathering raw information from various sources, which can include surveys, experiments, observations, administrative records, and digital sources (e.g., web data, transaction logs). *Example:* A teacher wants to analyse students' performance in a mathematics test. They collect marks out of 100 from each student in a class of 40.

But for collecting the data systematically one should have the knowledge of the various *Types of Data*, because understanding the type of data is crucial for selecting appropriate analysis methods. Data can be broadly classified as Qualitative or Categorical data type and Quantitative or Numerical data type.

The Qualitative (Categorical) Data is further classified as follows:

- Nominal: Categories with no inherent order (e.g., gender, religion)
- Ordinal: Categories with a logical order (e.g., education level: high school, undergraduate, postgraduate)

The Quantitative (Numerical) Data is further classified as follows:

- Discrete: Countable values (e.g., number of children)
- Continuous: Any value within a range (e.g., height, weight)

Now, after collection of suitable data, it is generally required to analyse the data where in the *frequency distribution* is the point to begin with. A frequency distribution shows how often each different value in a set of data occurs. It helps in identifying patterns such as concentration of data, outliers, and data spread.

Further, the Graphical Representation of Frequency Distribution i.e. Visualizing data helps you quickly identify patterns, trends, and outliers. Several types of graphs are used in descriptive analysis, a few are listed below:

- *Bar charts* are used for categorical data.
- *Histograms* show the distribution of continuous numerical data.
- *Pie charts* display proportions of categories.
- *Line graphs* show trends over time.

These graphs are quite useful in *summarizing the data, which involves* condensing the data into key statistics or visual formats that convey the main features. Summarizing the data helps you focus on important features without listing every data point. Summaries can be in the form of tables, graphs, or statistics like the average, range, and median etc., which are in fact the Measures of Central Tendency. These measures tell you what is typical or average in your dataset. The three main measures of central tendency are as follows:

- *Mean* is the arithmetic average. Add all values and divide by the number of values.
- *Median* is the middle value when data is sorted.
- *Mode* is the value that appears most frequently.

Each measure gives different insights into the dataset. The mean tells you the overall average, the median shows the central value, and the mode highlights the most frequent score.

Further, to look into the quality of decision one needs to *Measure the Dispersion or Variability.* These measures of dispersion or variability, describe the spread or variability in a dataset. The measures of dispersion are classified as follows:

- *Range:* Difference between the maximum and minimum values
- *Variance:* Average of the squared differences from the Mean
- *Standard Deviation (SD):* Square root of the variance; shows how much data deviates from the mean

A higher standard deviation indicates more variability in the data, while a lower one suggests the data points are closer to the mean.

These concepts are discussed in more detail in the subsequent sections, and as you practice these concepts with real-world data, you'll develop a strong foundation in data analysis and become better equipped to interpret and communicate your findings effectively.

☞ **Check Your Progress 1**

Q1. Explain the concept of Descriptive Data Analysis. Why is it important in data analysis?

……………………………………………………………………………………………

……………………………………………………………………………………………

Q2 Classify the types of data with suitable examples.

……………………………………………………………………………………………

……………………………………………………………………………………………

Q3. What are Measures of Central Tendency and Dispersion? Explain briefly.

……………………………………………………………………………………………

……………………………………………………………………………………………

1.3.1 DATA COLLECTION

Data collection is the first and one of the most important steps in any data analysis. It means carefully gathering the information you need to study a topic or answer a question. Whether you're doing research, running a project, or making business choices, the data you collect is the foundation for your work. If the data is not collected properly, your results might be wrong and lead to bad decisions. But if the data is collected correctly, it helps you understand the situation better and make smarter, more accurate decisions.

One needs to understand that this step is quite crucial, as it is just the beginning of any study/analysis, if error persists in the beginning itself i.e. at the stage of data collection itself, then it may amplify the errors and inaccuracy of results at later stages.

When working with data for analysis, it's important to know where your data is coming from, is it coming from the primary source or the secondary source? Depending on the source of data collection the data is broadly classified into two categories i.e. Primary data and Secondary Data. Knowing the difference between these two helps you decide how to collect the right kind of data for your project or research.

Before understanding the methods of gathering these types of data i.e. Primary data or Secondary data, lets understand the meaning of Primary data and Secondary data respectively.

Primary data is data that you collect yourself, directly from the source, for a specific purpose. This means it is original and has never been used before. You decide what information to collect, how to collect it, and from whom. Since you are directly involved in gathering the data, it is often **more accurate and relevant** to your needs.
However, collecting primary data can take a lot of time and effort. It might also be more expensive, especially if you're surveying a large number of people or conducting experiments.

Let's look at a few common methods used to collect primary data:

- **Surveys and Questionnaires**: These are sets of questions you ask people to get their responses. For example, a company may ask customers to fill out a feedback form after using a product.
- **Interviews**: In this method, you talk to people directly—either face-to-face or over a call—and ask specific questions. For instance, a researcher might interview farmers to learn about their crop practices.
- **Observations**: Here, you watch how people behave or how things happen and take notes. For example, a teacher may observe students during a group activity to understand how they interact.
- **Experiments**: You create a controlled environment to test how changes in one thing affect another. For example, a scientist might test how a new fertilizer affects plant growth.
- **Focus Groups**: This involves a small group of people discussing a topic while a moderator guides the conversation. For instance, a company may gather a group to talk about a new product idea and get their opinions.

Each of these methods gives you fresh, direct information that is well-suited to your specific needs.

Secondary Data is the information that has already been collected and published by someone else, usually for a different purpose. You use this existing data for your own study.

For example, if you are studying population growth, you might use census data already published by the government. This saves you time and money because the data is already available.

Secondary data can come from many sources, such as:

- **Government Reports and Publications**: These include census data, health statistics, or economic reports from official bodies like the Ministry of Health or the Reserve Bank.
- **Institutional Records**: Schools, colleges, hospitals, and companies often maintain records that can be useful. For example, student performance data from a school can help in analyzing academic trends.
- **Research Articles and Journals:** You can use published research papers to support your work or learn from previous studies.
- **Books, Magazines, and Newspapers**: These can offer useful facts, opinions, or background information.
- **Websites and Online Databases**: Many organizations make data freely available online. For example, economic data can be downloaded from the World Bank or websites like data.gov.in.
- **Company Reports:** Financial statements, annual reports, and sales data can help analyze business performance.

Based on the above information one may consolidate the ***differences between the Primary and Secondary Data****. Broadly the difference between the two are as follows*:

- Primary data is collected directly by you for your current study. It is original, often more accurate, but can take more time and effort to gather.
- Secondary data is collected by someone else, usually for a different purpose. It is easy to access and less costly, but may not fully meet your current needs.

In general, there is a question that among the primary data and the secondary data, Which One Should You Use? The answer to this query about the choice among primary and secondary data depends on your goals, time, and budget.

- If you need specific, detailed, and up-to-date information, and you have the resources to collect it yourself, go for primary data.
- If good-quality data already exists and is suitable for your purpose, secondary data can save you time and effort.

In many cases, a good research project uses both types of data—primary data to get fresh insights, and secondary data to support or compare with existing information.

Understanding the difference between primary and secondary data helps you plan your data collection more effectively. Primary data gives you control and relevance, while secondary data gives you speed and convenience. By choosing the right type of data—or combining both—you can ensure your analysis is strong, reliable, and meaningful.

☞ **Check Your Progress 2**

Q1. What is Data Collection? Why is it considered a crucial step in data analysis?

……………………………………………………………………………………………

……………………………………………………………………………………………

Q2. Differentiate between Primary Data and Secondary Data.

……………………………………………………………………………………………

……………………………………………………………………………………………

Q3. Explain any three methods of collecting Primary Data with examples.

……………………………………………………………………………………………

……………………………………………………………………………………………

1.3.2 TYPES OF DATA

Before you begin any kind of data analysis, it's important to understand the type of data you're working with. This is because different types of data require different methods of analysis, visual representation, and interpretation. We learned from the above section 1.4 that depending on the source of data collection the data is broadly classified in to two categories i.e. Primary data and Secondary Data, and we learned about them. Now we are in position to take a deeper dive and understand the various types of the data which contributes towards the data analysis. Although a brief introduction for the types of data was made in the introduction section 1.1, now we will elaborate and discuss these data types here.

Broadly, data can be classified into two main categories: qualitative (or categorical) data and quantitative (or numerical) data. Each of these has subtypes, and some data types have more structure than others.

Qualitative Data (Categorical Data): It describes qualities or characteristics. This type of data is non-numeric and is used to classify data into categories or groups. You cannot perform mathematical operations like addition or averaging on this data, but you can count how often each category occurs.

There are two main types of qualitative data:

a) Nominal Data: Nominal data refers to data that labels categories without any order. These categories are simply names or labels used to identify different groups. There is no ranking or natural order among the values. *Examples*: **Gender**: *Male, Female, Other*; **Blood Group**: *A, B, AB, O;* **Types of Fruits**: *Apple, Banana, Mango*. Here, one category isn't greater or smaller than the other—just different.

b) Ordinal Data: Ordinal data represents categories that do have a logical order, but the differences between the categories aren't measurable. You can rank them, but you can't tell exactly how much more or less one is compared to another. *Examples*:

Education Level: *High School, Undergraduate, Postgraduate;* **Customer Satisfaction**: *Very Unsatisfied, Unsatisfied, Neutral, Satisfied, Very Satisfied;* **Socioeconomic Status**: *Low, Middle, High.*

You can say that someone with a postgraduate degree has a higher qualification than someone with just high school, but you can't measure exactly how much more educated they are based on these labels.

Quantitative Data (Numerical Data): It consists of numbers and represents values that can be measured or counted. You can perform mathematical operations on this data, and it provides more depth for analysis.

Quantitative data is further divided into:

a) Discrete Data: Discrete data includes **countable values**. These are often whole numbers and usually answer questions like "how many?" *Examples*: Number of students in a class: 25 ; Number of books in a library: 5,000 ; Number of children in a family: 2, 3, 4

Discrete values are separate and distinct. You can't have 2.5 students or 3.8 cars.

b) Continuous Data: Continuous data can take **any value within a range**, and it's often used to measure things. It includes decimals and fractions. *Examples*: Height: 172.5 cm ; Weight: 68.2 kg ; Temperature: 36.6°C.

Since continuous data can take infinite values within a range, it provides a lot of flexibility for analysis.

Like *Nominal and Ordinal data* under the category of Qualitative data, we are also having the *Interval and Ratio data* under the category of Quantitative data.

The **Interval data** is numerical data where the difference between values is meaningful, and intervals are equal. However, it lacks a true zero point. This means that while you can add and subtract interval data, you cannot multiply or divide it meaningfully in terms of ratios. *Examples*: Temperature in Celsius or Fahrenheit ; IQ scores ; Dates on a calendar

Let's take temperature as an example. The difference between 10°C and 20°C is the same as between 20°C and 30°C—that is, 10°C. But you can't say 20°C is "twice as hot" as 10°C, because 0°C doesn't represent the absence of temperature—it's just another point on the scale.

The **Ratio data** is also numerical and has all the properties of interval data, but with a meaningful or absolute zero point. Because of this, all mathematical operations, including multiplication and division, are valid. *Examples*: Height (e.g., 180 cm) ; Weight (e.g., 70 kg) ; Age (e.g., 25 years) ; Income (e.g., ₹50,000)

With ratio data, you can say that a person who weighs 60 kg is twice as heavy as someone who weighs 30 kg. This is because zero weight means no weight, making the scale absolute.

Now let's look at a more precise classification based on the level of measurement, which includes Nominal, Ordinal, Interval, and Ratio data types. This classification helps determine what types of statistical techniques you can use.

Levels of Measurement: The Four Data Types

1. Nominal – Categories without order (qualitative)
2. Ordinal – Categories with order but no fixed interval (qualitative)
3. Interval – Numeric data with equal spacing, but no true zero (quantitative)
4. Ratio – Numeric data with equal spacing and a true zero (quantitative)

Important:

- It is to be noted that the interval data is almost always continuous. But the Ratio data can be either discrete or continuous, depending on whether you're counting or measuring something.
- Interval and Ratio data are two subtypes of quantitative (numerical) data. Both deal with numbers and allow for mathematical operations, but they differ in one key way: ratio data has a true zero point, while interval data does not.
- Both ratio and interval data can be either continuous or discrete, but they are most commonly associated with continuous data.

In this section we learned that understanding data types—nominal, ordinal, interval, and ratio—is a fundamental skill in data analysis. Qualitative data helps classify and group information, while quantitative data allows for precise measurement and mathematical analysis. The level of measurement determines what kind of analysis you can perform. By correctly identifying your data type, you ensure that your analysis is appropriate, accurate, and meaningful.

Whether you're just starting out or working on a complex project, always begin by asking: What type of data am I working with? The answer will guide the rest of your analysis.

☞ **Check Your Progress 3**

Q1. Explain the four levels of measurement in data analysis i.e. Nominal, Ordinal, Ratio & interval data. Why are they important?

……………………………………………………………………………………………

……………………………………………………………………………………………

1.3.3 FREQUENCY DISTRIBUTION OF A VARIABLE

A frequency distribution is a powerful tool that transforms raw data into an organized format, making it easier to analyse. Whether you use a simple table for small datasets or a grouped one for larger values, frequency distributions help identify patterns, trends, and key summary statistics like the most frequent (modal) class, range, and central values.

In short, frequency distribution is your first step toward understanding and visualizing data effectively.

When you collect a set of raw data, it can be messy and hard to understand at first glance—especially if the number of observations is large. To make sense of such data, one of the first steps in descriptive data analysis is to organize it into a frequency distribution.

A frequency distribution is a way to arrange data values to show how often each value or range of values (called a class) occurs. This helps reveal patterns, trends, and how the values are spread across the dataset.

Imagine you conducted a survey among 30 students to find out how many hours they studied in a week. The responses might look like this:

[4, 7, 6, 5, 4, 9, 10, 7, 6, 4, 8, 7, 6, 5, 5, 4, 8, 9, 6, 5, 6, 4, 7, 6, 5, 10, 8, 7, 6, 9]

Looking at this list directly can be overwhelming. It’s hard to draw conclusions from such raw data.

Let’s now understand how to create a Discrete F*requency Distribution* from the data above.

Step 1: Identify the data range:

Minimum value = 4
Maximum value = 10

So, we’ll group the data based on the hours studied, from 4 to 10.

Step 2: Count how many times each value occurs

Study Hours (x)	Frequency (f)
4	5
5	5
6	7
7	5
8	3
9	3
10	2

This table shows how many students studied for each number of hours.

From this table, you can quickly observe:

- Most students (7) studied for 6 hours.
- Fewest students (2) studied for 10 hours.
- The data is spread around 4–10 hours.

This kind of representation makes it easy to analyse the data and identify central tendencies (like mean or median), spot patterns, or decide the next steps in your analysis (e.g., plotting a graph or calculating measures of spread).

In the above example the dataset is quite limited so the creation of frequency distribution is quite simple and is achieved in just 2 steps, as discussed above.

But, If the data range is wide or there are too many individual values, we can group them into class intervals i.e. in case of larger datasets, we need to opt for *Grouped Frequency Distribution*, discussed below

Let's say you surveyed 50 people about their monthly spending on groceries in ₹ (rupees), and the values range from ₹500 to ₹2000. Instead of listing each value, you can group them like this:

Monthly Spending (₹)	**Frequency**
500–799	6
800–1099	10
1100–1399	15
1400–1699	12
1700–1999	7

This is a *grouped frequency distribution*. It's useful when individual values are not as important as the **range** they fall into.

To understand the Grouped Frequency Distribution in a better way, let's understand how to create it by an example in a stepwise manner.

Let's go through the steps one by one using an example:

Example Scenario: Suppose you conducted a test and collected the **marks obtained by 40 students** out of 100. Here's the (unordered) data:

72, 55, 48, 89, 90, 68, 77, 52, 69, 73, 64, 83, 95, 70, 66, 71, 60, 58, 85, 80, 74, 56, 63, 61, 86, 67, 84, 76, 54, 59, 62, 65, 79, 78, 50, 81, 75, 88, 53, 57

Step 1: Identify the Range

- Find the **lowest** and **highest** values.
  - Lowest = 48
  - Highest = 95
- **Range = 95 − 48 = 47**

Step 2: Decide the Number of Classes

A good rule of thumb is to use between **5 and 10 classes**, depending on the data size. For 40 values, let's go with **8 classes**.

Step 3: Calculate Class Width ; = Range/Number of Classes =47/8 ≈5.875

We round this up to the next whole number: **6**

So each class will cover an interval of **6 marks**.

Step 4: Create Class Intervals
Start from the lowest value (rounded down, if needed). Let's start at **48**. The intervals will be:

- 48–53
- 54–59
- 60–65
- 66–71
- 72–77
- 78–83
- 84–89
- 90–95

Note: The intervals are **continuous**, meaning there is no gap between them.

Step 5: Frequency Distribution Table (Completed)
Now, go through the dataset and count how many values fall into each class.

Class Interval	Frequency
48–53	4
54–59	6
60–65	6
66–71	6
72–77	6
78–83	5
84–89	5
90–95	2
Total	**40**

Each frequency value represents how many student marks fall within that class interval. The total (sum of frequencies) equals 40, matching the total number of students in the dataset.

Steps of Creating a continuous frequency distribution involves:

1. Finding the range of the data.
2. Deciding on the number of class intervals.
3. Calculating the class width.
4. Creating non-overlapping, continuous intervals.
5. Tallying the frequency of data in each class.

When studying frequency distribution, you may come across a few related terms like Relative and cumulative frequencies. Let's further extend our discussion, and understand the concept of Relative Frequency and Cumulative Frequency.

- **Relative Frequency**: The proportion or percentage of the total that each class represents.
    - Formula: Relative Frequency = (Class Frequency / Total Frequency)

- **Cumulative Frequency**: The running total of frequencies up to a certain class.
    - Useful for percentile and median calculations.

Enhanced Frequency Distribution Table
Given:
Total number of students = **40**
We'll compute:

- **Relative Frequency** = Frequency ÷ Total
- **Cumulative Frequency** = Running total of frequencies

Class Interval	Frequency (f)	Relative Frequency (f/40)	Cumulative Frequency
48–53	4	4 / 40 = 0.10	4
54–59	6	6 / 40 = 0.15	10
60–65	6	6 / 40 = 0.15	16
66–71	6	6 / 40 = 0.15	22
72–77	6	6 / 40 = 0.15	28
78–83	5	5 / 40 = 0.125	33
84–89	5	5 / 40 = 0.125	38
90–95	2	2 / 40 = 0.05	40
Total	40	1.00	—

Here, in the above table the Relative Frequency shows the proportion or percentage of students in each class. Whereas the Cumulative Frequency shows how many students scored up to and including a given class range.

1.3.4 GRAPHICAL REPRESENTATION OF FREQUENCY DISTRIBUTION

Graphical representation of frequency distribution is a method used to visually summarize and interpret data. Instead of examining rows and columns of numbers in a table, graphs provide a clear picture of how data is spread across different values or intervals. This makes it easier to identify trends, patterns, outliers, and the overall shape of the distribution. Several types of graphs are commonly used depending on the nature of the data and the purpose of the analysis.

One of the most frequently used graphs for continuous numerical data is the **histogram**. A histogram consists of adjacent (touching) bars where the horizontal axis (X-axis) represents the class intervals, and the vertical axis (Y-axis) shows the frequencies—the number of data values that fall within each interval. Each bar's height corresponds to the frequency for that class. For example, if students' exam marks are grouped into intervals like 40–49, 50–59, 60–69, and so on, a histogram can quickly show how many students fall into each score range. Since the data is continuous, there are no gaps between bars.

Steps to Create a Histogram for Grouped (Continuous) Data

Creating a histogram from grouped data involves a systematic process that helps visualize the distribution of a dataset. Let's walk through the steps using the example of marks obtained by 40 students, distributed across class intervals.

Step 1: Organize the Data

The first step is to prepare a frequency distribution table. This table should include clearly defined, non-overlapping class intervals and their corresponding frequencies. For example:

Class Interval	Frequency
48–53	4
54–59	6
60–65	6
66–71	6
72–77	6
78–83	5
84–89	5
90–95	2

Make sure that the class intervals are **continuous**, meaning there are no gaps between them, and that the frequencies are accurately counted from the raw data.

Step 2: Draw the Axes

Next, draw two perpendicular axes to create your graph. The **horizontal axis (X-axis)** represents the class intervals, such as 48–53, 54–59, and so on. These intervals should be equally spaced and in ascending order. The **vertical axis (Y-axis)** represents the frequency or the number of observations in each class. Choose a scale for the Y-axis that comfortably accommodates the highest frequency in your table; in this case, the highest frequency is 6.

Step 3: Draw Bars

Now, for each class interval, draw a bar such that:

- The **width** of each bar is equal (because our class intervals are of equal width — 6 marks in this example).

- The **height** of each bar corresponds to the **frequency** for that interval.

For example, the bar for the 48–53 interval should have a height of 4, while the bar for 54–59 should rise to 6, and so on. Since the data is continuous, ensure that the bars are **adjacent to each other without any gaps**.

Step 4: Label the Histogram

Once the bars are drawn, label your chart clearly. Give the histogram a descriptive **title**, such as "Histogram of Student Marks." The **X-axis** should be labeled as "Class Intervals" or "Marks," and the **Y-axis** should be labeled as "Frequency" or "Number of Students." This labeling helps in interpreting the graph easily.

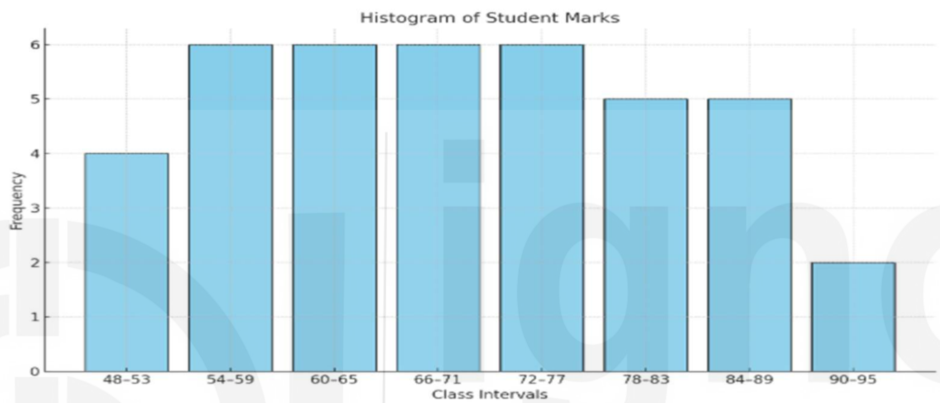

Here's the histogram representing the distribution of student marks based on the continuous frequency distribution table:

Step 5: Analyze the Histogram

After constructing the histogram, analyze its shape and pattern. Look for where most of the values are concentrated, which indicates the **central tendency** of the data. You can also assess whether the distribution is **symmetrical**, **skewed** to the left or right, or has multiple peaks (**bimodal**). Additionally, observe any **outliers** or gaps in the data. This visual analysis can provide quick insights into the nature of the dataset.

Another useful graphical tool is the **frequency polygon**, which is similar in purpose to a histogram but uses a line graph instead of bars. In a frequency polygon, midpoints of each class interval are plotted on the X-axis, and frequencies are plotted on the Y-axis. These points are then connected with straight lines to form a polygon. Often, additional points with zero frequency are added at the beginning and end to close the polygon. This graph is especially helpful when comparing multiple distributions on the same plot, making it easier to identify differences in shape and spread.

Steps to Create a Frequency Polygon

Step 1: Prepare the Frequency Distribution Table

Create a table with class intervals and their corresponding frequencies. Also, compute the **midpoints** of each class interval.

Class Interval	Frequency (f)	Midpoint (x)
48–53	4	50.5
54–59	6	56.5
60–65	6	62.5
66–71	6	68.5
72–77	6	74.5
78–83	5	80.5
84–89	5	86.5
90–95	2	92.5

Midpoint is calculated using the formula = (Lower Limit + Upper Limit)/2

Step 2: Add Extra Midpoints at Both Ends (Optional but Recommended)

To close the polygon at the baseline, add one class interval before the first and one after the last with **zero frequency**.

- Before 48–53: midpoint = 44.5, frequency = 0
- After 90–95: midpoint = 98.5, frequency = 0

Now, the extended list of midpoints and frequencies becomes:

Midpoint (x)	Frequency (f)
44.5	0
50.5	4
56.5	6
62.5	6
68.5	6
74.5	6
80.5	5
86.5	5
92.5	2
98.5	0

Step 3: Plot the Points

On graph paper or using software (e.g., Excel, Python, or any graphing tool):

- Plot each point with a midpoint **on the X-axis** and **frequency on the Y-axis**.
- For example, plot the points (50.5, 4), (56.5, 6), and so on.

Step 4: Connect the Points with Line Segments

- Draw straight lines to connect the plotted points in order from left to right.
- Make sure to include the starting and ending points with zero frequency to anchor the polygon to the baseline.

Step 5: Label and Title the Graph

- Add labels for axes: **X-axis**: Class Midpoints **Y-axis**: Frequency
- Title the graph: *Frequency Polygon for Student Marks*

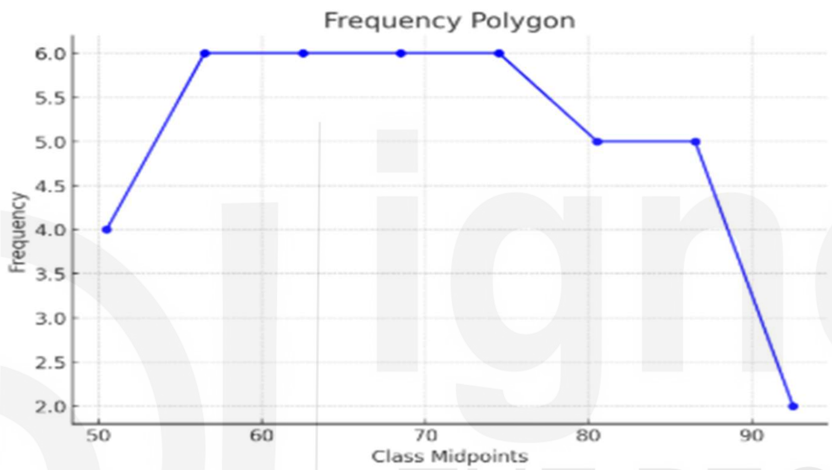

What Does a Frequency Polygon Show? : The Frequency Polygon allows you to see the **shape** of the distribution more clearly than a histogram, especially when comparing multiple datasets. Peaks and symmetry are easier to detect.

Now, For analyzing cumulative trends, the **ogive or cumulative frequency curve** is an effective choice. There are two types of ogives: the "less than" ogive, which plots cumulative frequencies against the upper boundaries of the class intervals, and the "more than" ogive, which uses the lower boundaries. For example, if you want to find how many students scored less than 70 marks, the "less than" ogive can give you a quick visual answer. Ogives are particularly useful for locating the median, quartiles, and percentiles in a data set.

Steps to Create an Ogive (Cumulative Frequency Curve)

Step 1: Prepare the Frequency Distribution Table

Use the same frequency table used for the histogram and compute the **cumulative frequencies**.

Class Interval	Upper Boundary	Frequency (f)	Cumulative Frequency (Less Than Type)
48–53	53	4	4
54–59	59	6	10
60–65	65	6	16
66–71	71	6	22
72–77	77	6	28
78–83	83	5	33
84–89	89	5	38
90–95	95	2	40

Cumulative Frequency is the running total of frequencies up to each class.
For example: 4 + 6 = 10, 10 + 6 = 16, and so on.

Step 2: Plot the Points
Each point is plotted with:

- **Upper Class Boundary on the X-axis**
- **Cumulative Frequency on the Y-axis**

So, plot the points:
(53, 4), (59, 10), (65, 16), (71, 22), (77, 28), (83, 33), (89, 38), (95, 40)

You may optionally begin from a lower value such as (47, 0) to anchor the curve to the X-axis.

Step 3: Connect the Points Smoothly

- Join the plotted points using a smooth curve or straight lines.
- Ensure the curve starts from zero and rises as you move right, because cumulative frequency always increases or remains constant.

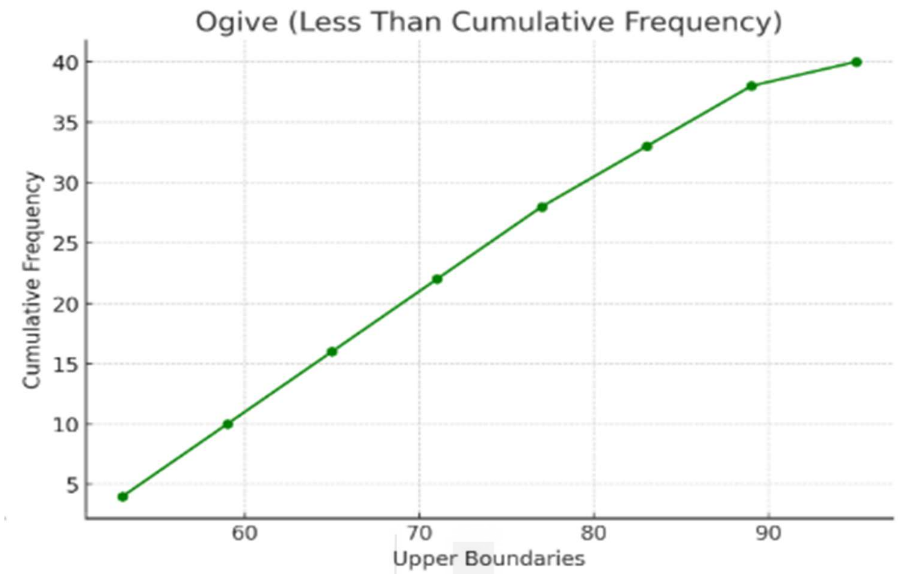

Step 4: Label and Title the Graph

- **X-axis**: Upper Boundaries of Class Intervals **Y-axis**: Cumulative Frequency
- **Title**: *Ogive (Less Than Cumulative Frequency Curve) for Student Marks*

What Does Ogive (Cumulative Frequency Curve) Show? : The Ogive helps you to Understand how data accumulates across intervals. Also it estimates the median, quartiles, and percentiles. And helps to quickly answer questions like "How many students scored less than 70?"

To further analyze distribution and detect spread or outliers, one can use a box plot, also known as a **box-and-whisker plot**. This graphical representation displays the five-number summary of a dataset: minimum, first quartile (Q1), median (Q2), third quartile (Q3), and maximum. The box represents the interquartile range (middle 50% of the data), and the lines extending from the box (whiskers) indicate variability outside the upper and lower quartiles. Any values lying outside these whiskers may be considered outliers. For instance, a box plot of students' marks can show not only the median score but also whether any students scored exceptionally low or high.

Steps to Create a Box-and-Whisker Plot (Box Plot)

A Box-and-Whisker Plot visually displays the **distribution**, **spread**, and **skewness** of a dataset through five summary statistics:

- **Minimum**
- **First Quartile (Q1)**
- **Median (Q2)**
- **Third Quartile (Q3)**
- **Maximum**

Step 1: Reconstruct the Raw Data from Grouped Data

Since a Box Plot needs raw data or an approximation of it, reconstruct a **pseudo dataset** using the **midpoints** of each class interval and the corresponding **frequency**. Here's how:

Class Interval	Frequency	Midpoint
48–53	4	50.5
54–59	6	56.5
60–65	6	62.5
66–71	6	68.5
72–77	6	74.5
78–83	5	80.5
84–89	5	86.5
90–95	2	92.5

Total observations (N) = 40

Construct Pseudo Raw Data: i.e. expand the data based on the frequency:

50.5 × 4 → positions 1–4
56.5 × 6 → positions 5–10
62.5 × 6 → positions 11–16
68.5 × 6 → positions 17–22
74.5 × 6 → positions 23–28
80.5 × 5 → positions 29–33
86.5 × 5 → positions 34–38
92.5 × 2 → positions 39–40

Step 2: Arrange the Data in Ascending Order

Arrange the values from smallest to largest. Since the data is already binned using midpoints, and each midpoint is repeated according to its frequency, it's already structured to represent an ordered dataset.

Step 3: Find the Five-Number Summary i.e. Min., Max, Q1, Q2, and Q3

Using the pseudo raw data given above:

- **Minimum**: Smallest value = 50.5
- **Maximum**: Largest value = 92.5
- **Median (Q2)**: Middle value (20th and 21st observation average) ≈ 68.5
- **First Quartile (Q1)**: 25th percentile (10th observation) ≈ 58.0
- **Third Quartile (Q3)**: 75th percentile (30th observation) ≈ 80.5

So, the five-number summary is:

- **Min** = 50.5
- **Q1** = 58.0
- **Median (Q2)** = 68.5
- **Q3** = 80.5
- **Max** = 92.5

Let's understand how to determine the Quartiles Q1,Q2 and Q3 ; especially in the context of a dataset like the one we've been working with (40 student marks). You can use the same method for any dataset.

So, What Are Quartiles? we have 3 quartiles in general Q1, Q2, and Q3; their respective meanings are as follows:

- **Q1 (First Quartile)**: The value that separates the **lowest 25%** of the data.
- **Q2 (Median)**: The middle value; separates the **lower 50%** from the **upper 50%**.
- **Q3 (Third Quartile)**: The value that separates the **lowest 75%** from the top 25%.

Now, we understand How to Determine Q1, Q2 and Q3 (Manually)

Assume you have **N = 40** data points, arranged in **ascending order**.

Step A: Find the Positions

- Q1 position = (N+1)/4=(41)/4=10.25 th item
- Q2 position = (N+1)/2 = (41)/2 = 20.5th item
- Q3 position = 3(N+1)/4 =3(41)/4 =30.75 th item

These are **positions**, not values yet.

Step B: Estimate the Quartile Values

To get the values at 10.25th and 30.75th positions, you interpolate between the actual values at the 10th, 11th (for Q1), 20th, 21st (for Q2), and 30th, 31st (for Q3).

Suppose the ordered pseudo data looks like:

(Refer to Pseudo Raw Data given above)

50.5 × 4 → positions 1–4
56.5 × 6 → positions 5–10
62.5 × 6 → positions 11–16
68.5 × 6 → positions 17–22
74.5 × 6 → positions 23–28
80.5 × 5 → positions 29–33

Finding Quartile-1(Q1), from the pseudo raw data given we can see that there are

4 values of 50.5

6 values of 56.5 → positions 5 to 10
6 values of 62.5 → positions 11 to 16

So the **10th value = 56.5**, and **11th value = 62.5**
Thus, $Q1=56.5+0.25\times(62.5-56.5)=56.5+1.5=58.0$

(Q2)Median position$=(N+1)/2=41/2=20.5$th value
We want the **value at the 20.5th position** in the ordered pseudo data.

Now from the pseudo raw data given we can see that both
20^{th} value = 68.5 and 21^{st} value = 68.5

Thus, $Q2=(68.5+68.5)/2=68.5$
Since the 20.5th value lies between 68.5 and 68.5, the median is **68.5**

Similarly, **30th value = 80.5**, **31st value = 86.5**
$Q3=80.5+0.75\times(86.5-80.5)=80.5+4.5=85.0$

So, **Q1 ≈ 58.0**, **Q2 ≈ 68.5** and **Q3 ≈ 85.0** in this case.

Step 4: Draw the Box-and-Whisker Plot

- **Draw a horizontal number line** that covers the range of your data (e.g., from 45 to 100).
- **Mark the five-number summary** values on the line.
- **Draw a rectangular box** from Q1 to Q3 .
- **Draw a vertical line** inside the box at the median .
- **Draw whiskers** (horizontal lines) from:
  - Min to Q1
  - Q3 to Max

Step 5: Label the Plot

- Label the five-number summary on the axis.
- Title your plot appropriately: *"Box-and-Whisker Plot of Student Marks"*

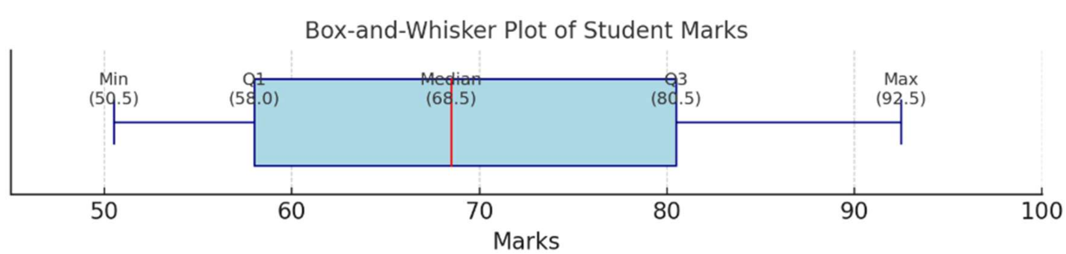

Here is the **Box-and-Whisker Plot** based on the student marks data. It visually displays the five-number summary:

Minimum: 50.5 ; **Q1 (First Quartile)**: 58.0 ; **Median (Q2)**: 68.5 ;
Q3 (Third Quartile): 80.5 ; **Maximum**: 92.5

This plot helps you understand the **spread**, **center**, and **symmetry** of the data, as well as spot potential outliers (though none appear here).
In short, The box plot helps to:

- Visualize the **spread and central tendency**.
- Detects **symmetry** or **skewness**.
- Identify **potential outliers** (values far from the whiskers).
- Quickly compare multiple datasets if needed.

Each of these graphical methods serves a specific purpose and is chosen based on the data type and the kind of insight the analyst wants to gain. Histograms, frequency polygons, and ogives are best suited for continuous numerical data. Bar charts and pie charts are ideal for categorical data. Box plots provide a deeper look into data distribution and help identify variability and outliers.

When dealing with categorical or discrete data, bar charts are commonly used. In a **bar chart**, each category is represented by a separate bar, and unlike a histogram, the bars do not touch. The X-axis lists the categories (such as blood groups: A, B, AB, and O), while the Y-axis represents the frequency of each category. For example, a bar chart can display how many people in a survey belong to each blood group. Thus, Bar charts are useful for showing comparisons among **discrete categories** using rectangular bars. The height (or length) of each bar represents the value of the corresponding category.

Example: Create a bar chart for the data collected under a survey which shows the number of students preferring different programming languages:

Language	Number of Students
Python	30
Java	20
C++	10
JavaScript	15
R	5

Stepwise – Solution to create Bar chart:

Step 1: Choose Axes : Decide which variable will go on each axis:

- **X-axis**: Categories (e.g., programming languages)
- **Y-axis**: Frequencies or counts (e.g., number of students)

Step 2: Draw the Axes and Label Them : Draw two perpendicular lines:

- Label the X-axis with the category names (Python, Java, etc.).
- Label the Y-axis with a numerical scale (0 to 35, for example, with intervals of 5).

Step 3: Draw Bars : Draw a **separate vertical bar for each category**, ensuring that:

- Bars are of **equal width**
- There is **equal spacing** between bars
- Heights correspond exactly to the values

For instance:

- Python bar will be tallest (30 units high)
- R bar will be shortest (5 units high)

Step 4: Add Title and Color (Optional) : Title the chart: *"Students' Preferred Programming Languages"*. Optionally, use different colors for bars to enhance readability.

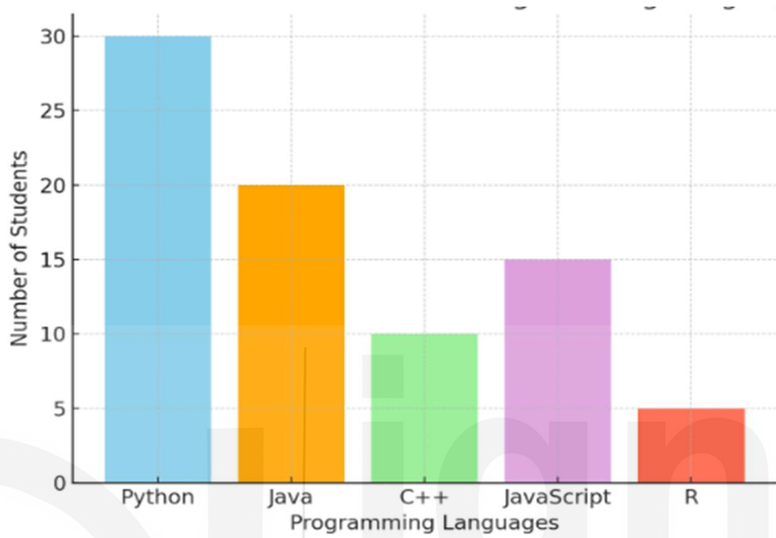

Another widely used representation for categorical data is the **pie chart**, which shows data as slices of a circular pie. Each slice represents a category, and the size of the slice corresponds to the proportion of that category in the whole dataset. The angle of each slice is calculated by the formula: (Category Frequency / Total Frequency) × 360°.

So, a pie chart is a circular graph divided into slices to illustrate numerical proportions. It is best used when you want to show **percentage contributions** of each category to a whole.

Let's understand how to generate a pie chart, the stepwise solution to the generation of pie-chart for the above example is as follows

Stepwise – Solution to create Pie chart:

Step 1: Collect Data: Use the same example as above:

Language	Number of Students
Python	30
Java	20
C++	10
JavaScript	15
R	5

Step 2: Calculate Total: Add up all the values to get the **total number of students**.

Total=30+20+10+15+5=80

Step 3: Calculate Percentage or Angle for Each Category :

To calculate the angle for each slice of the pie:

Angle=(Value/Total)×360∘

For example:

- Python: (30/80)×360=135∘
- Java: (20/80)×360=90∘
- C++: (10/80)×360=45∘
- JavaScript: (15/80)×360=67.5∘
- R: (5/80)×360=22.5∘

Step 4: Draw the Circle and Plot Slices

- Draw a circle using a compass or software.
- Use a protractor to mark each angle from the center.
- Label each slice with its category name and percentage.

Step 5: Add Title and Color (Optional)

Title it something like *"Language Preferences Among Students"*

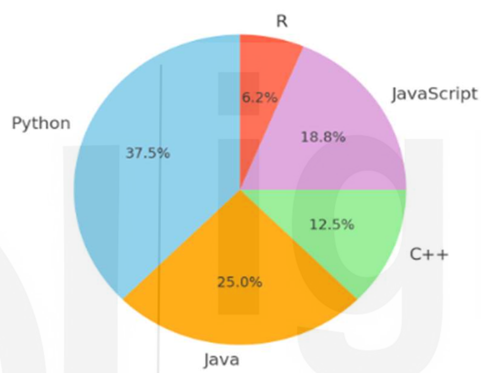

Both Bar Chart and Pie-Chart are powerful tools for visualizing categorical data

- Use **bar charts** to compare discrete categories clearly and directly.
- Use **pie charts** to show the **proportion** each category contributes to the total.

Understanding these graphical tools is essential for effective data analysis, as they transform raw numbers into meaningful visual stories that are much easier to interpret and communicate.

1.3.5 MEASUREMENT OF CENTRAL TENDENCY

Measures of central tendency help us summarize a large amount of data using a single representative value. This makes it easier to understand and analyse the overall trend in the data. For instance, it's not practical to remember the individual incomes of millions of people in India. However, knowing the **average income** gives us a general idea of the income level of the population as a whole.

These measures also make it easier to **compare different data sets**. For example, the **average sales** in April can be compared with those of previous months to understand whether performance has improved or declined.

A **good measure of central tendency** is one that effectively represents a dataset with a single, meaningful value. To be truly useful, such a measure should have certain key characteristics. First, it should be **easy to understand**, so that people without technical backgrounds can interpret it. It should

also be **simple to calculate**, using basic mathematical operations. A reliable measure must be **based on all the observations** in the dataset, not just a few selected ones, to ensure it truly reflects the data. Additionally, the measure should be **uniquely defined**, meaning that it always gives the same result for a given dataset. It should also allow for **further algebraic treatment**, so it can be used in more complex analyses and calculations. Finally, it should be **resistant to extreme values (outliers)**; that is, it should not be heavily influenced by unusually large or small numbers that might distort the overall picture.
In practice, several types of measures of central tendency are commonly used in fields like **business, economics, and industry**. These include:

- **Arithmetic Mean** – the most common average.
- **Weighted Arithmetic Mean** – used when different values have different levels of importance.
- **Median** – the middle value when the data is arranged in order.
- **Quartiles** – values that divide the data into equal-sized intervals, such as quartiles or percentiles.
- **Mode** – the value that appears most frequently.
- **Geometric Mean** – useful in cases involving rates of growth or ratios.
- **Harmonic Mean** – often applied in averaging ratios or rates, like speed.

Understanding these measures and their characteristics helps in selecting the right one depending on the nature and context of the data.

Let's Begin with **Calculation of the Central Tendency for the Ungrouped Data:**

1. Arithmetic Mean: The arithmetic mean, often referred to as the mean or average, is the most widely used and easily understood measure of central tendency. In statistical terms, the word "average" generally refers to any value that represents the center of a data set. Specifically, the arithmetic mean is calculated by adding all the values in the dataset and then dividing the total by the number of observations. It can be expressed symbolically as follows:

$$\underline{X} = \frac{\Sigma x}{N}$$

Here, $\sum$**x** represents the total of all observed values, while **N** stands for the number of observations in the dataset.
Let's take an example to understand this better. Suppose The marks obtained by a student in five subjects are: 60, 70, 80, 90, 100

To find the **arithmetic mean**, we add all the marks and divide by the number of students

Thus, the arithmetic mean is $\underline{X}$ = (60 + 70 + 80 + 90 + 100) / 5 = 400 / 5 = 80 i.e. the average marks are 80

2. **Weighted Arithmetic Mean**: Used when some values carry more importance or weight than others.

Example:
Grades and credit weights: Math (90, weight 4), Science (80, weight 3), English (70, weight 2)
Weighted Mean = (90×4 + 80×3 + 70×2) / (4 + 3 + 2) = 740 / 9 ≈ 82.22

3. **Median:** The median is the middle value in an ordered dataset.
Example 1: Data: 12, 15, 20, 25, 30 → Median = 20
Example 2: Data: 10, 20, 30, 40 → Median = (20 + 30)/2 = 25

4. **Quartiles:** Quartiles divide ordered data into four equal parts.
Example: Data: 10, 20, 30, 40, 50, 60, 70, 80
Q1 = 25, Q2 (Median) = 45, Q3 = 65

5. **Mode:** The mode is the value that appears most frequently.
Example: Data: 3, 4, 4, 5, 6, 6, 6, 7 → Mode = 6

6. **Geometric Mean**: Useful for multiplicative data like growth rates.
Example: Growth rates: 10%, 20%, 30% → Convert to factors: 1.10, 1.20, 1.30
Geometric Mean = $\sqrt[3]{(1.10\times1.20\times1.30)} \approx \sqrt[3]{1.716} \approx 1.19$ or 19% growth per year

7. **Harmonic Mean:** Best used when averaging rates like speed.
Example: Speed to and from a place: 60 km/h and 40 km/h
Harmonic Mean = 2×60×40 / (60 + 40) = 4800 / 100 = 48 km/h

So far, we have learned how to calculate the arithmetic mean for **ungrouped data**. Now, let's understand how this changes when dealing with **grouped data**. In a frequency distribution, the data values are grouped into class intervals. Since we don't know the exact values within each class, we use the **midpoint** of each class interval as a representative value for that class. Based on this approach, the **arithmetic mean for grouped data** is calculated using these midpoints. It is defined as follows:

$$\underline{X} = \frac{\Sigma f . x}{N}$$

Where X is midpoint of various classes, f is the frequency for corresponding class and N is the total frequency, i.e. N = $\sum f$.

1.3.6 MEASUREMENT OF DISPERSION

In above section 1.3.5 we learned about the measure of central tendency of data, now we extend our discussion towards the understanding of the dispersion of data and how it is measured. It is to be noted that the degree to which numerical data tend to spread about an average value is called the *variation or dispersion of data.*
Actually, there are two basic kinds of a measure of dispersion

(i) Absolute measures and
(ii) Relative measures.

The *absolute measures of dispersion* are used to measure the variability of a given data expressed in the same unit, while the *relative measures* are used to compare the variability of two or more sets of observations.

Following are the different measures of dispersion:

1. Range
2. Quartile Deviation

3. Mean Deviation
4. Standard Deviation and Variance

*Before extending our discussion for the different measures of dispersion, firstly we need to understand the **significance of Measures of dispersion**, actually they are needed for the following purposes:*

- They show how much the data values vary or spread around the average.
- They help determine how reliable or representative an average value is.
- Less variation means the average represents the data well; more variation means it may be unreliable.
- They help identify the causes of variation so that it can be controlled (e.g., quality control in production).
- They allow comparison of two or more data sets in terms of consistency or uniformity.
- Smaller dispersion indicates greater uniformity and consistency in data.
- They form the basis for many statistical methods such as correlation, hypothesis testing, analysis of variance, and quality control techniques.

Further The characteristics of a good measure of dispersion are similar to those of a good measure of central tendency. Thus, it is reiterated that an ideal measure of dispersion should have the following qualities:

- It should be easy to understand.
- It should be simple to calculate.
- It should be clearly and precisely defined.
- It should consider all observations in the dataset.
- It should be suitable for further algebraic or mathematical analysis.
- It should remain stable across different samples.
- It should not be excessively influenced by extreme values.

Let's begin our discussion on the different measures of dispersion i.e. Range, Quartile Deviation, Mean Deviation and Standard Deviation and Variance; range is the first one to begin with

1. Range: Range is the simplest measure of dispersion. It represents the difference between the maximum and minimum values in a dataset.

Range = Maximum value (X_{max}) – Minimum value (X_{min})

Example: for the Data: 5, 10, 15, 20, 25; the maximum value is 25 and the minimum value is 5 thus the Range = 25 – 5 = 20

The Merits and Demerits of Range are as follows :

Merits of Range

1. It is the simplest to understand;

2. It can be visually obtained since one can detect the largest and the smallest observations easily and can take the difference without involving much calculations; and

3. Though it is crude, it has useful applications in areas like order statistics and statistical quality control.

Demerits of Range

1. It utilizes only the maximum and the minimum values of variable in the series and gives no importance to other observations;
2. It is affected by fluctuations of sampling;
3. It is not very suitable for algebraic treatment;
4. If a single value lower than the minimum or higher than the maximum is added or if the maximum or minimum value is deleted range is seriously affected; and
5. Range is the measure having units of the variable and is not a pure number.

That's why sometimes coefficient of range is calculated by

$$\text{Coefficient of Range: } \frac{(Xmax - Xmin)}{(Xmax + Xmin)}$$

2. Quartile Deviation: Quartile Deviation measures dispersion using the middle 50% of the data. It is also called the Semi-Interquartile Range. (Q3 – Q1) gives the inter quartile range. The semi inter quartile range which is also known as Quartile Deviation (QD) is given by

$$\text{Quartile Deviation (QD)} = (Q_3 - Q_1) / 2$$

Relative measure of Q.D. known as Coefficient of Q.D. and is defined as

$$\text{Coefficient of QD} = \frac{(Q3 - Q1)}{(Q3 + Q1)}$$

Example-1: For the Data: 2, 4, 6, 8, 10, 12, 14; the first Quartile $Q_1 = 4$, and the third Quartile are $Q_3 = 12$. Therefore, QD = (12 – 4) / 2 = 4

Example-2: Find the Quartile Deviation for the following data.

Class Interval	Frequency
0–10	3
10–20	5
20–30	7
30–40	9
40–50	4

Solution:

Class Interval	Frequency (f)	Cumulative Frequency (cf)
0 – 10	3	3
10 – 20	5	8
20 – 30	7	15
30 – 40	9	24
40 – 50	4	28

Step 1: Calculate Total Frequency (N)= 3+5+7+9+4 = 28

Step 2: Determine Quartile Positions for Q1 = N/4 = 7 and for Q3 = 3N/4 = 21

Step 3: Determine Quartile Classes

Q_1 class → Cumulative frequency just greater than 7 → 10–20

Q_3 class → Cumulative frequency just greater than 21 → 30–40

i.e. Q1 class = 10–20 and Q3 class = 30–40

Step 4: We know the Formula for Quartile (grouped data) is Q = L + ((kN/4 – cf) / f) × h

Where:

L= lower limit of quartile class

cf= cumulative frequency before quartile class

f= frequency of quartile class

h= class width

Step 5: Now Calculate Q1 and Q3

Calculate Q_1: For class 10–20 we have $L = 10, cf = 3, f = 5, h = 10$

Therefore Q1 = 10 + ((7 – 3) / 5) × 10 = 18

Calculate Q_3: For class 30–40 we have $L = 30, cf = 15, f = 9, h = 10$

Therefore Q3 = 30 + ((21 – 15) / 9) × 10 = 36.67

Step 6: Finally calculate Quartile Deviation (QD) = (Q3 – Q1) / 2

Quartile Deviation (QD) = (36.67 – 18) / 2 = 9.34

3. Mean Deviation:

Mean deviation is defined as the average of the sum of the absolute values of deviations from any arbitrary value such as mean, median, mode, etc. It is often suggested to calculate mean deviation from the median because it gives the least value when measured from the median.

The deviation of an observation X_i from the assumed mean A is defined as: (xi − A)

Therefore, mean deviation can be defined as: MD = (1/n) Σ |xi − A|

Note : It is to be noted that the quantity Σ |xi − A| is minimum when A is the median.

Mean deviation from mean is defined as: MD = (1/n) Σ |xi − x̄|

Example: Given Data: 2, 4, 6, 8 has Mean = 5,

therefore, Mean Deviation = (3 + 1 + 1 + 3) / 4 = 2

Mean deviation from median is defined as: MD = $(1/n)\ \Sigma\ |x_i - \text{Median}|$

For a frequency distribution, the formulae are:

Mean Deviation from Mean: MD = $\Sigma f_i |x_i - \bar{x}| / \Sigma f_i$

Mean Deviation from Median: MD = $\Sigma f_i |x_i - \text{Median}| / \Sigma f_i$

where all symbols have their usual meanings.

Example: For the data given below,

x:	1	2	3	4	5	6	7
f:	3	5	8	12	10	7	5

Determine

(i) Mean deviation from mean and

(ii) Mean deviation from Median

Solution :

<u>(i) Mean Deviation from Mean</u>

Step 1: Calculate Mean

$\Sigma f = 50$

$\Sigma fx = (1\times3)+(2\times5)+(3\times8)+(4\times12)+(5\times10)+(6\times7)+(7\times5) = 3+10+24+48+50+42+35 = 212$

Therefore, Mean ($\bar{x}$) = $\Sigma fx / \Sigma f = 212 / 50 = 4.24$

Step 2: Calculate $|x - \bar{x}|$ and $f|x - \bar{x}|$

x	f	$\lvert x-4.24\rvert$	$f\lvert x-4.24\rvert$
1	3	3.24	9.72
2	5	2.24	11.20
3	8	1.24	9.92
4	12	0.24	2.88
5	10	0.76	7.60
6	7	1.76	12.32
7	5	2.76	13.80

Find $\Sigma f |x - \bar{x}| = 67.44$

Therefore, Mean Deviation from Mean MD(mean) = $\Sigma f|x-\bar{x}| / \Sigma f = 67.44 / 50 = 1.348$

<u>(ii) Mean Deviation from Median</u>

Step 1: Find Median

Compute Cumulative Frequency:

x:	1	2	3	4	5	6	7
f:	3	5	8	12	10	7	5
cf:	3	8	16	28	38	45	50

As $N = 50 \rightarrow N/2 = 25$

Thus, the 25th observation lies at $x = 4$
Hence, Median = 4

Step 2: Calculate |x − Median| and f | x – Median |

x	f	$\lvert x-4 \rvert$	$f\lvert x-4 \rvert$
1	3	3	9
2	5	2	10
3	8	1	8
4	12	0	0
5	10	1	10
6	7	2	14
7	5	3	15

As $\Sigma f \; \lvert x - \text{Median} \rvert = 66$

Mean Deviation from Median

$$MD(\text{median}) = \Sigma f \lvert x - \text{Median} \rvert / \Sigma f$$
$$= 66 / 50$$
$$= 1.32$$

Thus, we got, Mean Deviation from Mean = 1.348 and Mean Deviation from Median = 1.32

The mean deviation also possesses some merits and demerits; they are as follows:

Merits of Mean Deviation:

1. It utilizes all the observations.
2. It is easy to understand and calculate.
3. It is not much affected by extreme values.

Demerits of Mean Deviation:

1. Negative deviations are converted into positive values.
2. It is not suitable for algebraic treatment.
3. It cannot be calculated for open-end classes.

4. Standard Deviation and Variance

Variance and Standard Deviation are important measures of dispersion used in statistics and data science. They describe how much the data values deviate from the mean.

formally, Variance is the average of the squares of deviations of observations from the arithmetic mean, and Standard deviation is the positive square root of variance.

The formula for variance and standard deviation are as follows :

Formula (Ungrouped Data): Variance $(\sigma^2) = \Sigma (x - \bar{x})^2 / n$
Formula: (Ungrouped Data): Standard deviation $(\sigma) = \sqrt{[\Sigma (x - \bar{x})^2 / n]}$

Formula (Frequency Distribution): Variance $(\sigma^2) = \Sigma f (x - \bar{x})^2 / \Sigma f$

Formula (Frequency Distribution): Standard deviation $(\sigma) = \sqrt{[\Sigma f (x - \bar{x})^2 / \Sigma f]}$

Example (Ungrouped Data) : Determine the variance in the data: 2, 4, 6, 8

Solution:

Step 1: Calculate Mean ($\bar{x}$) = (2 + 4 + 6 + 8) / 4 = 5

Step 2: Find Deviations

x	$x - \bar{x}$	$(x - \bar{x})^2$
2	-3	9
4	-1	1
6	1	1
8	3	9

Calculate $\Sigma(x - \bar{x})^2 = 20$

Step 3: Find Variance $(\sigma^2) = \Sigma(x - \bar{x})^2 / n$

$\sigma^2 = 20 / 4 = 5$

If Variance = 5 then Standard Deviation $(\sigma) = \sqrt{5} \approx 2.24$

INTERPRETATION: Standard deviation shows that the data values deviate from the mean by about 2.24 units.

Above we have discussed how to calculate the variance for a single variable. If there are two or more populations and the information about the means and variances of those populations are available then we can obtain the combined variance of several populations.

If $n_1, n_2, \ldots, n_k$ are the sizes $x_1, x_2, \ldots, x_k$ are the means and $\sigma_1^2, \sigma_2^2, \ldots, \sigma_k^2$ are the variances of k populations, then the combined variance is given by the Formula for Variance of Combined Series:

$$\sigma^2 = \frac{[\, n_1(\sigma_1{}^2 + d_1{}^2) + n_2(\sigma_2{}^2 + d_2{}^2) + \ldots + n_k(\sigma_k{}^2 + d_k{}^2)\,]}{(n_1 + n_2 + \ldots + n_k)}$$

Where:

$d_i = \bar{x}_i - \bar{x}$

Mean of Combined Series $(\bar{x}) = (\, \Sigma\, n_i\, \bar{x}_i \,) / (\, \Sigma\, n_i \,)$

Here:

σ^2 = variance of the combined series

σ_i^2 = variance of the i-th series

n_i = number of observations in the i^{th} series

$\bar{x}_i$ = mean of the i^{th} series

$\bar{x}$ = mean of the combined series

d_i = deviation of individual mean from combined mean

This formula is used when the means and variances of two or more series are known and the combined variance is required.

Example: Suppose a series of 100 observations has mean 50 and variance 20. Another series of 200 observations has mean 80 and variance 40. Find the combined variance of the given series.

Solution:

Series	**(n)**	**Mean**	**Variance**
1	100	50	20
2	200	80	40

Given:
First Series: $n_1 = 100$; $\bar{x}_1 = 50$; $\sigma_1^2 = 20$
Second Series: $n_2 = 200$; $\bar{x}_2 = 80$; $\sigma_2^2 = 40$

Step 1: Calculate Combined Mean $\bar{x} = (n_1\bar{x}_1 + n_2\bar{x}_2) / (n_1 + n_2)$

$= (100\times50 + 200\times80) / 300$

$= (5000 + 16000) / 300$

$= 21000 / 300 = 70$

Step 2: Calculate Deviations from Combined Mean

$d_1 = \bar{x}_1 - \bar{x} = 50 - 70 = -20$

$d_1^2 = 400$

$d_2 = \bar{x}_2 - \bar{x} = 80 - 70 = 10$

$d_2^2 = 100$

Step 3: Apply Combined Variance Formula

$\sigma^2 = [\, n_1 (\sigma_1^2 + d_1^2) + n_2 (\sigma_2^2 + d_2^2) \,] / (n_1 + n_2)$

$\sigma^2 = [\, 100(20 + 400) + 200(40 + 100) \,] / 300$

$\sigma^2 = (42000 + 28000) / 300$

$\sigma^2 = 70000 / 300$

$\sigma^2 = 233.33$

Thus, Combined Variance = 233.33

The Variance and Standard deviation also possess some merits and demerits; they are as follows:

Variance:

- **Merits:** Variance is a precisely defined measure of dispersion that considers every observation in the dataset, ensuring comprehensive coverage of the data. Its mathematical formulation makes it suitable for algebraic manipulation and further statistical analysis. By squaring deviations, it effectively eliminates the issue of negative values and appropriately emphasizes larger deviations. Owing to these properties, variance has become one of the most widely used and accepted measures of dispersion in statistical studies.

- **Demerits:** Variance may not be appropriate in situations where the mean is not a suitable measure of central tendency, such as in distributions with open-ended classes; in such cases, quartile deviation is often preferred. It is not a unit-free measure, and its computation can be tedious for large datasets, often requiring the use of calculators or computers. Additionally, since variance is expressed in squared units of the original variable, interpreting the magnitude of dispersion becomes less intuitive compared to standard deviation.

Standard Deviation:

- **Merits:** Standard deviation is a well-defined and reliable measure of dispersion that takes into account all observations in a dataset. It is mathematically tractable, allowing for extensive algebraic and statistical applications. The squaring of deviations helps remove negative signs while preserving the influence of extreme values. Because it is expressed in the same unit as the original data, standard deviation is widely regarded as the most practical and popular measure of dispersion.

- **Demerits:** Standard deviation may be unsuitable when the mean does not adequately represent the data, particularly in the presence of open-ended classes, where quartile deviation can be more appropriate. It is not a unit-free measure, which limits comparisons across datasets with different units. Although conceptually straightforward, its calculation for large or complex datasets can be time-consuming and typically requires computational tools.

Note: Standard deviation is more suitable for algebraic and mathematical operations than mean deviation. Mean deviation directly ignores the signs of deviations, which results in the loss of information carried by positive and negative differences. In contrast, standard deviation retains the effect of signs initially, though they are removed automatically through squaring. However, mean deviation is still preferable in situations where the median is a better measure of central tendency than the mean.

There are two more aspects of dispersion one is Root Mean Square Deviation and Other is Coefficient of Variation, both are discussed below:

Root Mean Square Deviation (RMSD): As discussed earlier, standard deviation is the positive square root of the average of the squares of deviations taken from the mean. But, if the deviations are taken from an assumed mean instead of the actual mean, the resulting measure is called *Root Mean Square Deviation.*

Formula (Ungrouped Data): $RMSD = \sqrt{[\,\Sigma(x_i - A)^2 / n\,]}$

where: A = assumed mean and n = number of observations

For a frequency distribution, the formula is: $RMSD = \sqrt{[\,\Sigma f_i (x_i - A)^2 / \Sigma f_i\,]}$

Note: When the assumed mean is equal to the actual mean ($A = \bar{x}$), the Root Mean Square

Deviation becomes equal to the Standard Deviation.

The Coefficient of Variation (CV) is defined as a relative measure of dispersion used to compare the variability of two or more data series. The data series having a smaller coefficient of variation is considered more consistent.

Formula: $CV = (\sigma / \bar{x}) \times 100$

Note: *One should be careful while interpreting the coefficient of variation.*

For example, the series 10, 10, 10 has zero standard deviation and hence CV is also zero. Similarly, the series 50, 50, 50 also has zero standard deviation and CV equal to zero. Although both series have the same CV, the second series has a higher mean and is therefore considered more consistent.

Example: Suppose batsman A has mean 50 with Standard deviation (SD) 10. Batsman B has mean 30 with Standard deviation (SD) 3. What do you infer about their performance?

Solution: To compare the performance and consistency of the two batsmen, we use the Coefficient of Variation (CV). We Know, $CV = \frac{Standard\ Deviation}{Mean}$

Calculations for Coefficient of Variation for Batsman A & B:

- Coefficient of Variation for Batsman A (CV_A). $= \frac{10}{50} = 0.20$
- Coefficient of Variation for Batsman B (CV_B). $= \frac{3}{30} = 0.10$

Interpretation:

- Batsman A has a higher mean score (50), so he is a better run scorer.
- Batsman B has a lower coefficient of variation (0.10) compared to A (0.20), so he is more consistent.

Final Inference:

- Batsman A is the better performer in terms of average runs.
- Batsman B is the more consistent batsman.

1.3.7 SUMMARIZATION OF DATA

Descriptive summarization of data refers to the process of organizing, simplifying, and presenting raw data in a meaningful manner so that its essential characteristics can be clearly understood. Instead of focusing on individual observations, descriptive summarization highlights overall patterns, typical values, variability, and the general nature of the data distribution. It forms the foundation of statistical analysis and is widely used in data science to gain an initial understanding of datasets.

The main parameters used for descriptive summarization are broadly grouped into the following categories.

1. Measures of Central Tendency
2. Measures of Dispersion (Variability)
3. Measures of Position (Location)
4. Measures of Shape of Distribution
5. Measures of Relative Variability
6. Frequency-Based Parameters
7. Graphical & Visual Summaries

In the above sections we have already discussed in detail the parameters under these categories. Although a brief discussion on the parameters in each category is given below:

1. Measures of Central Tendency

Measures of central tendency describe the central or most representative value of a dataset. The mean provides the arithmetic average of all observations and is useful for numerical data that does not contain extreme values, although it is sensitive to outliers. The median represents the middle value of an ordered dataset and is more reliable when the data is skewed or contains extreme values. The mode indicates the most frequently occurring value and is especially useful for categorical data or for identifying the most common observation. Together, these measures offer a single value that summarizes the centre of the data.

The parameters listed below, Measures of Central Tendency and describe the central or typical value around which the data is clustered.

(a) Mean

- The arithmetic average of all observations.
- Useful for numerical data without extreme values.
- Sensitive to outliers.

(b) Median

- The middle value when data is arranged in ascending or descending order.
- More reliable for skewed distributions or data with extreme values.

(c) Mode

- The most frequently occurring value.
- Useful for categorical data and identifying the most common observation.

The purpose of Measures of central tendency is to provide a single representative value for the dataset.

2. Measures of Dispersion (Variability)

Measures of dispersion explain how far the data values are spread around the central value. The range gives a basic idea of spread by calculating the difference between the maximum and minimum values. Quartile deviation focuses on the dispersion of the middle 50% of the data and is less affected by extreme values. Mean deviation calculates the average of absolute deviations from a central value and is simple to understand, though less useful for advanced analysis. Variance measures the average of squared deviations from the mean and plays a significant role in theoretical and statistical modeling. Standard deviation, being the square root of variance, is the most widely used measure of dispersion because it is expressed in the same unit as the data. These measures help evaluate consistency, reliability, and variability within the dataset.

The parameters listed below, describe Measures of Dispersion (Variability) i.e. how spread out the data values are around the central value.

(a) Range

- Difference between maximum and minimum values.
- Provides a quick idea of data spread.

(b) Quartile Deviation (Interquartile Range)

- Measures dispersion of the middle 50% of data.
- Less affected by extreme values.

(c) Mean Deviation

- Average of absolute deviations from mean or median.
- Simpler to understand but less useful for advanced analysis.

(d) Variance

- Average of squared deviations from the mean.
- Used heavily in theoretical and statistical modeling.

(e) Standard Deviation

- Square root of variance.
- Most widely used measure of dispersion in data science.

The Purpose of Measures of dispersion is to help assess consistency, reliability, and variability of data.

3. Measures of Position (Location)

Measures of position indicate the relative standing of individual data values within a dataset. Quartiles divide the data into four equal parts, deciles divide it into ten equal parts, and percentiles divide it into one hundred equal parts. These measures help in understanding how a particular observation compares with the rest of the data and are useful in ranking and comparative analysis.

The parameters listed below, describe Measures of Position (Location) and indicate the relative position of a data value within a dataset

(a) Quartiles

- Divide data into four equal parts (Q1, Q2, Q3).

(b) Deciles

- Divide data into ten equal parts.

(c) Percentiles

- Divide data into one hundred equal parts.

The Purpose of Measures of position is to provide help to compare individual observations relative to the dataset.

4. Measures of Shape of Distribution

Measures of shape describe the overall form and pattern of the data distribution. Skewness indicates the degree of asymmetry in the distribution, showing whether the data is skewed to the right or left. Kurtosis measures the peakedness or flatness of the distribution and provides insight into the presence of extreme values. Together, these measures offer a deeper understanding of the nature of the data beyond central values and variability.

The parameters listed below, describe Measures of Shape of Distribution and describe the pattern and form of the data distribution followed by the data in any dataset

(a) Skewness

- Indicates the **symmetry or asymmetry** of the distribution.
- Positive skew: longer right tail
- Negative skew: longer left tail

(b) Kurtosis

- Measures the **peakedness or flatness** of the distribution.
- Helps identify the presence of extreme values.

The Purpose of Measures of shape is to provide insight into the nature of data distribution beyond averages.

5. Measures of Relative Variability

Measures of relative variability allow comparison between datasets that differ in units or scale. The coefficient of variation, which relates the standard deviation to the mean, is a unit-free measure that helps determine which dataset is more consistent or stable. This is particularly useful when comparing variability across different contexts.

The parameters listed below, describe Measures of Relative Variability and allow comparison between datasets with different units or scales.

Coefficient of Variation (CV)

- Ratio of standard deviation to mean.
- Unit-free measure of consistency.

The Purpose of Relative measures is to provide help to identify which dataset is more consistent or stable.

6. Frequency-Based Parameters

Frequency-based parameters summarize how often data values occur. These include frequency distributions, relative frequencies, and cumulative frequencies. Such parameters help reveal patterns of data concentration, distribution trends, and the overall structure of the dataset.

Frequency-Based Parameters summarize how often values occur, and they include following:

- Frequency tables
- Relative frequency
- Cumulative frequency

The Purpose of Frequency-Based Parameters is to help in understanding the distribution patterns and data concentration.

7. Graphical and Visual Summaries

Graphical and visual summaries complement numerical measures in descriptive summarization. Tools such as bar charts, histograms, pie charts, box plots, and line graphs visually represent data patterns, trends, and outliers. These visual techniques make complex datasets easier to interpret and communicate.

Although Graphical and Visual Summaries are numerical in nature, descriptive summarization is often supported by visual tools like:

- Histograms
- Bar charts
- Pie charts
- Box plots
- Line graphs

The Purpose of Graphical and Visual summaries is to make patterns, trends, and outliers easier to interpret.

Thus, we learned that the Descriptive summarization of data relies on a combination of measures of central tendency, dispersion, position, shape, relative variability, and frequency, supported by graphical representations. Together, these parameters provide a comprehensive statistical overview of a dataset, enabling better understanding, comparison, and informed decision-making in statistics and data science.

1.4 SUMMARY

This Unit-1 titled "Descriptive Data Analysis" is the foundational step in understanding any dataset before applying advanced statistical techniques. It begins with data collection, where information is gathered either as primary data (collected directly) or secondary data (already available). Once collected, it is essential to identify the type of data—qualitative (nominal, ordinal) or quantitative (discrete, continuous, interval, ratio)—as this determines the appropriate methods of analysis. The data is then organized using frequency distributions, which help in identifying patterns, concentrations, and irregularities within the dataset.

To make data more understandable, it is presented through graphical representations such as bar charts, histograms, pie charts, and line graphs, enabling quick visual interpretation of trends and patterns. Further, the dataset is summarized using measures of central tendency like mean, median, and mode, which provide insights into the typical or average values. Alongside this, measures of dispersion such as range, variance, and standard deviation help in understanding the spread or variability of the data, indicating how much the values deviate from the average.

Finally, data summarization combines these techniques to present information in a concise and meaningful way through tables, graphs, and key statistical values. This process allows analysts to draw initial insights, detect patterns, and prepare the data for further analysis. Overall, mastering these concepts equips learners with essential skills to interpret data accurately and make informed decisions in academic, research, and real-world contexts.

1.5 SOLUTIONS/ANSWERS

☞ **Check Your Progress 1**

Q1. Explain the concept of Descriptive Data Analysis. Why is it important in data analysis?
Answer: Refer to section 1.3

Q2 Classify the types of data with suitable examples.
Answer: Refer to section 1.3

Q3. What are Measures of Central Tendency and Dispersion? Explain briefly.
Answer: Refer to section 1.3

☞ **Check Your Progress 2**

Q1. What is Data Collection? Why is it considered a crucial step in data analysis?
Answer: Refer to section 1.3.1

Q2. Differentiate between Primary Data and Secondary Data.
Answer: Refer to section 1.3.1

Q3. Explain any three methods of collecting Primary Data with examples.
Answer: Refer to section 1.3.1

☞ **Check Your Progress 3**

Q1. Explain the four levels of measurement in data analysis i.e. Nominal, Ordinal, Ratio & interval data. Why are they important?
Answer: Refer to section 1.3.2

1.6 FURTHER READINGS

1. *Handbook of Statistical Analysis and Data Mining Applications* by Robert Nisbet, John Elder, and Gary Miner, published by Elsevier (2nd Edition, 2017, ISBN: 978-0124166455).
2. *Data Mining and Analysis: Fundamental Concepts and Algorithms* by Mohammed J. Zaki and Wagner Meira Jr., published by Cambridge University Press (1st Edition, 2014, ISBN: 978-0521766333).
3. *Core Concepts in Data Analysis: Summarization, Correlation and Visualization* by *Boris Mirkin, published by Springer (1st Edition, 2011, ISBN: 978-0857292872).*
4. *Applied Predictive Modeling, Max Kuhn and Kjell Johnson, 1st Edition, Springer, 2013.*
5. *An Introduction to Statistical Learning with Applications in R, Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani, 1st Edition, Springer, 2013.*
6. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Trevor Hastie, Robert Tibshirani and Jerome Friedman, 2nd Edition, Springer, 2009.*
7. *An Introduction to Statistical Learning with Applications in R by Gareth James, Daniela Witten, Trevor Hastie, and Robert Tibshirani, published by Springer (2nd Edition, 2021, ISBN: 978-1071614174).*
8. *Time Series Analysis: Forecasting and Control by George E. P. Box, Gwilym M. Jenkins, Gregory C. Reinsel, and Greta M. Ljung, published by Wiley (5th Edition, 2015, ISBN: 978-1118675021)*

UNIT 2 DIAGNOSTIC DATA ANALYSIS

Structure

2.1 INTRODUCTION

In our earlier unit-1 we learned about the Descriptive Data Analysis, here in this unit we are extending our discussion for Diagnostic Data Analysis, by the end of this unit we will understand that Descriptive data analysis and diagnostic data analysis are two closely related but conceptually different stages of data analysis, each serving a distinct purpose in understanding data.

Descriptive data analysis focuses on *what has happened* in the data. Its primary objective is to summarize, organize, and present raw data in a meaningful and understandable way. In this type of analysis, data is condensed using measures such as mean, median, mode, range, variance, standard deviation, percentages, and frequency distributions. Visual tools like tables, bar charts, pie charts, histograms, and line graphs are commonly used to highlight patterns, trends, and overall behavior of the dataset. For example, in a business context, descriptive analysis may show total sales for each month, average customer spending, or the distribution of product defects. However, descriptive analysis does not explain *why* these patterns occur; it only provides a clear snapshot of the current or past state of the data.

In contrast, **diagnostic data analysis** goes a step further by addressing the question of *why something happened*. It builds upon descriptive results and attempts to identify the underlying causes, relationships, and drivers behind observed patterns or anomalies. Diagnostic analysis uses more advanced techniques such as data mining, correlation analysis, drill-down analysis, segmentation, and comparative analysis. For instance, if descriptive analysis shows a sudden drop in sales in a particular month, diagnostic analysis investigates possible reasons—such as changes in customer behavior, pricing, competition, supply issues, or marketing performance. By examining relationships between variables (for example, advertising spend and sales, or temperature and product demand), diagnostic analysis provides insights into cause-and-effect or strong associations.

In practice, diagnostic analysis cannot exist without descriptive analysis, because one must first understand *what* is happening before exploring *why* it is happening. Together, they form a logical progression in data analysis: descriptive analysis provides clarity and structure to data, while diagnostic analysis transforms that clarity into actionable insights by uncovering causes and relationships.

From a learning perspective, the key difference lies in intent and depth. Descriptive analysis answers *"What is happening or what has happened?"* and is mainly concerned with summarization and presentation. Diagnostic analysis answers *"Why did it happen?"* and emphasizes reasoning, explanation, and interpretation. While descriptive analysis relies mostly on basic statistical measures and visualization, diagnostic analysis heavily depends on data mining algorithms (such as classification, clustering, and association rule mining) and correlation techniques (such as Pearson, Spearman, or Kendall correlation) to uncover hidden patterns and relationships.

It can be understood that, at the core of diagnostic analysis lie **data mining techniques** and **correlation analysis**, which together enable analysts to explore large datasets, detect hidden patterns, and identify meaningful relationships among variables.

Let's extend our discussion for Diagnostic Data Analysis at large, in the subsequent section.

2.2 EXPECTED LEARNING OUTCOMES

After completing this unit, you should be able to:

- Understand Diagnostic Analysis
- Understand Data Mining
- Understanding Correlation Analysis

2.3 DIAGNOSTICS ANALYSIS

Diagnostic analysis is a crucial stage in data analytics that goes beyond simply describing what has happened and focuses on understanding **why it happened**. While descriptive analysis summarizes historical data using measures such as mean, median, charts, and tables, diagnostic analysis investigates the underlying causes, patterns, and relationships hidden within the data. It plays a key role in transforming raw data into meaningful insights that support informed decision-making.

In the context of **Data Mining**, diagnostic analysis involves systematically exploring large datasets to discover significant patterns, trends, anomalies, and associations. Data mining techniques such as classification, clustering, association rule mining, and anomaly detection help analysts drill down into data to identify factors that influence outcomes. These techniques enable organizations to diagnose performance issues, customer behavior changes, operational inefficiencies, and risk factors by uncovering relationships that are not immediately visible through simple summaries.

Correlation Analysis is a fundamental statistical tool used in diagnostic analysis to measure the strength and direction of relationships between variables. By analyzing correlation coefficients, analysts can determine whether variables move together positively, negatively, or not at all. Correlation analysis helps answer diagnostic questions such as whether an increase in one variable is associated with a change in another, and to what extent. Although correlation does not imply causation, it provides strong evidence for identifying potential drivers that require further investigation.

Together, Data Mining and Correlation Analysis form the backbone of diagnostic analytics. Data mining techniques help extract and structure relevant information from complex datasets, while correlation analysis quantifies relationships among variables to support logical reasoning. In diagnostic analysis, these tools are widely applied across domains such as business intelligence, healthcare, finance, marketing, and operations to identify root causes of problems, explain performance variations, and support strategic decisions. The interrelation between Descriptive analysis, Diagnostic Analytics, Data Mining and Correlation Analysis will be more clear from the examples given below:

Example-1: To understand how data mining and correlation analysis relate to diagnostic analysis, consider a practical example from a business context. Suppose an e-commerce company observes a sudden drop in monthly sales. A descriptive analysis may report figures such as total sales, average order value, or number of customers, confirming that sales have indeed declined. However, it does not explain the reason behind this decline. At this stage, diagnostic analysis is applied to investigate the underlying causes.

Using data mining techniques, the company can explore large volumes of transactional and customer data to discover hidden patterns. For example, clustering algorithms may group customers based on purchasing behavior and reveal that a specific customer segment has sharply reduced its purchases. Classification algorithms may identify that customers exposed to delayed deliveries or negative product reviews are more likely to stop buying. Association rule mining might show that customers who experienced payment failures are also likely to abandon future purchases. These insights help diagnose *why* sales dropped rather than merely stating *that* they dropped.

At the same time, correlation analysis plays a crucial role in diagnostic analysis by quantifying the strength and direction of relationships between variables. For instance, the company may calculate the correlation between delivery time and repeat purchases and find a strong negative correlation, indicating that longer delivery times are associated with fewer repeat orders. Similarly, a high negative correlation between product return rates and customer satisfaction scores may suggest that frequent returns contribute to declining loyalty. While correlation does not imply causation, it provides strong evidence pointing toward potential causes that warrant deeper investigation.

Example-2 : In another example from healthcare, a hospital may observe an increase in patient readmission rates. Diagnostic analysis supported by data mining can identify patterns such as specific treatments, age groups, or hospital departments associated with higher readmissions. Correlation analysis may reveal that readmission rates are strongly correlated with shorter initial hospital stays or lack of follow-up care. Together, these techniques help healthcare administrators diagnose operational or clinical factors contributing to the problem.

Thus, data mining contributes to diagnostic analysis by automatically discovering patterns, trends, anomalies, and groupings in large and complex datasets, while correlation analysis supports diagnostic reasoning by measuring relationships among variables and highlighting influential factors. In diagnostic analysis, these techniques are often used together: data mining identifies what patterns exist, and correlation analysis helps explain how variables are related.

From the above examples it can be concluded that the diagnostic data analysis relies heavily on data mining and correlation analysis to move from observation to explanation. By uncovering hidden patterns and examining relationships among variables, these techniques enable analysts to answer critical "why" questions, making diagnostic analysis an essential step for informed decision-making and problem solving.

Overall, diagnostic analysis bridges the gap between **"what happened"** and **"why it happened."** By leveraging data mining methods and correlation analysis, analysts gain deeper insights into data behavior, enabling organizations to move from observation to explanation and lay the foundation for *predictive and prescriptive analytics to be discussed in subsequent units i.e. Unit 3 and Unit 4.*

So, we learned that, at the core of diagnostic analysis lie **data mining techniques** and **correlation analysis**, which together enable analysts to explore large datasets, detect hidden patterns, and identify meaningful relationships among variables.

In our subsequent section we are going to discuss the role of Data Mining in performing Diagnostic data analysis.

☞ Check Your Progress 1

Q1. Differentiate between Descriptive Data Analysis and Diagnostic Data Analysis.

…………………………………………………………………………………………

…………………………………………………………………………………………

Q2. What is Diagnostic Data Analysis? Explain its importance in decision-making.

…………………………………………………………………………………………

…………………………………………………………………………………………

Q3. Explain the role of Data Mining and Correlation Analysis in Diagnostic Analysis with examples.

…………………………………………………………………………………………

…………………………………………………………………………………………

2.3.1 DATA MINING

Data mining provides the computational and algorithmic foundation for diagnostic analysis. It involves extracting useful patterns, relationships, and knowledge from large and complex datasets that are otherwise difficult to analyse manually. We learned from the examples cited above in section 2.3 that, while performing diagnostic analysis, data mining helps to:

- **Identify root causes** behind outcomes such as declining sales, system failures, or customer churn
- **Discover hidden patterns** that are not immediately visible through simple summaries
- **Segment the data** into meaningful groups to compare behavior across categories
- **Detect anomalies and exceptions** that may explain unusual events

Several data mining algorithms are particularly useful in performing diagnostic analysis a few are listed below:

1. **Classification algorithms** (Decision Trees, Naïve Bayes, k-Nearest Neighbours):
   Used to determine factors responsible for categorical outcomes, such as loan default or disease diagnosis.
2. **Clustering algorithms** (K-Means, Hierarchical Clustering, DBSCAN):
   Help identify groups with similar characteristics to compare behaviors and diagnose differences.
3. **Association Rule Mining** (Apriori, FP-Growth):
   Discovers relationships between variables, such as product combinations influencing sales or symptoms linked to diseases.
4. **Anomaly / Outlier Detection**:
   Identifies unusual patterns that may signal errors, fraud, or rare events influencing results.
5. **Regression Analysis**:
   Helps quantify the impact of independent variables on a dependent variable, supporting cause-effect investigation.

Note: *In this section we are going to discuss the above-mentioned algorithms in brief, as there will be detailed discussion on these algorithms in other courses of this programme.*

Now in the remaining part of this section we are going to discuss the various data mining algorithms, which are particularly useful in performing diagnostic analysis, and the same are mentioned above.

1) **Classification algorithms** are an important group of data mining techniques used in diagnostic data analysis to identify and explain the factors responsible for **categorical outcomes**. These algorithms assign data instances to predefined classes based on input

features and are widely used to diagnose outcomes such as *loan default vs. non-default*, *disease present vs. absent*, or *spam vs. non-spam*.

- **Decision Trees** classify data by repeatedly splitting it into smaller subsets based on the most influential attributes. Each internal node represents a decision condition, and each leaf node represents a final class. In diagnostic analysis, decision trees are especially useful because they clearly show which factors lead to a particular outcome. For example, in loan default analysis, a decision tree may reveal that low income combined with poor credit history and high loan amount leads to a higher probability of default.

- **Naïve Bayes** is a probabilistic classification algorithm based on Bayes' theorem and the assumption of independence among predictors. Despite this simplifying assumption, it performs well in many real-world problems. In diagnostic contexts, Naïve Bayes helps estimate the probability of a categorical outcome given certain conditions. For instance, in disease diagnosis, it can compute the likelihood of a disease based on symptoms such as fever, blood pressure, and test results.

- **k-Nearest Neighbours (k-NN)** is a distance-based classification algorithm that assigns a class to a data point based on the majority class among its *k* closest neighbors. It is particularly useful when class boundaries are complex and non-linear. In diagnostic analysis, k-NN can help identify outcomes by comparing new cases with historically similar cases, such as predicting customer churn by analyzing behavior patterns of similar past customers.

Overall, classification algorithms play a vital role in diagnostic analysis by **explaining why a specific categorical outcome occurs** and identifying the key factors influencing that outcome.

2) **Clustering algorithms** are unsupervised data mining techniques used in diagnostic data analysis to identify **groups of data points with similar characteristics** without predefined class labels. These algorithms help analysts compare behaviors across different groups and diagnose why certain patterns or outcomes differ within a dataset.

- **K-Means clustering** partitions data into a predefined number of clusters by minimizing the distance between data points and their cluster centroids. In diagnostic analysis, K-Means is useful for segmenting entities such as customers, products, or regions to identify groups with similar behavior. For example, customer clusters based on spending patterns can help diagnose why certain segments respond differently to marketing strategies.

- **Hierarchical clustering** builds a tree-like structure (dendrogram) that shows how data points are progressively grouped based on similarity. It allows analysts to examine clustering at different levels of granularity. In diagnostic analysis, hierarchical clustering is helpful for understanding nested relationships, such as grouping patients based on medical profiles to diagnose variations in treatment outcomes.

- **DBSCAN (Density-Based Spatial Clustering of Applications with Noise)** groups data points based on density rather than distance. It is particularly effective in identifying irregularly shaped clusters and detecting noise or outliers. In diagnostic analysis, DBSCAN helps identify unusual behaviors or rare patterns, such as detecting abnormal transaction clusters that may explain fraud or system anomalies.

Overall, clustering algorithms support diagnostic analysis by revealing **hidden group structures**, enabling meaningful comparisons, and helping explain differences in behavior across segments.

3) **Association Rule Mining** is a data mining technique used in diagnostic data analysis to uncover **hidden relationships and co-occurrence patterns** among variables in large datasets. It helps explain why certain events or outcomes occur together by identifying rules of the form *if–then*, such as *if a customer buys item A, then they are likely to buy item B*. This makes association rule mining especially useful in understanding interactions between factors that influence outcomes.

   - The **Apriori algorithm** identifies frequent item sets by repeatedly scanning the dataset and using the principle that all subsets of a frequent itemset must also be frequent. In diagnostic analysis, Apriori helps determine which combinations of variables are strongly associated. For example, in retail analytics, Apriori can reveal that customers who purchase smartphones and earphones are also likely to purchase protective cases, helping diagnose sales patterns and cross-selling opportunities.
   - The **FP-Growth (Frequent Pattern Growth)** algorithm improves efficiency by avoiding repeated database scans. It uses a compact tree structure called an FP-tree to extract frequent item sets more quickly. In diagnostic contexts, FP-Growth is particularly effective for large datasets, such as medical records, where it can identify combinations of symptoms or treatments that commonly occur together, aiding in disease diagnosis and treatment evaluation.

   Overall, association rule mining plays a key role in diagnostic analysis by revealing **interdependencies between variables**, enabling analysts to understand how combinations of factors contribute to observed outcomes.

4) **Anomaly or Outlier Detection** is an important data mining technique used in diagnostic data analysis to identify **unusual, rare, or unexpected patterns** in data that deviate significantly from normal behavior. Such anomalies often indicate errors, fraud, system failures, or exceptional events that strongly influence overall results. By detecting these irregularities, analysts can diagnose underlying issues that may not be visible through average-based or summary statistics.

Several algorithms are commonly used for anomaly detection. **Statistical methods**, such as Z-score and Interquartile Range (IQR), identify outliers by measuring how far a data point lies from the central tendency of the dataset. These methods are simple and effective when data follows a known distribution. **Distance-based techniques**, including k-Nearest Neighbors (k-NN), detect anomalies by identifying points that are far away from their nearest neighbors, making them useful when abnormal behavior differs significantly from the majority of observations.

Density-based algorithms, such as DBSCAN and Local Outlier Factor (LOF), identify anomalies based on regions of low data density. These methods are particularly effective in complex datasets where normal behavior forms dense clusters and anomalies appear in sparse regions. **Isolation Forest**, a tree-based algorithm, isolates anomalies by randomly partitioning the data; since anomalies are easier to separate, they require fewer splits, making this method efficient for large datasets.

In diagnostic analysis, anomaly detection helps explain unusual outcomes, such as sudden spikes in transactions indicating fraud, abnormal sensor readings pointing to equipment failure, or rare medical cases affecting health outcomes. By identifying and analyzing these outliers, organizations can better understand exceptional events and take corrective or preventive actions.

5) **Regression Analysis** is a fundamental analytical technique used in diagnostic data analysis to **quantify the relationship between a dependent variable and one or more independent variables**. It helps investigate cause–effect relationships by estimating how changes in explanatory variables influence an outcome. By providing numerical coefficients and statistical significance measures, regression analysis enables analysts to identify key drivers behind observed trends and variations.

Several regression algorithms are commonly used in diagnostic analysis. **Simple Linear Regression** examines the effect of a single independent variable on a dependent variable, such as assessing how advertising spend influences sales. **Multiple Linear Regression** extends this approach to multiple predictors, allowing analysts to evaluate the combined impact of factors like price, income, and promotion on demand. **Logistic Regression** is used when the dependent variable is categorical, such as predicting loan default or disease occurrence, by estimating probabilities rather than continuous values.

Advanced regression techniques include **Ridge Regression** and **Lasso Regression**, which address multicollinearity and improve model stability by adding regularization terms. **Polynomial Regression** captures non-linear relationships between variables, while **Elastic Net** combines features of Ridge and Lasso for better variable selection. In diagnostic analysis, these models help isolate influential factors, measure their relative importance, and support evidence-based explanations of outcomes.

Overall, regression analysis plays a critical role in diagnostic analytics by translating relationships into quantifiable effects, thereby strengthening cause–effect investigation and decision-making.

Now it is the time to understand the difference between Correlation analysis and regression analysis, as they are the two closely related statistical techniques, but they serve different purposes and are used in different analytical contexts.

Regression analysis, on the other hand, focuses on examining the **cause–effect relationship** between variables. It clearly distinguishes between a dependent variable (outcome) and one or more independent variables (predictors). Regression analysis not only identifies relationships but also quantifies the impact of independent variables on the dependent variable by estimating regression coefficients. These coefficients indicate how much the dependent variable is expected to change for a unit change in an independent variable. Unlike correlation, regression analysis supports prediction and forecasting through regression equations and models. The relationship in regression is directional, moving from independent variables to the dependent variable, and the variables are not interchangeable.

Correlation analysis is primarily concerned with measuring the **degree and direction of association** between two variables. It answers the question of whether variables are related and how strongly they move together. The relationship measured through correlation is symmetrical, meaning there is no distinction between dependent and independent variables—both variables are treated equally. The result of correlation analysis is a correlation coefficient, usually ranging between –1 and +1, where positive values indicate a direct relationship, negative values indicate an inverse relationship, and values close to zero indicate little or no linear association. Correlation analysis does not explain cause-and-effect relationships and cannot be used for prediction; instead, it serves as an exploratory tool to identify relationships that may warrant further investigation, especially in diagnostic data analysis.

In practical applications, correlation analysis is often used as a preliminary step to detect associations between variables, while regression analysis is applied afterward to model and explain these relationships in detail. Together, they complement each other in diagnostic data analysis by first identifying related variables and then quantifying and validating their influence on outcomes.

In the subsequent section we are going to extend our discussion for Diagnostic Analysis to the level of correlation analysis.

☞ Check Your Progress 2

Q1. What is the role of Data Mining in Diagnostic Data Analysis?

…….

…….

Q2. Explain any three data mining algorithms used in diagnostic analysis with examples

……….…

…………………………………………………………………………………………………….………

Q3. Differentiate between Correlation Analysis and Regression Analysis.

……….…

…………………………………………………………………………………………………….………

2.3.2 CORRELATION ANALYSIS

Correlation analysis is a core statistical technique in **diagnostic data analysis**, whose primary objective is to explain **why certain outcomes occur** by examining relationships among variables. While descriptive analysis tells us *what has happened*, diagnostic analysis seeks to uncover *why it happened*, and correlation analysis provides one of the most powerful tools for this investigation.

You have already learned various statistical measures such as measures of central tendency, measures of dispersion, moments, skewness, and kurtosis, all of which analyze variables individually. However, in many practical situations, it is necessary to examine two variables simultaneously in order to understand the relationship between them.

This section introduces the concept of correlation, which focuses on studying the linear relationship between two or more variables. For the sake of basic understanding, one can say that "*when two variables are related in such a way that change in the value of one variable affects the value of another variable, then variables are said to be correlated or there is correlation between these two variables.*"

You will learn how to compute the correlation coefficient under different conditions and understand its properties. Therefore, before beginning this section, it is recommended that you revise the concepts of arithmetic mean and variance, as they will help you grasp the fundamentals of correlation more effectively.

Further, in diagnostic data analysis, the basic understanding of correlation is essential, as it enables analysts to move beyond simple averages and summaries and investigate the interdependencies that exist within the data.

At the most fundamental level, correlation analysis is used to measure the degree and direction of association between two variables. It helps determine whether two variables move together, whether a change in one variable is accompanied by a change in the other, and whether the relationship between them is positive, negative, or nonexistent. The strength of this relationship is represented numerically by a correlation coefficient, which generally lies between –1 and +1. A value of +1 indicates a perfect positive relationship (i.e. *Positive Correlation*), –1 indicates a perfect negative relationship (i.e. *Negative Correlation*), and 0 signifies the absence of a linear relationship.

It can be understood that according to the direction of change in variables there are two **types of correlation:**

1. **Positive Correlation:** The Correlation between two variables is said to be positive if the values of the variables deviate in the same direction i.e. if the values of one variable increase (or decrease) then the values of other variables also increase (or decrease).

2. **Negative Correlation:** The Correlation between two variables is said to be negative if the values of variables deviate in opposite directions i.e. if the values of one variable increase (or decrease) then the values of other variables decrease (or increase).

So far as the graphical interpretation is concerned, the scatter diagram is quite useful to understand the type of relationship between the variables. It can be understood that the Scatter diagram is a statistical tool for determining the potentiality of correlation between dependent variable and independent variable. The scatter diagram does not talk about the exact relationship between two variables but it indicates whether they are correlated or not.

Let (X_i, Y_i) where i = 1, 2, …, n be the bivariate distribution. If the values of the dependent variable Y_i are plotted against corresponding values of the independent variable X_i in the XY plane, such a diagram of dots is called a scatter diagram or dot diagram. *It is to be noted that scatter diagrams are not suitable for large numbers of observations.*

In order to interpret from the scatter diagram, one needs to analyse the pattern followed by the dots in the scatter diagram, and interpretations can be made on the basis of the following:

1. If dots are in the shape of a line and the line rises from left bottom to the right top, then correlation is said to be perfectly positive, as shown below.

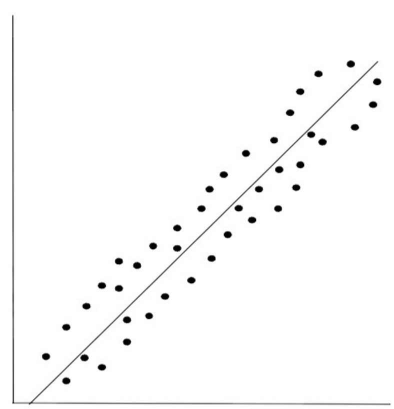

2. If dots in the scatter diagram are in the shape of a line and line moves from left top to right bottom, then correlation is perfectly negative, as shown below.

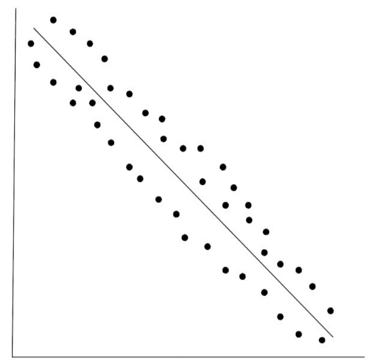

3. If dots of scatter diagram do not show any trend there is no correlation between the variables, as shown below.

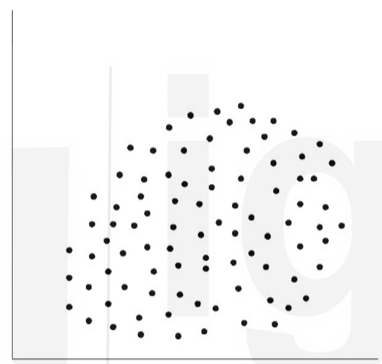

Coefficient of Correlation: Scatter diagram tells us whether variables are correlated or not. But it does not indicate the extent of which they are correlated. Coefficient of correlation gives the exact idea of the extent of which they are correlated. Coefficient of correlation measures the intensity or degree of linear relationship between two variables.

There are some Common Correlation Measures Used in Diagnostic Analysis, as diagnostic analysis progresses, different correlation measures are chosen based on data type and complexity, a few are listed below:

- **Pearson Correlation:** Used to measure linear relationships between continuous variables. It is widely applied in business, finance, and engineering diagnostics.
- **Spearman Rank Correlation**: Used when relationships are monotonic but not necessarily linear, or when data is ordinal. It is valuable in diagnostic analysis when assumptions of normality are violated.
- **Kendall's Tau:** Focuses on rank concordance and is more robust for small datasets or data with ties.

At an intermediate level, analysts learn to select the appropriate correlation measure based on data distribution, scale, and analytical Objectives.

Now we will discuss these Common Correlation Measures Used in Diagnostic Analysis, let's begin with Pearson Correlation.

Pearson Correlation:

Pearson correlation is one of the most widely used statistical techniques in data science for measuring the strength and direction of a linear relationship between two continuous variables. In diagnostic data analysis, it plays an important role in identifying variables that are closely associated with an outcome and may act as potential drivers behind observed patterns. Pearson correlation focuses on how two variables vary together and helps determine whether an increase or decrease in one variable is accompanied by a corresponding change in another variable, as well as the nature and intensity of this relationship.

The result of Pearson correlation is expressed through the Pearson correlation coefficient, usually denoted by r. The value of r lies between –1 and +1. A value of +1 indicates a perfect positive linear relationship, meaning that both variables move in the same direction in exact proportion. A value of –1 indicates a perfect negative linear relationship, where one variable increases while the other decreases in exact proportion. A value of 0 suggests that there is no linear relationship between the variables. In diagnostic analysis, values of r closer to ±1 highlight strong associations that may require deeper investigation.

The Pearson correlation coefficient is calculated using the formula

$$r = \frac{\sum(x-\bar{x})(y-\bar{y})}{\sqrt{\sum(x-\bar{x})^2\sum(y-\bar{y})^2}}$$

where x and y represent individual observations of the two variables, and $\bar{x}$ and $\bar{y}$ represent their respective means. The numerator of this formula measures the extent to which the two variables vary together, while the denominator standardizes this covariance by accounting for the variability of each variable individually. This standardization ensures that the coefficient is unit-free and confined within the range of –1 to +1.

Alternate formula (in terms of Variance and covariance of the data): If **X** and **Y** are two random variables, then the **correlation coefficient** between **X** and **Y** is denoted by **r** and is defined as:

$$r = Corr(X,Y) = \frac{Cov(X,Y)}{\sqrt{V(X)\,V(Y)}}$$

Here, **Corr(X, Y)** indicates the correlation coefficient between the two variables **X** and **Y**. Where **Cov(X, Y)** is the covariance between **X** and **Y**, defined as:

$$Cov(X,Y) = \frac{1}{n}\sum_{i=1}^{n}\,(x_i-\bar{x})(y_i-\bar{y})$$

And **V(X)**, the variance of **X**, is defined as:

$$V(X) = \frac{1}{n}\sum_{i=1}^{n}\,(x_i-\bar{x})^2$$

And **V(Y)**, the variance of **Y**, is defined as:

$$V(Y) = \frac{1}{n}\sum_{i=1}^{n} (y_i - \bar{y})^2$$

To illustrate the concept, consider a dataset where advertising expenditure and sales revenue are observed together.

Example: Suppose a data scientist wants to diagnose whether **advertising spend** is related to **sales**.

Advertising Spend (X)	Sales (Y)
10	40
20	50
30	60
40	70
50	80

Step 1: Calculate the means

$$\bar{x} = \frac{10 + 20 + 30 + 40 + 50}{5} = 30 \qquad \bar{y} = \frac{40 + 50 + 60 + 70 + 80}{5} = 60$$

Step 2: Create deviation table

X	Y	$x - \bar{x}$	$y - \bar{y}$	$(x - \bar{x})(y - \bar{y})$	$(\bar{x})^2$	$(\bar{y})^2$
10	40	-20	-20	400	400	400
20	50	-10	-10	100	100	100
30	60	0	0	0	0	0
40	70	10	10	100	100	100
50	80	20	20	400	400	400

$$\sum(x - \bar{x})(y - \bar{y}) = 1000 \sum(x - \bar{x})^2 = 1000 \sum(y - \bar{y})^2 = 1000$$

Step 3: Apply Pearson correlation formula

$$r = \frac{1000}{\sqrt{1000 \times 1000}} = \frac{1000}{1000} = 1$$

As the computed value of *r* is 1, it indicates a strong positive linear relationship, suggesting that higher advertising spend is strongly associated with higher sales.

In diagnostic data analysis, such a result helps analysts identify advertising as a potential driver of sales performance, although it does not by itself confirm causation.

Example: Given Data (5 days), A data analyst wants to check whether Digital Ad Spend (X) is linearly related to Daily Sales (Y) for a small online store.

Day	Ad Spend X (₹ '000)	Sales Y (₹ '000)
1	2	20
2	4	22
3	6	25
4	8	29
5	10	35

Solution: Using the correlation formula: $r = \boldsymbol{Corr}(X, Y) = \frac{\boldsymbol{Cov}(X,Y)}{\sqrt{V(X)\,V(Y)}}$ following Tasks are to be performed to calculate the coefficient of correlation (r), using the above formula $\bar{x}$ $\bar{y}$

1. Compute $\bar{x}$ and $\bar{y}$
2. Compute $Cov(X, Y)$
3. Compute $V(X)$and $V(Y)$
4. Find r and interpret it for diagnostic analysis (does spend relate to sales?)

Step 1: Means

$$\bar{x} = \frac{2+4+6+8+10}{5} = \frac{30}{5} = 6 \quad ; \quad \bar{y} = \frac{20+22+25+29+35}{5} = \frac{131}{5} = 26.2$$

Step 2: Build calculation table

X	**Y**	$x-\bar{x}$	$y-\bar{y}$	$(x-\bar{x})(y-\bar{y})$	$(\bar{x})^2$	$(\bar{y})^2$
2	20	-4	-6.2	24.8	16	38.44
4	22	-2	-4.2	8.4	4	17.64
6	25	0	-1.2	0	0	1.44
8	29	2	2.8	5.6	4	7.84
10	35	4	8.8	35.2	16	77.44

Now sum:

$$\sum(x-\bar{x})(y-\bar{y}) = 24.8 + 8.4 + 0 + 5.6 + 35.2 = 74 \sum(x-\bar{x})^2 = 16 + 4 + 0 + 4 + 16 = 40 \sum(y-\bar{y})^2 = 38.44 + 17.64 + 1.44 + 7.84 + 77.44 = 142.8$$

Step 3: Covariance and variances (using 1/n as in your formula)

$$Cov(X, Y) = \frac{1}{n}\sum(x-\bar{x})(y-\bar{y}) = \frac{74}{5} = 14.8 V(X) = \frac{1}{n}\sum(x-\bar{x})^2 = \frac{40}{5} = 8 V(Y) = \frac{1}{n}\sum(y-\bar{y})^2 = \frac{142.8}{5} = 28.56$$

Step 4: Correlation

$$r = \frac{14.8}{\sqrt{8 \times 28.56}} = \frac{14.8}{\sqrt{228.48}} = \frac{14.8}{15.115} \approx 0.98$$

Since r is close to +1, there is a very strong positive linear relationship between Ad Spend and Sales. This supports a diagnostic insight that higher ad spend is strongly associated with higher sales

Pearson correlation is especially advantageous in diagnostic analysis because it helps analysts move beyond simple averages and summaries to examine interdependencies between variables. It is often used as a preliminary diagnostic tool to filter relevant variables before applying more advanced techniques such as regression analysis or causal modelling. By highlighting which variables are strongly associated, *Pearson correlation supports hypothesis generation and guides deeper analytical investigation.*

We can say that the *Pearson correlation is highly relevant in diagnostic analysis* because of the following reasons:

- *Identifies variables that move together*
- *Helps narrow down potential explanatory factors*
- *Supports hypothesis generation*
- *Acts as a foundation for regression and predictive modelling*

However, for an expert-level understanding, it is important to recognize the *limitations of Pearson correlation*. It captures only linear relationships, is sensitive to outliers, and does not establish cause-and-effect relationships. Therefore, in diagnostic data analysis, Pearson correlation should be interpreted carefully and used in combination with domain knowledge, regression analysis, and other data mining techniques.

Limitations of Pearson correlation (Diagnostic Perspective),For expert diagnostic use, it is important to note that:

- Pearson correlation captures only linear relationships
- It is sensitive to outliers
- It does not imply causation

Correlation vs. Causation in Diagnostic Analysis: At an expert level, understanding the limitations of correlation it is important to understand that the Correlation analysis:

- Identifies associations, not causality
- Cannot, by itself, prove that one variable cause another
- Must be interpreted in conjunction with domain knowledge, regression analysis, experiments, or causal models

It is to be noted that in diagnostic data analysis, correlation is used as **evidence**, not as proof. Strong correlations suggest where causes may exist, but further analytical or experimental validation is required.

Therefore, Pearson correlation should be used as a diagnostic indicator, not as final proof.

It is learned from the above discussion that the Pearson correlation is a foundational tool for diagnostic data analysis in data science. It provides a clear and quantitative measure of linear association between variables, helping analysts identify influential factors, understand data behavior, and lay the groundwork for more advanced analytical methods.

Spearman Rank Correlation:
Spearman rank correlation is a non-parametric statistical measure used to assess the strength and direction of association between two variables based on their ranks rather than their actual values. It is especially useful in diagnostic data analysis when the data does not satisfy the assumptions required for Pearson correlation, such as normality or linearity, or when the data is ordinal in nature.
Spearman rank correlation examines whether **two variables move together in a consistent (monotonic) manner**, meaning that as one variable increases, the other either consistently increases or consistently decreases. Unlike Pearson correlation, Spearman correlation does not require the relationship to be linear; it only needs to be monotonic.

In diagnostic data analysis, Spearman correlation is valuable when:

- Data contains **outliers**
- Data is **ordinal or ranked**
- The relationship is **non-linear but monotonic**
- Exact numerical values are less reliable than relative ordering

The Spearman rank correlation coefficient is denoted by **ρ (rho)**.

The value of Spearman's rank correlation coefficient lies between **–1 and +1**:

- **ρ = +1** → Perfect positive monotonic relationship
- **ρ = –1** → Perfect negative monotonic relationship
- **ρ = 0** → No monotonic relationship

In diagnostic analysis, higher absolute values of ρ indicate stronger associations that may point to important explanatory variables.

When there are **no tied ranks**, the Spearman rank correlation coefficient is calculated using the formula:

$$\rho = 1 - \frac{6\sum d^2}{n(n^2 - 1)}$$

Explanation of Symbols:

- d= difference between the ranks of corresponding observations
- n= number of observations
- $\sum d^2$= sum of squares of rank differences

This formula measures how closely the rankings of two variables agree.

Note: When tied ranks are present, an adjusted formula using Pearson correlation on ranks is applied, it is out of the scope and we are not going to discuss this case here.

Example: Given Data:

Employee	Experience (Years)	Performance Score
A	2	60
B	4	65
C	6	70
D	8	85
E	10	90

A data scientist wants to diagnose whether employee experience is related to job performance ratings. *Where the performance is ranked subjectively.*

Solution: Since performance is ranked subjectively, Spearman rank correlation is appropriate.

Step 1: Assign Ranks

Employee	Experience	Rank X	Performance	Rank Y
A	2	1	60	1
B	4	2	65	2
C	6	3	70	3
D	8	4	85	4
E	10	5	90	5

Step 2: Compute Rank Differences

Employee	Rank X	Rank Y	d = X − Y	d²
A	1	1	0	0
B	2	2	0	0
C	3	3	0	0
D	4	4	0	0
E	5	5	0	0

$$\sum d^2 = 0$$

Step 3: Apply the Formula

$$\rho = 1 - \frac{6 \times 0}{5(5^2 - 1)} = 1$$

A Spearman correlation coefficient of **ρ = 1** indicates a **perfect positive monotonic relationship** between experience and performance. In diagnostic data analysis, this suggests that higher experience consistently corresponds to higher performance ratings, making experience a strong diagnostic factor.

Example: Given Data

Employee	Training Hours (X)	Performance Score (Y)
A	12	80
B	20	82
C	16	85
D	25	90
E	10	70
F	18	78

A data scientist wants to examine whether there is a relationship between hours of online training completed (X) and performance ranking (Y) of employees.

Solution: Since performance is ranked and the relationship may not be perfectly linear, Spearman rank correlation is appropriate.

Employee	Training Hours (X)	Performance Score (Y)
A	12	80
B	20	82
C	16	85
D	25	90
E	10	70
F	18	78

Step 1: Assign Ranks : **Rank X**

X	Rank X
10	1
12	2
16	3
18	4
20	5
25	6

Rank Y

Y	Rank Y
70	1
78	2
80	3
82	4
85	5
90	6

Step 2: Compute d and d²

Employee	Rank X	Rank Y	d	d²
A	2	3	-1	1
B	5	4	1	1
C	3	5	-2	4
D	6	6	0	0
E	1	1	0	0
F	4	2	2	4

$$\sum d^2 = 1 + 1 + 4 + 0 + 0 + 4 = 10$$

Step 3: Apply Spearman Rank Correlation Formula

$$\rho = 1 - \frac{6\sum d^2}{n(n^2 - 1)} \rho = 1 - \frac{6 \times 10}{6(36 - 1)} = 1 - \frac{60}{210} = 1 - 0.286 = 0.714$$

A Spearman correlation of 0.714 indicates a strong positive monotonic relationship, but not a perfect one. This suggests that higher training hours generally lead to better performance, but other factors (such as experience, motivation, or role complexity) also influence performance. This insight is valuable for diagnostic analysis, as it highlights training as an important but not exclusive driver of performance.

Kendall's Tau:

Kendall's Tau correlation is a **non-parametric measure of association** used to determine the **strength and direction of relationship between two variables based on the relative ordering (ranks) of data rather than their actual values**. It is especially useful in **diagnostic data analysis** when the dataset is small, contains tied ranks, or when the relationship between variables is ordinal and monotonic rather than strictly linear.

Kendall's Tau focuses on the **consistency of ordering between two variables**. Instead of measuring how far values deviate from the mean (as in Pearson correlation), Kendall's Tau evaluates whether the **relative rankings of observations agree or disagree** across two variables.

In diagnostic data analysis, this is particularly important when:

- Data is **ordinal or ranked**
- There are **many tied values**
- The sample size is **small**
- Robustness and interpretability are preferred over sensitivity to outliers

Kendall's Tau is denoted by **τ (tau)**.

For any pair of observations $(y_i)$and (y_j):

- The pair is **concordant** if:
  - $x_i > x_j$and $y_i > y_j$, or
  - $x_i < x_j$and $y_i < y_j$
- The pair is **discordant** if:
  - $x_i > x_j$and $y_i < y_j$, or
  - $x_i < x_j$and $y_i > y_j$

Kendall’s Tau is based on comparing the **number of concordant and discordant pairs** in the dataset.

The basic formula for Kendall’s Tau (τ) is:

$$\tau = \frac{C - D}{\frac{1}{2}n(n-1)}$$

Where:

- C= number of concordant pairs
- D= number of discordant pairs
- n= number of observations
- $\frac{1}{2}n(n-1)$= total number of possible pairs

This formula measures how much agreement exists between the rankings of two variables.

The value of Kendall’s Tau lies between **–1 and +1**:

- **τ = +1** → Perfect agreement in rankings
- **τ = –1** → Perfect disagreement in rankings
- **τ = 0** → No association in rankings

In diagnostic data analysis, higher absolute values of τ indicate stronger ordinal association and suggest meaningful diagnostic relationships.

Example: Given Data

Project	Complexity Rank (X)	Completion Time Rank (Y)
A	1	2
B	2	1
C	3	4
D	4	3

A data scientist wants to diagnose whether **project complexity ranking** is related to **time-to-completion ranking** for completed projects, Using Kendall’s Tau correlation.

Solution:

Step 1: List all possible pairs

For $n = 4$, total pairs:

$$\frac{4(4-1)}{2} = 6$$

Pairs:

- (A, B)
- (A, C)
- (A, D)
- (B, C)
- (B, D)
- (C, D)

Step 2: Determine Concordant and Discordant Pairs

Pair	X Order	Y Order	Result
A, B	$1 < 2$	$2 > 1$	Discordant
A, C	$1 < 3$	$2 < 4$	Concordant
A, D	$1 < 4$	$2 < 3$	Concordant
B, C	$2 < 3$	$1 < 4$	Concordant
B, D	$2 < 4$	$1 < 3$	Concordant
C, D	$3 < 4$	$4 > 3$	Discordant

- Concordant pairs (C) = 4
- Discordant pairs (D) = 2

Step 3: Apply Kendall's Tau Formula

$$\tau = \frac{C - D}{\frac{1}{2}n(n-1)} = \frac{4-2}{6} = \frac{2}{6} = 0.33$$

A Kendall's Tau value of **0.33** indicates a **moderate positive ordinal association** between project complexity and completion time. In diagnostic data analysis, this suggests that **higher complexity projects tend to take longer**, but the relationship is not perfectly consistent, indicating the influence of other factors such as team size, experience, or tools used.

Important Note: In this example, the data consists of **ranked variables**—project complexity rank and completion time rank. Since the information is ordinal rather than continuous, Karl Pearson's coefficient of correlation is not appropriate. Pearson correlation requires numerical data measured on an interval or ratio scale and assumes a linear relationship and approximate normality. In this case, the actual numerical distances between ranks are not meaningful; only the **order of observations** matters. Therefore, Pearson correlation would give misleading results.

While **Spearman rank correlation** is suitable for ranked data, it measures association based on **differences in ranks** and treats all rank differences equally. Spearman's method is more sensitive to larger rank gaps and assumes that the monotonic relationship between variables is relatively smooth. In small datasets, such as the one used in the example, Spearman correlation can exaggerate the strength of association because a single change in rank can significantly affect the coefficient. This makes it less reliable for fine-grained diagnostic interpretation when the number of observations is limited.

Kendall's Tau, on the other hand, is especially well suited for this example because it is based on the **comparison of concordant and discordant pairs**. Instead of focusing on the magnitude of rank differences, Kendall's Tau evaluates how consistently the ordering of one variable agrees with the ordering of the other. This pairwise comparison approach provides a more **robust and intuitive measure of association**, particularly for small samples and diagnostic analysis.

From a diagnostic perspective, Kendall's Tau is more meaningful because it directly answers the question: *What proportion of pairs are ordered consistently across the two variables?* In the given example, this helps assess how reliably higher project complexity is associated with longer completion time, without being overly influenced by individual rank shifts. Moreover, Kendall's Tau handles tied ranks and ordinal inconsistencies more effectively, making it a preferred choice when data quality or scale is limited.

Overall, Kendall's Tau is more relevant than Spearman rank correlation and Karl Pearson's coefficient in this example because the data is ordinal, the sample size is small, and the diagnostic goal is to assess **consistency of ordering rather than strength of linear or monotonic relationships**.

It is important to note that Kendall's Tau is particularly valuable because it is robust to the presence of outliers, handles tied ranks more effectively than Spearman's rank correlation, performs well even with small datasets, and offers a probability-based interpretation of agreement between variables. When compared with Spearman correlation, Kendall's Tau is generally more conservative and easier to interpret, while Spearman is more sensitive to large differences in ranks. Additionally, Kendall's Tau tends to perform better in situations involving small samples and tied observations, as it directly measures the degree of pairwise agreement between the rankings of two variables.

Overall, we learned that Kendall's Tau correlation is a powerful and robust tool for **diagnostic data analysis**, particularly when dealing with ranked, ordinal, or non-linear data. By focusing on concordant and discordant pairs, it provides a clear and intuitive measure of association that supports root-cause investigation and informed analytical reasoning.

☞ Check Your Progress 3

Q1. What is Correlation Analysis? Explain the meaning and range of the correlation coefficient. Also, Differentiate between Positive and Negative Correlation with examples.

……
……

Q2. Compare Pearson, Spearman, and Kendall correlation methods.

……
……

2.5 SUMMARY

Diagnostic Data Analysis is an advanced stage of data analysis that focuses on answering the question **"why did something happen?"**. While descriptive analysis explains what has occurred, diagnostic analysis goes deeper to identify the underlying causes, patterns, and relationships within the data. A key component of this process is **Data Mining**, which provides the computational tools and algorithms required to extract meaningful patterns from large and complex datasets. Techniques such as classification, clustering, association rule mining, anomaly detection, and regression analysis help analysts identify root causes, segment data into meaningful groups, detect unusual patterns, and quantify the impact of various factors on outcomes. These methods enable organizations to move from observation to explanation, making data-driven decision-making more effective.

Another fundamental component of diagnostic analysis is **Correlation Analysis**, which examines the strength and direction of relationships between variables. It helps determine whether changes in one variable are associated with changes in another, using measures such as Pearson, Spearman, and Kendall correlation coefficients. The correlation coefficient ranges from –1 to +1, indicating negative, positive, or no relationship. While correlation does not imply causation, it plays a crucial role in identifying potential drivers and guiding further investigation. Together, data mining and correlation analysis form the backbone of diagnostic analytics, enabling analysts to uncover hidden relationships, explain observed patterns, and generate actionable insights for solving real-world problems and improving decision-making.

2.5 SOLUTIONS/ANSWERS

☞ **Check Your Progress 1**

Q1. Differentiate between Descriptive Data Analysis and Diagnostic Data Analysis.
Solution: Refer to section 2.1 & 2.3
Q2. What is Diagnostic Data Analysis? Explain its importance in decision-making.
Solution: Refer to section 2.1 & 2.3
Q3. Explain the role of Data Mining and Correlation Analysis in Diagnostic Analysis with examples.
Solution: Refer to section 2.1 & 2.3

☞ **Check Your Progress 2**

Q1. What is the role of Data Mining in Diagnostic Data Analysis?
Solution: Refer to section 2.3.1
Q2. Explain any three data mining algorithms used in diagnostic analysis with examples
Solution: Refer to section 2.3.1
Q3. Differentiate between Correlation Analysis and Regression Analysis.
Solution: Refer to section 2.3.1

☞ **Check Your Progress 3**

Q1. What is Correlation Analysis? Explain the meaning and range of the correlation coefficient. Also, Differentiate between Positive and Negative Correlation with examples.
Solution: Refer to section 2.3.2
Q2. Compare Pearson, Spearman, and Kendall correlation methods.
Solution: Refer to section 2.3.2

2.6 FURTHER READINGS

1. *Handbook of Statistical Analysis and Data Mining Applications by Robert Nisbet, John Elder, and Gary Miner, published by Elsevier (2nd Edition, 2017, ISBN: 978-0124166455).*
2. *Data Mining and Analysis: Fundamental Concepts and Algorithms by Mohammed J. Zaki and Wagner Meira Jr., published by Cambridge University Press (1st Edition, 2014, ISBN: 978-0521766333).*
3. *Core Concepts in Data Analysis: Summarization, Correlation and Visualization by Boris Mirkin, published by Springer (1st Edition, 2011, ISBN: 978-0857292872).*
4. *Applied Predictive Modeling, Max Kuhn and Kjell Johnson, 1st Edition, Springer, 2013.*
5. *An Introduction to Statistical Learning with Applications in R, Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani, 1st Edition, Springer, 2013.*
6. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Trevor Hastie, Robert Tibshirani and Jerome Friedman, 2nd Edition, Springer, 2009.*
7. *An Introduction to Statistical Learning with Applications in R by Gareth James, Daniela Witten, Trevor Hastie, and Robert Tibshirani, published by Springer (2nd Edition, 2021, ISBN: 978-1071614174).*
8. *Time Series Analysis: Forecasting and Control by George E. P. Box, Gwilym M. Jenkins, Gregory C. Reinsel, and Greta M. Ljung, published by Wiley (5th Edition, 2015, ISBN: 978-1118675021)*

UNIT 3 PREDICTIVE DATA ANALYSIS

Structure

3.1 INTRODUCTION

Descriptive, diagnostic, predictive, and prescriptive analyses together form a **continuum of data analytics**, where each stage builds logically upon the previous one to support increasingly informed decision-making. Understanding this interconnection is essential for students of data science, as it clarifies how raw data is progressively transformed into actionable insights and optimal decisions.

Descriptive analysis is the foundation of this continuum. It focuses on summarizing and organizing historical data to answer the question *"What has happened?"*. Using measures of central tendency, dispersion, frequency distributions, and visualizations, descriptive analysis provides a clear and structured view of past data. However, it does not explain the reasons behind observed patterns.

Building on this foundation, **diagnostic analysis** seeks to answer *"Why did it happen?"*. It examines relationships, dependencies, and underlying causes using techniques such as correlation analysis, data mining, and drill-down analysis. Diagnostic analysis relies heavily on descriptive results, as understanding what occurred is a prerequisite for investigating why it occurred.

Once the past has been described and explained, **predictive data analysis** extends the analytical process by addressing the question *"What is likely to happen next?"*. Predictive analysis uses historical data patterns uncovered through descriptive and diagnostic analyses to build statistical and machine learning models, such as regression, time series forecasting, and classification. These models estimate future outcomes and trends, enabling anticipation of events rather than mere reaction to them.

Finally, **prescriptive data analysis** represents the most advanced stage in the analytics continuum. It answers the question *"What should be done?"* by combining predictions with optimization, simulation, and decision models. Prescriptive analysis uses the outputs of predictive models to recommend actions that lead to optimal or desired outcomes under given constraints. The prescriptive data analysis is the forte of the subsequent unit of this block.

Overall, descriptive analysis provides understanding, diagnostic analysis provides explanation, predictive analysis provides foresight, and prescriptive analysis provides guidance. Having already covered descriptive and diagnostic analyses in earlier units, the current unit on predictive data analysis naturally follows by focusing on forecasting future outcomes. This progression sets the stage for the subsequent unit on prescriptive data analysis, where data-driven recommendations and decision strategies are developed.

Let's begin this unit on Predictive data analysis.

3.2 EXPECTED LEARNING OUTCOMES

After completing this unit, you should be able to:

- Understand Predictive Analysis
- Understand Regression Analysis
- Understand Time Series Analysis
- Understand Classification
- Understand Clustering

3.3 PREDICTIVE ANALYSIS

Predictive data analysis is an advanced stage of data analytics that focuses on answering the question "What is likely to happen in the future?" by using historical data, statistical techniques, and machine learning models. While descriptive and diagnostic analyses deal with understanding past data and identifying the reasons behind observed outcomes, predictive analysis builds upon these insights to forecast future trends, behaviours, and events. For students of data science, predictive analysis forms a crucial bridge between understanding data and applying it to real-world decision-making.

At the core of predictive data analysis lies the idea that patterns observed in historical data can be used to make informed predictions about future outcomes. By identifying relationships, trends, and structures within data, predictive models estimate the likelihood of future events and quantify uncertainty. These predictions support decision-making in diverse domains such as business forecasting, healthcare risk assessment, financial modelling, demand planning, and recommendation systems.

The Predictive Data Analytics involves the following:

- Regression Analysis,
- Time Series Analysis,
- Classification,
- Clustering

Regression analysis is one of the fundamental techniques used in predictive analysis. It models the relationship between a dependent variable and one or more independent variables to predict continuous outcomes, such as sales, prices, or demand. Regression not only provides predictions but also helps quantify the impact of each predictor, making it both a predictive and explanatory tool.

Time series analysis is another important component of predictive data analysis, particularly when data is collected over time. It focuses on identifying temporal patterns such as trends, seasonality, and cycles, and uses these patterns to forecast future values. Time series models are widely used in applications like stock price forecasting, weather prediction, and sales demand estimation.

Classification techniques are used in predictive analysis when the outcome variable is categorical. These methods assign new observations to predefined classes based on learned patterns from historical data. Classification models are commonly applied in areas such as fraud detection, disease diagnosis, customer churn prediction, and credit risk assessment, where the goal is to predict class membership rather than numerical values.

Clustering, although primarily an unsupervised learning technique, plays a supporting role in predictive analysis by discovering hidden group structures within data. By segmenting data into meaningful clusters, clustering helps improve predictive performance, personalize predictions, and identify patterns that influence future behavior. For example, customer segmentation through clustering can enhance targeted marketing predictions.

In this unit, students will learn how predictive data analysis integrates regression analysis, time series analysis, classification, and clustering to build models that anticipate future outcomes. Emphasis is placed on understanding the assumptions, strengths, and limitations of each technique, as well as their practical relevance in real-world data science applications. This unit equips learners with the foundational knowledge required to design, evaluate, and interpret predictive models effectively, preparing them for advanced analytics and machine learning tasks.

☞ **Check Your Progress 1**

Q1. Explain the continuum of data analytics from descriptive to prescriptive analysis.

……………………………………………………………………………………………

……………………………………………………………………………………………

Q2. What is Predictive Data Analysis? How does it differ from descriptive and diagnostic analysis?

…………………………………………………………………………………………………

…………………………………………………………………………………………………

Q3. List and briefly explain the main techniques used in Predictive Data Analysis.…………………………………………………………………………………………………

…………………………………………………………………………………………………

3.3.1 REGRESSION ANALYSIS

Regression analysis is one of the most important and widely used tools in **predictive data analysis**. It helps data scientists **model relationships between variables** and use those relationships to **predict future outcomes**. In predictive analytics, regression answers the core question: *"Given what we know from historical data, what value is likely to occur next?"*

At its core, regression analysis studies how a **dependent variable (output)** changes with respect to one or more **independent variables (inputs)**. By learning this relationship from past data, a regression model can be used to estimate or forecast unknown or future values. Because of this capability, regression forms the backbone of prediction tasks in domains such as sales forecasting, demand estimation, price prediction, risk assessment, and performance analysis.

Regression analysis is highly valued in predictive analytics because it produces predictions that are easy to interpret and understand, clearly explains how different predictors influence the outcome, and supports forecasting as well as scenario-based analysis. Additionally, regression forms the foundation for many machine learning algorithms, which is why regression models are often developed first before progressing to more complex predictive techniques.

Regression analysis serves predictive data analysis in the following ways:

- It **quantifies relationships** between variables
- It **estimates future values** of the dependent variable
- It **measures the impact** of each predictor
- It provides **interpretable models**, which is important for business and policy decisions

For example, a data scientist may use regression to predict future sales based on advertising spend, forecast house prices using location and size, or estimate patient recovery time using medical indicators.

Regression analysis can be classified into several types based on the nature of the dependent variable and the complexity of relationships; a few are listed below:

1. Simple Linear Regression
2. Multiple Linear Regression

3. Polynomial Regression
4. Logistic Regression (Although called regression, it is primarily a **classification-based predictive technique**.)
5. Regularized Regression (Ridge and Lasso)

We will briefly discuss each of the above-mentioned type of regression techniques, to begin with the Simple linear regression:

1. Simple Linear Regression: Simple linear regression is the most basic form of regression analysis. It models the relationship between **one independent variable** and **one dependent variable**, assuming a linear relationship.

The general form of the simple linear regression equation is:

$$Y = a + bX$$

$\bar{y}$ $\bar{x}$

Where:

- Y is the dependent variable
- X is the independent variable
- a is the intercept (value of Y when $X = 0$)
- b is the slope (change in Y for a unit change in X)

The slope b is calculated as:

$$b = \frac{\sum(x - \bar{x})(y - \bar{y})}{\sum(x - \bar{x})^2}$$

and the intercept a is:

$$a = \bar{y} - b\,\bar{x}$$

Example: Suppose a data scientist wants to predict **sales (Y)** based on **advertising spend (X)**.

X (Ad Spend)	**Y (Sales)**
10	40
20	50
30	60
40	70
50	80

Mean values: $\bar{x} = 30$, $\bar{y} = 60$

Using the formulas:

$$b = 1, a = 30$$

Regression equation:

$$Y = 30 + 1X$$

This model can now be used for prediction. If advertising spend is ₹60, predicted sales will be:

$$Y = 30 + 60 = 90$$

This illustrates how regression converts historical patterns into predictive capability.

2. Multiple Linear Regression: Multiple linear regression extends simple linear regression by using **two or more independent variables** to predict a single dependent variable. This is more realistic for real-world predictive problems.

The general equation is:

$$Y = a + b_1X_1 + b_2X_2 + \cdots + b_nX_n$$

Where:

- $X_1, X_2, \ldots, X_n$ are independent variables
- $b_1, b_2, \ldots, b_n$ measure the impact of each variable

Multiple regression is widely used in predictive analytics when outcomes depend on several factors simultaneously, such as predicting house prices using size, location, and age of property.

Note: In the case of multiple linear regression, the problem involves predicting a continuous outcome using more than one influencing factor.

In the problem given below, the objective is to estimate the price of a house by considering multiple independent variables such as the size of the house and the number of bedrooms. Since house price is influenced by several factors simultaneously, this problem requires a regression model that can capture the combined effect of multiple predictors on a single dependent variable, making multiple linear regression an appropriate choice.

Example: Given the Regression equation $Y = 20 + 0.05X_1 + 10X_2$ Predict the price of a house of size 100 sq. m with 3 bedrooms.

Solution:

Here, we are to Predict house price (Y) based on size (X_1) and number of bedrooms (X_2).

This is the case of multiple linear regression, because the problem involves predicting a continuous outcome i.e. house price (Y) using more than one influencing factor i.e. size (X_1) and number of bedrooms (X_2).

Thus, the Multiple regression Model to be used will be $Y = a + b_1X_1 + b_2X_2$

Given Regression Equation

$$Y = 20 + 0.05X_1 + 10X_2$$

Where:

- X_1= House size (sq. meters)
- X_2= Number of bedrooms
- Y= Price (in lakhs)

Hence, the prediction for a house of 100 sq. m with 3 bedrooms:

$$Y = 20 + 0.05(100) + 10(3)Y = 20 + 5 + 30 = 55$$

The predicted house price is ₹55 lakhs. Size and number of bedrooms both positively influence price, with bedrooms having a stronger impact.

3. Polynomial Regression: Polynomial regression is used when the relationship between variables is **non-linear** but still follows a predictable curve.

The general form is:

$$Y = a + bX + cX^2 + dX^3 + \cdots$$

Although it is called non-linear, polynomial regression is linear in terms of parameters and is often used to capture growth, saturation, or curved trends in predictive analysis, such as predicting energy consumption or population growth.

Note: For polynomial regression, the problem focuses on predicting a continuous variable when the relationship between the predictor and the outcome is non-linear.

Consider the example given below, the aim is to estimate electricity consumption based on temperature. As temperature increases, electricity usage does not increase at a constant rate but follows a curved pattern. Polynomial regression is suitable here because it can model such non-linear relationships by incorporating higher-degree terms of the independent variable.

Example: Given the Polynomial regression model as $Y = 10 + 2X + 0.5X^2$, Predict electricity consumption (Y) based on temperature (X) for 4 units of Temperature.

Solution: Here the relationship is non-linear (parabolic) because $Y = 10 + 2X + 0.5X^2$

As the Given Equation is $Y = 10 + 2X + 0.5X^2$ therefore Prediction For X = 4 units will be:

$$Y = 10 + 2(4) + 0.5(16)Y = 10 + 8 + 8 = 26$$

Electricity consumption increases at an increasing rate as temperature rises, showing a non-linear predictive relationship.

4. Logistic Regression: Logistic regression is used when the dependent variable is **categorical**, usually binary (0 or 1). Instead of predicting a numerical value, it predicts the **probability of an event**.

The logistic regression model is:

$$P(Y = 1) = \frac{1}{1 + e^{-(a+b\ \ )}}$$

Logistic regression is widely used in predictive data analysis for:

- Credit default prediction
- Disease diagnosis
- Fraud detection
- Customer churn prediction

Note:

- Although called regression, it is primarily a **classification-based predictive technique**.
- The logistic regression problem involves predicting a categorical outcome rather than a numerical value.

Important: The objective of the problem given below is to determine whether a customer is likely to churn or not based on usage behavior. Since the outcome has two possible categories—churn or no churn—logistic regression is used to estimate the probability of churn. This makes it a valuable predictive tool for classification problems where decision-making is based on likelihood rather than exact numerical prediction.

Example: Given Equation the equation for logistic regression as below:

$$P(Y = 1) = \frac{1}{1 + e^{-(-4+0.8X)}}$$

Predict whether a customer will churn (Yes = 1 / No = 0) based on monthly usage (X).
Model

Solution: In the given equation $P(Y = 1) = \frac{1}{1+e^{-(a+bX)}}$; it is to be noted that $a + bX$ is the equation of simple linear regression only where a & b are the intercept and slope respectively.

Now, the Given Equation is $P(Y = 1) = \frac{1}{1+e^{-(-4+0.8X)}}$ we need to make prediction for X = 6:

$$P = \frac{1}{1 + e^{-(-4+4.8)}} = \frac{1}{1 + e^{-0.8}} \approx 0.69$$

The outcome can be interpreted that there is a 69% probability that the customer will churn. Logistic regression predicts probabilities, making it ideal for classification-based prediction.

5. Regularized Regression (Ridge and Lasso)

In predictive modeling with many variables, issues such as multicollinearity and overfitting may arise. Regularized regression techniques address these problems.

- **Ridge Regression** adds a penalty on squared coefficients
- **Lasso Regression** adds a penalty on absolute coefficients and performs variable selection

These techniques improve predictive accuracy and model stability, especially in high-dimensional data science problems.

In regularized regression, the problem centers on predicting a continuous outcome when there are many independent variables that may be correlated with each other.

For Example: The goal is to estimate an exam score using multiple predictors such as study hours, coaching hours, online practice, and attendance. Because these predictors may overlap in the information they provide, regularization techniques like Ridge and Lasso regression are applied to control model complexity, reduce overfitting, and improve prediction accuracy. Ridge regression shrinks coefficients to manage multicollinearity, while Lasso regression additionally helps identify and retain only the most important predictors by eliminating less relevant ones.

We restrict our discussion on Regularized Regression (Ridge and Lasso) up to this only.

Now we are going to extend our discussion to the Time series analysis in the subsequent section of this unit.

☞ Check Your Progress 2

Q1. What is Regression Analysis? Explain its role in Predictive Data Analysis

……………………………………………………………………………………………………

……………………………………………………………………………………………………

Q2. Differentiate between Simple Linear Regression and Multiple Linear Regression

……………………………………………………………………………………………………

……………………………………………………………………………………………………

Q3. Explain Logistic Regression and Polynomial Regression with their use cases

……………………………………………………………………………………………………

……………………………………………………………………………………………………

3.3.2 TIME SERIES ANALYSIS

Time Series Analysis is a statistical and analytical technique used to study data points collected sequentially over time, with the objective of identifying patterns such as trend, seasonality, and cyclic behaviour, and using these patterns to forecast future values. Unlike other forms of data analysis where the order of observations may not matter, time series analysis treats time as a critical dimension, meaning that the sequence and spacing of observations directly influence the analysis and results.

In the context of predictive data analytics, time series analysis plays a central role because it focuses explicitly on predicting future outcomes based on historical time-dependent data. By learning from past behaviour, time series models enable data scientists to forecast variables such as future sales, demand, stock prices, energy consumption, and website traffic. Predictive analytics relies on these forecasts to anticipate events, plan resources, and make proactive decisions.

There is a strong conceptual and methodological connection between regression analysis and time series analysis in predictive data analytics. Regression analysis models the relationship between a dependent variable and one or more independent variables, and this framework can be extended to time-based data. In time series analysis, regression is often used by incorporating time or lagged variables as predictors. For example, future sales can be predicted using past sales values, seasonal indicators, or external variables such as promotions and economic indicators. Models such as trend regression, distributed lag models, and time series regression with exogenous variables (ARIMAX) illustrate how regression principles are embedded within time series forecasting.

However, time series analysis goes beyond traditional regression by explicitly accounting for autocorrelation, where current values depend on past values of the same variable. Techniques like ARIMA combine regression-like components with moving averages of past errors to improve predictive accuracy. Thus, while regression focuses on explaining relationships between variables at a point in time, time series analysis emphasizes modeling temporal dependence for forecasting.

Overall, we can say that the time series analysis is a key component of predictive data analytics, designed specifically for forecasting time-dependent data. Regression analysis and time series analysis are closely connected, with regression concepts forming the foundation of many time series models. Together, they enable data scientists to build robust predictive models that capture both relationships among variables and their evolution over time.

For better clarity on Time Series Analysis firstly we need to understand few basic concepts, terms and definitions

Let's begin by understanding what is a Time Series? one can express that - "a time series is a sequence of observations recorded over time at regular intervals such as daily, monthly, quarterly, or yearly".

Unlike cross-sectional data, time series data is ordered in time, and the temporal order is crucial for analysis and modeling. Examples include daily stock prices, monthly sales figures, and annual rainfall data.

A time series typically consists of four main components :

1. The Trend (T) represents the long-term movement or direction of the data, such as a steady increase in sales over years.
2. The Seasonal component (S) captures regular and repeating patterns within a fixed period, such as higher sales during festive seasons.
3. The Cyclical component (C) reflects long-term fluctuations caused by economic or business cycles that do not follow a fixed period. T
4. The Irregular or Random component (I) accounts for unpredictable variations due to random or unforeseen events, such as natural disasters or sudden market shocks.

Common examples of time series data include daily temperature readings, monthly electricity consumption, quarterly GDP growth rates, annual population figures, hourly website traffic, and weekly product demand. In data science, time series data is widely used for forecasting, monitoring performance, and detecting anomalies over time.

Time series analysis involves studying past observations to understand underlying patterns and to forecast future values. It focuses on identifying trend, seasonality, and other structural components in the data. Techniques used in time series analysis range from simple methods like moving averages and exponential smoothing to advanced models such as ARIMA and state-space models. In predictive data analysis, time series modeling plays a vital role in forecasting future trends, supporting planning, and enabling proactive decision-making across domains such as finance, economics, healthcare, and operations.

A systematic learning path for time series analysis begins with basic methods such as moving averages and exponential smoothing, progresses to intermediate techniques like Auto Correlation, Holt–Winters and time series decomposition, advances to statistical models such as ARIMA and SARIMA, and finally machine learning and deep learning approaches like LSTM, GRU, and transformer-based models. This progression enables learners to move from simple forecasting to advanced predictive time series modeling in data science.

In this section we restrict our discussion to a few of the mentioned techniques, and they are listed below:

1. Moving Averages
2. Exponential Smoothening
3. Auto Correlation

Let's Begin with the first one i.e. the technique of Moving Averages for Time Series Analysis

Moving averages are one of the simplest and most widely used techniques in time series analysis, especially at the initial stage of predictive data analytics. They are primarily used to smooth short-term fluctuations in time series data and highlight the underlying trend. Because of their simplicity and interpretability, moving averages are often the first forecasting tool introduced to students of data science.

The moving average method works by **averaging a specified number of recent observations** to reduce random noise and irregular variations in the data. As the averaging window moves forward, older values are dropped and newer values are included. This process smooths the series and makes long-term patterns easier to identify.

Moving averages are particularly useful when:

- The data shows **random fluctuations**
- The goal is **trend identification**
- Short-term forecasting is required
- The time series does not have strong seasonality

However, moving averages are less effective when data contains strong seasonal patterns unless special seasonal moving averages are used.

Formally one can define Moving Averages as Below:

A **moving average** is defined as the average of a fixed number of consecutive observations in a time series, calculated repeatedly by moving the window forward one time period at a time. Each new average replaces the oldest observation with the newest one, which is why it is called a "moving" average.
Mathematically, for a time series $Y_1, Y_2, \ldots, Y_n$, the **k-period moving average** at time $t$is given by:

$$Moving\ Average_t = \frac{Y_t + Y_{t-1} + \cdots + Y_{t-k+1}}{k}$$

There are various types of Moving Averages:

1. Simple Moving Average (SMA)
2. Weighted Moving Average and
3. Exponential Moving averages

The most commonly used type is the **Simple Moving Average (SMA)**, where all observations within the window receive equal weight. More advanced variants, such as **Weighted and Exponential Moving Averages**, assign greater importance to recent observations, but the simple moving average is typically used for introductory and diagnostic analysis.

Example (Simple Moving Average): Consider the following **monthly sales data (in units)** for a product:

Month	Sales
Jan	120
Feb	130
Mar	125
Apr	140
May	150
Jun	160

Compute a **3-month simple moving average** and interpret the result.

Solution:

Step-by-Step Calculation

- **Moving Average for March:** $MA_{Mar} = \frac{120+130+12}{3} = \frac{375}{3} = 125$
- **Moving Average for April:** $MA_{Apr} = \frac{130+125+14}{3} = \frac{395}{3} \approx 131.67$
- **Moving Average for May:** $MA_{May} = \frac{125+140+150}{3} = \frac{415}{3} \approx 138.33$
- **Moving Average for June:** $MA_{Jun} = \frac{140+150+160}{3} = \frac{450}{3} = 150$

Moving Average Table

Month	Sales	3-Month Moving Average
Jan	120	–
Feb	130	–
Mar	125	125.00
Apr	140	131.67
May	150	138.33
Jun	160	150.00

The 3-month moving averages smooth out short-term fluctuations in monthly sales and reveal a **clear upward trend** over time. While individual sales figures vary from month to month, the moving average steadily increases from 125 in March to 150 in June, indicating consistent growth in demand. This insight is valuable for predictive data analysis, as it suggests that future sales are likely to continue rising if current conditions remain unchanged.

However, it is important to note that moving averages **lag behind actual data**, meaning they respond slowly to sudden changes. Therefore, while they are excellent for trend analysis and short-term forecasting, they may not capture abrupt shifts in the time series.

Note: Moving averages are a foundational technique in time series analysis and predictive data analytics. They provide a simple yet powerful way to smooth data, identify trends, and support short-term forecasts. For data science students, understanding moving averages is essential before progressing to more advanced forecasting techniques such as exponential smoothing and ARIMA models.

Weighted Moving Average:

Weighted Moving Averages (WMA) in Time Series Analysis

Weighted Moving Averages are an extension of simple moving averages and are widely used in **time series analysis and predictive data analytics** to smooth data while giving **more importance to recent observations**. They are especially useful when recent data points are believed to better reflect current patterns or future behavior.

Definition of Weighted Moving Average

A **Weighted Moving Average (WMA)** is defined as the average of a fixed number of consecutive observations in a time series, where **different weights are assigned to each observation**, typically giving higher weights to more recent data points. The sum of the weights is usually equal to 1 (or 100%).

Mathematically, for a k-period weighted moving average:

$$WMA_t = \frac{w_1 Y_t + w_2 Y_{t-1} + \cdots + w_k Y_{t-k+1}}{w_1 + w_2 + \cdots + w_k}$$

Where:

- $Y_t, Y_{t-1}, \ldots$are observed values
- $w_1, w_2, \ldots, w_k$are the assigned weights
- $w_1 > w_2 > \cdots > w_k$in most practical cases

Description of the Weighted Moving Average Method

The weighted moving average method works by assigning **unequal importance to observations** within the moving window. Unlike simple moving averages, where each observation contributes equally, WMA reflects the belief that **recent values are more relevant for understanding current trends**. As the window moves forward in time, older observations drop out and newer ones enter the calculation with higher influence.

Weighted moving averages are particularly useful when:

- Recent data is more indicative of future behavior
- The time series shows gradual changes rather than abrupt shifts
- Short-term forecasting is required

However, selecting appropriate weights requires judgment and domain knowledge, making WMA slightly more complex than simple moving averages.

Example: Consider the following **monthly demand data** for a product:

Month	Demand
Jan	100
Feb	110
Mar	105
Apr	120
May	130

Compute a **3-month weighted moving average** using the following weights:

- Most recent month = 3
- Second recent month = 2
- Third recent month = 1

Step-by-Step Calculation

Weighted Moving Average for April

$$WMA_{Apr} = \frac{3(120) + 2(105) + 1(110)}{3 + 2 + 1} = \frac{360 + 210 + 110}{6} = \frac{680}{6} \approx 113.33$$

Weighted Moving Average for May

$$WMA_{May} = \frac{3(130) + 2(120) + 1(105)}{6} = \frac{390 + 240 + 105}{6} = \frac{735}{6} = 122.50$$

Weighted Moving Average Table

Month	Demand	3-Month WMA
Jan	100	–
Feb	110	–
Mar	105	–
Apr	120	113.33
May	130	122.50

Interpretation of the Outcome

The weighted moving averages smooth out irregular month-to-month fluctuations while responding **more quickly to recent increases in demand** compared to a simple moving average. For example, although demand jumps from 120 in April to 130 in May, the WMA reflects this upward trend more strongly because higher weight is assigned to recent months. This makes WMA particularly suitable for **short-term forecasting and trend monitoring** in predictive data analysis.

However, like all moving average methods, WMA still **lags behind actual values** and may not respond immediately to sudden structural changes or shocks in the data.

Thus, it may be concluded that the Weighted moving averages provide a balance between simplicity and responsiveness in time series analysis. By assigning greater importance to recent observations, they offer improved smoothing and better short-term predictive insight than simple moving averages. For data science students, WMA serves as an important stepping stone toward more advanced forecasting methods such as exponential smoothing and ARIMA models.

Note: Exponential Smoothing and Exponential Moving Average (EMA) are closely related but not exactly the same. They are often used interchangeably in practice, but technically they differ in purpose and formulation. EMA is mainly used for smoothing historical data, whereas exponential smoothing is used as a forecasting technique in predictive data analysis.

☞ **Check Your Progress 3**

Q1. What is Time Series Analysis? Why is it important in Predictive Data Analysis?

………
………

Q2. Explain the components of a Time Series.

………
………

Q3. What is Moving Average in Time Series Analysis? Explain its purpose with an example.

………
………

3.3.3 CLASSIFICATION

Classification is one of the most widely used techniques in **Predictive Data Analysis**. It is a supervised learning approach in which a model is trained using a labelled dataset to predict the category or class of new observations. In classification problems, the target variable is categorical, meaning it represents predefined classes such as *yes/no, fraud/not fraud, spam/not spam,* or *high/medium/low risk*. The objective of classification is to learn patterns from historical data so that the model can accurately assign new data points to the appropriate class. This technique is widely applied in areas such as medical diagnosis, credit risk analysis, customer churn prediction, fraud detection, and email spam filtering.

The classification process generally involves several stages, including **data collection, preprocessing, feature selection, model training, evaluation, and prediction**. During preprocessing, missing values, noise, and inconsistencies in the dataset are handled to improve data quality. Feature selection or feature engineering is then used to identify the most relevant variables that influence the prediction outcome. Once the data is prepared, a classification algorithm is trained on historical data to learn relationships between the input features and the target class. The trained model is later used to classify unseen data and generate predictions.

Several classification techniques are commonly used in predictive analytics. **Decision Trees** are one of the most intuitive approaches, where the data is divided into branches based on feature values, leading to a final classification decision. **Logistic Regression** is another popular method that estimates the probability of a binary outcome using a logistic function. **Naïve Bayes**, based on Bayes' theorem, is widely used in text classification and spam filtering because of its efficiency and ability to handle large datasets. **K-Nearest Neighbours (KNN)** classifies data points by comparing them with the closest training examples in the feature space. Advanced classification techniques include **Support Vector Machines (SVM)** and **Artificial Neural Networks (ANNs)**. SVM works by identifying the optimal hyperplane that separates different classes with the maximum margin, making it effective for high-dimensional data. Neural networks, particularly deep learning models, can learn complex patterns and nonlinear relationships in large datasets. These models are widely used in applications such as image recognition, speech processing, and predictive maintenance.

Another important concept in classification is **model evaluation**. To ensure that a classification model performs well, several evaluation metrics are used, such as **accuracy, precision, recall, F1-score, and confusion matrix**. Accuracy measures the overall correctness of predictions, while precision and recall provide deeper insights into the model's ability to correctly identify specific classes. In many predictive analytics tasks, especially those involving imbalanced datasets, relying solely on accuracy may not be sufficient, and other metrics become crucial for proper evaluation.

We learned that Classification is a fundamental technique in **predictive data analysis** used to assign observations into predefined categories or classes based on historical data. It belongs to the family of **supervised learning methods**, where a predictive model is trained using labeled datasets and then used to classify new, unseen data. The objective of classification is to learn a mapping between **input features (independent variables)** and **target classes (dependent variables)**. Common applications include fraud detection, medical diagnosis, credit scoring, spam filtering, and customer churn prediction. Several algorithms are used for classification; among the most widely used are **Naïve Bayes, Logistic Regression, and Decision Trees**.

Now, we will discuss these fundamental algorithms, let's begin with the **Naïve Bayes.**

Naïve Bayes is a probabilistic classification technique based on **Bayes' Theorem**. It assumes that the predictors (features) are **conditionally independent** given the class label. Despite this simplifying assumption, the method performs well in many real-world problems, particularly in text classification and spam filtering.
The classifier calculates the **posterior probability** of each class and assigns the class with the highest probability.

Formula: $P(A \mid B) = \frac{P(A) \cdot P(B|A)}{P(B)}$

Where,

- $P(A \mid B)$: Posterior probability (probability of A given B)
- $P(A)$: Prior probability (initial belief about A)
- $P(B \mid A)$: Likelihood (probability of observing B if A is true)
- $P(B)$: Evidence (overall probability of B)

Note: *Posterior probability is the probability of an event or hypothesis after considering new evidence. Further, we need to understand the difference between Prior & Posterior probabilities, Prior probability is your initial belief but the posterior probability is your belief after evidence. Also, posterior probability changes as new data arrives, making it more realistic than static probabilities, and it allows continuous learning and updating of models.*

Example

*Imagine a forest with **20% oak trees** and **80% maple trees**.*

- *Prior probability:* $P(Oak) = 0.2$
- *Suppose you observe a leaf shape that is **90% likely if oak** and **10% likely if maple**.*
- *Posterior probability of the tree being oak given the leaf shape:*

$$P(Oak \mid Leaf) = \frac{0.2 \cdot 0.9}{(0.2 \cdot 0.9) + (0.8 \cdot 0.1)} = 0.69$$

*So, after seeing the leaf, the probability of the tree being oak rises from **20% to 69%***

Let's, attempt **one more example**, exhibiting Email Spam Classification by using Naïve Bayes classification.

Example: Consider a binary classification problem where an email must be classified as Spam or Not Spam.

- Prior probabilities: $P(Spam) = 0.4$, $P(Not\ Spam) = 0.6$.
- Likelihoods:
  - $P(\text{“}discount\text{”} \mid Spam) = 0.7$
  - $P(\text{“}discount\text{”} \mid Not\ Spam) = 0.1$

Using Bayes' theorem, calculate the posterior probability that an email is Spam given that it contains the word "*discount*". Based on your calculation, state the class label assigned to the email.

Solution: *Here, we want to classify an email as **Spam** or **Not Spam** based on the presence of the word "discount".*

Step 1: Prior Probabilities

- $P(Spam) = 0.4$ *(40% of emails are spam)*
- $P(Not\ Spam) = 0.6$

Step 2: Likelihoods

- $P(\text{“}discount\text{”} \mid Spam) = 0.7$*(70% of spam emails contain “discount”)*
- $P(\text{“}discount\text{”} \mid Not\ Spam) = 0.1$

Step 3: Evidence

$$P(\text{“}discount\text{”}) = (0.4 \cdot 0.7) + (0.6 \cdot 0.1) = 0.28 + 0.06 = 0.34$$

Step 4: Posterior Probabilities

- *For Spam:*

$$P(Spam \mid \text{“}discount\text{”}) = \frac{0.4 \cdot 0.7}{0.34} = \frac{0.28}{0.34} \approx 0.82$$

- *For Not Spam:*

$$P(Not\ Spam \mid \text{“}discount\text{”}) = \frac{0.6 \cdot 0.1}{0.34} = \frac{0.06}{0.34} \approx 0.18$$

The **interpretation** of the results says that before seeing the word “discount,” we thought 40% of emails were spam. But, after seeing “discount,” the posterior probability jumps to 82% spam. The classifier would label this email as Spam.

Now, we are going to extend our discussion to Logistic regression as the next classification technique.

Logistic Regression *is a widely used classification technique that models the probability of a binary outcome using the* ***logistic (sigmoid) function****. Unlike linear regression, which predicts continuous values, logistic regression predicts* ***probabilities between 0 and 1****, which are then converted into class labels.*

The ***logistic regression probability function*** *can be written properly as:*

$$P(Y = 1 \mid X) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X)}}$$

Where,

- $P(Y = 1 \mid X)$→ *Probability that the dependent variable Y equals 1 given predictor X*
- β_0→ *Intercept (constant term)*
- β_1→ *Coefficient of the predictor variable*
- X→ *Independent variable*
- e→ *Euler's number (≈ 2.718)*

Note: This equation converts the linear combination $(\beta_1 X)$into a probability value between 0 and 1 using the sigmoid (logistic) function. If the resulting probability is greater than a threshold (commonly 0.5), the observation is classified as Class 1, otherwise Class 0.

Further, in theoretical discussions of ***Logistic Regression*** *within Predictive Data Analysis, the model is often expressed in terms of* ***log-odds (logit form)*** *rather than the probability form. As per, the* ***Log-Odds (Logit) Form of Logistic Regression formula :***

$$log\left(\frac{P(Y = 1 \mid X)}{1 - P(Y = 1 \mid X)}\right) = \beta_0 + \beta_1 X$$

where,

- $P(Y = 1 \mid X)$→ *Probability that the dependent variable* $Y = 1$*given predictor* X
- $1 - P(Y = 1 \mid X)$→ *Probability that* $Y = 0$
- $\frac{P(Y=1|X)}{1-P(Y=1|X)}$→ ***Odds*** *of the event occurring*
- $log\ log\left(\frac{P}{1-P}\right)$ → ***Log-odds (logit transformation)***
- β_0→ *Intercept (constant term)*
- β_1→ *Coefficient of predictor X*

Note: It can be observed that the Logistic regression assumes that the log-odds of the outcome are linearly related to the predictors. Further, A one-unit increase in Xchanges the log-odds of the event by β_1, and the Exponentiating the coefficient gives the odds ratio: e^{β_1} which represents how the odds of the outcome change for a one-unit increase in X.

Also, we need to understand, *Why Logit Form is Used?* Actually it converts probabilities (0–1 range) into unbounded values (−∞ to +∞), and allows the relationship between predictors and the response variable to be modelled linearly. Also, it makes parameter estimation easier using Maximum Likelihood Estimation (MLE).

Example: *Suppose a logistic regression model predicts whether a student will* ***Pass (1)*** *or* ***Fail (0)*** *based on hours of study, the given Model equation is* $z = -3 + 0.8X$

Solution: For a student studying ***5 hours,*** $z = -3 + 0.8(5) = -3 + 4 = 1$
Thus, based on the ***logistic regression probability function****:*

$$P(Y = 1 \mid X) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X)}}$$

hence, the Probability:

$$P = \frac{1}{1+e^{-1}} P \approx 0.731$$

It can be interpreted from the results that the probability of passing is **0.731 (73.1%)**. Since the probability is **greater than 0.5**, the model predicts that the student will **pass**. Thus, logistic regression converts predictor variables into probabilities and classifies observations based on a decision threshold.

Now we are going to discuss the ***Decision Tree classification,*** the Decision Tree Classification is a tree-structured classification model where decisions are made through a series of hierarchical rules. The dataset is recursively split based on features that provide the maximum information gain or minimum impurity.
*Decision Trees use metrics such as **Entropy** and **Information Gain** to determine the best split. The **Entropy formula** used in **Decision Tree classification (information theory)** can be written properly as:*

$$Entropy(S) = -\sum_{i=1}^{n} p_i log_{\,2}(p_i)$$

Where,

- *S→ The dataset or sample under consideration*
- *n→ Number of classes in the dataset*
- *p_i→ Probability of occurrence of class $i$in dataset S*
- *$log_{\,2}$→ Logarithm to base 2*

Note: Entropy measures the degree of impurity or uncertainty in the dataset. Entropy = 0 implies that the Dataset is perfectly pure (all observations belong to one class), and Entropy = 1 (for binary classification) implies Dataset is maximally impure (classes equally distributed).
*This measure is used in **Decision Tree algorithms (like ID3, C4.5)** to select the best attribute for splitting the data through **Information Gain**,*

$$Gain(S, A) = Entropy(S) - \sum \frac{|S_v|}{|S|} Entropy(S_v)$$

Where

- *S= dataset*
- *A= attribute*
- *S_v= subset after split*

Example: *Suppose we classify whether a person **plays tennis** based on weather.*

Weather	***Play***
Sunny	*No*
Sunny	*No*
Overcast	*Yes*
Rain	*Yes*

Solution :

Step 1: Calculate Entropy: Entropy of the whole dataset

There are 4 records in total:

- *Yes = 2*
- *No = 2*

So,

$$p(Yes) = \frac{2}{4} = 0.5 p(No) = \frac{2}{4} = 0.5$$

Using the entropy formula:

$$Entropy(S) = -\sum_{i=1}^{n} p_i log_{\,2}(p_i) Entropy(S) = -(0.5 log_{\,2} 0.5)$$

Since,

$$log_{\,2} 0.5 = -1 Entropy(S) = -(0.5(-1) + 0.5(-1)) Entropy(S) = -(-0.5 - 0.5) Entropy(S) = 1$$

Step 2: Calculate Entropy after Split on Weather: *Now split the dataset by the attribute* ***Weather****.*

a) Sunny

Records:

Weather	***Play***
Sunny	*No*
Sunny	*No*

Here:

- *Yes = 0*
- *No = 2*

$$p(Yes) = 0, p(No) = 1 Entropy(Sunny) = -(1 log_{\,2} 1) = 0$$

b) Overcast

Records:

Weather	***Play***
Overcast	*Yes*

Here:

- *Yes = 1*
- *No = 0*

$$Entropy(Overcast) = 0$$

c) Rain

Records:

Weather	***Play***
Rain	*Yes*

Here:

- *Yes = 1*

- *No = 0*

$$Entropy(Rain) = 0$$

Step 3 : Weighted Entropy after split

Formula:

$$Entropy\ after\ split = \sum \frac{|S_v|}{|S|} Entropy(S_v)$$

Here:

- *Sunny subset size = 2*
- *Overcast subset size = 1*
- *Rain subset size = 1*
- *Total size = 4*

So,

$$Entropy\ after\ split = \frac{2}{4}(0) + \frac{1}{4}(0) + \frac{1}{4}(0) Entropy\ after\ split = 0$$

Step 4: calculate the Information Gain:

Formula:

$$Information\ Gain(S, Weather) = Entropy(S) - Entropy\ after\ split$$

Substituting the values we get:

$$Information\ Gain(S, Weather) = 1 - 0 = 1$$

The results can be interpreted as follows:

From the calculations, the entropy of the original dataset was found to be 1, which indicates maximum uncertainty. This occurs because the dataset contains an equal number of positive and negative outcomes (2 Yes and 2 No). In such a situation, it is difficult to predict the class of a new observation without using additional attributes. When the dataset is split using the attribute Weather, each resulting subset becomes perfectly pure. The Sunny subset contains only "No" outcomes, while both Overcast and Rain subsets contain only "Yes" outcomes. Because each subset contains observations from only one class, the entropy of each subset becomes 0, meaning there is no uncertainty within those groups.

The weighted entropy after the split is therefore 0, since all subsets are perfectly pure. When this value is substituted into the information gain formula, the Information Gain = 1 − 0 = 1. This is the maximum possible information gain, indicating that the attribute Weather completely removes uncertainty from the dataset.

The result implies that Weather is the most informative attribute for predicting whether a person will play tennis in this example. In decision tree construction, the attribute with the highest information gain is selected as the root node, so Weather would be chosen as the first splitting criterion in the decision tree.

Overall, the analysis shows that the Weather attribute perfectly separates the classes, producing a fully accurate classification rule for this dataset. In practical predictive data analysis problems, such perfect splits are rare; however, the same principle is applied to select the attribute that most effectively reduces uncertainty at each step of the decision tree.

It can be concluded that the Weather attribute perfectly classifies the dataset, resulting in zero entropy after the split. Therefore, the decision tree will use Weather as the root node, allowing clear classification of future observations. Here, the original dataset has entropy 1, which means it is fully impure / uncertain. Now, after splitting on Weather, entropy becomes 0, which means each subset is perfectly pure. Therefore, Weather is an ideal attribute for splitting this dataset $Weather\ perfectly\ classifies\ whether\ the\ person\ plays\ tennis\ or\ not.$

With this we are going to conclude this section, here we learned that classification plays a critical role in predictive data analysis by enabling organizations to make informed decisions based on historical data patterns. By applying appropriate classification algorithms and evaluation techniques, predictive models can effectively categorize future observations and support decision-making processes in fields such as business analytics, healthcare, finance, and cybersecurity. As datasets continue to grow in size and complexity, classification techniques combined with modern machine learning approaches will remain a key component of predictive analytics systems.

Now we are going to extend our discussion to next topic, “Clustering”

☞ Check Your Progress 4

Q1. What is Classification in Predictive Data Analysis? Explain its key characteristics.

……
……

Q2. Explain the steps involved in the classification process.

……
……

Q3. Discuss any two classification algorithms with examples.

……
……

3.3.4 CLUSTERING

Clustering is an important technique used in predictive data analysis to group similar data points into clusters based on shared characteristics or patterns. Unlike classification, clustering is an unsupervised learning method, meaning that it does not rely on predefined class labels. Instead, it identifies natural groupings within the dataset by measuring similarities or distances between observations. The main objective of clustering is to discover hidden structures in large datasets, which can help analysts understand patterns, segment data, and support decision-making processes in fields such as marketing, healthcare, finance, and social sciences.

In predictive analytics, clustering is often used as a preliminary data exploration technique. By grouping similar observations together, analysts can identify patterns that may not be immediately visible in raw data. For example, businesses frequently use clustering to perform customer segmentation, where customers with similar purchasing behaviors, preferences, or demographics are grouped together. These clusters can then be used to design targeted marketing strategies, predict customer behavior, or improve service personalization. Similarly, clustering is applied in fraud detection, document classification, and image analysis to identify meaningful patterns in complex datasets.

Several clustering techniques are commonly used in predictive data analysis. One of the most widely used methods is K-Means clustering, which partitions the dataset into a predefined number of clusters (k). The algorithm works by assigning data points to the nearest cluster centroid and iteratively updating the centroid positions until the clusters stabilize. K-Means is computationally efficient and works well for large datasets, although it requires the number of clusters to be specified in advance and may be sensitive to initial centroid selection.

Another important clustering approach is Hierarchical Clustering, which builds a hierarchy of clusters either by progressively merging smaller clusters (agglomerative approach) or by dividing larger clusters into smaller ones (divisive approach). This technique produces a dendrogram, a tree-like diagram that illustrates the relationships between clusters at different levels of similarity. Hierarchical clustering is particularly useful when the number of clusters is unknown, as it allows analysts to visualize the cluster structure and choose an appropriate level of grouping.

A more advanced clustering technique is Density-Based Spatial Clustering of Applications with Noise (DBSCAN). Unlike K-Means, DBSCAN identifies clusters based on the density of data points rather than distance to centroids. It is especially effective for detecting clusters of arbitrary shapes and for identifying outliers or noise points in the dataset. Because of these characteristics, DBSCAN is widely used in applications such as geographic data analysis, anomaly detection, and network intrusion detection.

In this section we are going to discuss one of the most widely used clustering algorithms is **K-Means clustering**. The primary objective of the K-Means algorithm is to partition a dataset into **K distinct clusters**, where each data point belongs to the cluster whose centroid (mean point) is closest to it. The algorithm works iteratively by first selecting initial centroids, assigning each observation to the nearest centroid, and then recalculating the centroid of each cluster until the assignments stabilize. This process minimizes the variation within each cluster and ensures that similar observations are grouped together.

The basic Characteristics of K-Means are:

1. First, it is computationally efficient and works well with large datasets, making it widely used in predictive analytics.
2. Second, the algorithm requires the number of clusters $k$to be specified beforehand, which may sometimes require experimentation or techniques such as the Elbow Method to determine the optimal number of clusters.
3. Third, K-Means performs best when clusters are spherical and similar in size, as it relies on distance-based calculations.
4. Finally, the algorithm is sensitive to the initial selection of centroids, which can affect the final clustering result.

The Mathematical Formulation of K-Means states that the K-Means algorithm minimizes the within-cluster sum of squares (WCSS), also known as the clustering objective function.The K-Means objective function (Within-Cluster Sum of Squares) can be written properly as:

$$J = \sum_{i=1}^{k} \sum_{x \in C_i} \parallel x - \mu_i \parallel^2$$

Where,

- J→ Objective function (total within-cluster variance)
- k→ Number of clusters
- C_i→ Set of data points belonging to cluster i
- x→ A data point in cluster C_i
- μ_i→ Centroid (mean) of cluster i
- $\parallel x - \mu_i \parallel^2$→ Squared Euclidean distance between the data point and its cluster centroid

Note: The K-Means algorithm tries to minimize J. This means it aims to place data points into clusters such that the distance between each point and its cluster centroid is as small as possible, resulting in compact and well-separated clusters.

The **Euclidean distance** commonly used to measure similarity between points is:

$$d(x, y) = \sqrt{\sum_{j=1}^{n} (x_j - y_j)^2}$$

This distance measure helps determine which centroid a data point should belong to during the clustering process.

Let's understand K-Means algorithm with an Example, given below:

Example: Consider a simple dataset representing customer purchasing behavior based on two variables: **Annual Spending Score** and **Income Level**. Suppose the following four data points are observed:

Customer	Income (X)	Spending Score (Y)
A	1	1
B	1.5	2
C	5	6
D	6	7

Assume we want to divide the dataset into **K = 2 clusters**.

Solution:

Step 1: Determine Initial Centroids

Let the initial centroids be:

$$C_1 = (1,1) \quad , \quad C_2 = (5,6)$$

Step 2: Calculate Distance between Each Point and Centroids

Example for point B (1.5,2):

Distance to C_1:

$$d = \sqrt{(1.5-1)^2 + (2-1)^2} d = \sqrt{0.25+1} = \sqrt{1.25} \approx 1.118$$

Distance to C_2:

$$d = \sqrt{(1.5-5)^2 + (2-6)^2} d = \sqrt{12.25+16} = \sqrt{28.25} \approx 5.31$$

Thus, point **B belongs to Cluster 1**.

Applying the same calculations to all points gives:

Cluster 1 → A, B

Cluster 2 → C, D

Step 3: Recalculate Centroids

Cluster 1 centroid:

$$\mu_1 = \left(\frac{1+2}{2}\right) \mu_1 = (1.25,1.5)$$

Cluster 2 centroid:

$$\mu_2 = \left(\frac{6+7}{2}\right) \mu_2 = (5.5,6.5)$$

Since the assignments remain the same after re-computation, the algorithm converges.

It can be Interpretated from the Results that The final clustering divides the dataset into two meaningful groups. Cluster 1 contains customers with lower income and lower spending scores (A and B), while Cluster 2 contains customers with higher income and higher spending scores (C and D). This indicates that the algorithm successfully identified two natural customer segments. In predictive data analysis, such clustering results can be used for applications such as customer targeting, personalized marketing strategies, and behavioural analysis, where businesses design different strategies for different customer groups.

We learned that the K-Means clustering is a powerful unsupervised learning technique widely used in predictive data analysis to identify hidden patterns in datasets. By minimizing the within-cluster variance and grouping similar observations together, the algorithm enables analysts to discover meaningful structures within data. When combined with other analytical techniques, clustering helps organizations better understand their data, improve prediction models, and support data-driven decision-making.

Finally the unit concludes with the understanding that clustering plays a significant role in predictive data analysis by enabling the discovery of meaningful patterns and relationships within unlabelled datasets. Techniques such as K-Means, Hierarchical Clustering, and DBSCAN provide different approaches for grouping data depending on the structure and complexity of the dataset. By applying clustering methods, organizations can gain valuable insights, enhance predictive models, and support data-driven decision making across various domains.

☞ **Check Your Progress 5**

Q1. What is Clustering in Predictive Data Analysis? Explain its purpose

………
………

Q2. Differentiate between Classification and Clustering.

………
………

Q3. Explain the role of clustering in predictive data analysis with an example.

………
………

3.4 SUMMARY

This unit focused on the concepts of **Predictive Data Analysis**, which represents an important stage in the broader continuum of data analytics that includes descriptive, diagnostic, predictive, and prescriptive analysis. While descriptive and diagnostic analytics help understand past data and identify the reasons behind observed outcomes, predictive analysis aims to answer the question **"What is likely to happen in the future?"** by using historical data, statistical techniques, and machine learning models. This unit explains how predictive models identify patterns, trends, and relationships within data and use them to forecast future outcomes. Key techniques discussed in the unit include **Regression Analysis, Time Series Analysis, Classification, and Clustering**, which together form the foundation of predictive modelling in data science. These methods enable organizations to anticipate trends, estimate future values, and support proactive decision-making in areas such as business forecasting, healthcare analysis, financial modelling, and customer behaviour prediction.

The unit further explains each predictive technique in detail. **Regression analysis** is used to model relationships between variables and predict continuous outcomes such as sales or prices. **Time series analysis** focuses on data collected over time and identifies patterns such as trends, seasonality, and cycles to forecast future values. **Classification techniques**, including Naïve Bayes, Logistic Regression, and Decision Trees, are used to assign observations to predefined categories based on learned patterns from historical data. In contrast, **clustering** is an unsupervised learning technique that groups similar data points together to uncover hidden patterns and segments within datasets. Through these methods, predictive data analysis transforms historical information into actionable insights, allowing analysts and organizations to build models that anticipate future behaviour and improve data-driven decision-making.

3.5 SOLUTIONS/ANSWERS

☞ **Check Your Progress 1**

Q1. Explain the continuum of data analytics from descriptive to prescriptive analysis.
Solution : Refer to section 3.1 & 3.3
Q2. What is Predictive Data Analysis? How does it differ from descriptive and diagnostic analysis?
Solution: Refer to section 3.1 & 3.3
Q3. List and briefly explain the main techniques used in Predictive Data Analysis.
Solution: Refer to section 3.1 & 3.3

☞ **Check Your Progress 2**

Q1. What is Regression Analysis? Explain its role in Predictive Data Analysis
Solution: Refer to section 3.3.1
Q2. Differentiate between Simple Linear Regression and Multiple Linear Regression
Solution: Refer to section 3.3.1
Q3. Explain Logistic Regression and Polynomial Regression with their use cases
Solution: Refer to section 3.3.1

☞ **Check Your Progress 3**

Q1. What is Time Series Analysis? Why is it important in Predictive Data Analysis?
Solution: Refer to section 3.3.2
Q2. Explain the components of a Time Series.
Solution: Refer to section 3.3.2

Q3. What is Moving Average in Time Series Analysis? Explain its purpose with an example.
Solution: Refer to section 3.3.2

☞ **Check Your Progress 4**

Q1. What is Classification in Predictive Data Analysis? Explain its key characteristics.
Solution: Refer to section 3.3.3
Q2. Explain the steps involved in the classification process.
Solution: Refer to section 3.3.3
Q3. Discuss any two classification algorithms with examples.
Solution: Refer to section 3.3.3

☞ **Check Your Progress 5**

Q1. What is Clustering in Predictive Data Analysis? Explain its purpose
Solution: Refer to section 3.3.4
Q2. Differentiate between Classification and Clustering.
Solution: Refer to section 3.3.4
Q3. Explain the role of clustering in predictive data analysis with an example.
Solution: Refer to section 3.3.4

3.6 FURTHER READINGS

1. *Handbook of Statistical Analysis and Data Mining Applications* by Robert Nisbet, John Elder, and Gary Miner, published by Elsevier (2nd Edition, 2017, ISBN: 978-0124166455).
2. *Data Mining and Analysis: Fundamental Concepts and Algorithms* by Mohammed J. Zaki and Wagner Meira Jr., published by Cambridge University Press (1st Edition, 2014, ISBN: 978-0521766333).
3. *Core Concepts in Data Analysis: Summarization, Correlation and Visualization* by *Boris Mirkin, published by Springer (1st Edition, 2011, ISBN: 978-0857292872).*
4. *Applied Predictive Modeling, Max Kuhn and Kjell Johnson, 1st Edition, Springer, 2013.*
5. *An Introduction to Statistical Learning with Applications in R, Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani, 1st Edition, Springer, 2013.*
6. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Trevor Hastie, Robert Tibshirani and Jerome Friedman, 2nd Edition, Springer, 2009.*
7. *An Introduction to Statistical Learning with Applications in R by Gareth James, Daniela Witten, Trevor Hastie, and Robert Tibshirani, published by Springer (2nd Edition, 2021, ISBN: 978-1071614174).*
8. *Time Series Analysis: Forecasting and Control by George E. P. Box, Gwilym M. Jenkins, Gregory C. Reinsel, and Greta M. Ljung, published by Wiley (5th Edition, 2015, ISBN: 978-1118675021)*

UNIT-4 PRESCRIPTIVE DATA ANALYSIS

4.0 INTRODUCTION

Data analysis has evolved into a structured continuum, where each stage builds upon the insights of the previous one. At the foundation lies **descriptive analysis**, which summarizes historical data to answer the question *"What happened?"*. This is followed by **diagnostic analysis**, which probes deeper into causal relationships to explain *"Why did it happen?"*. Progressing further, **predictive analysis** leverages statistical modelling, machine learning, and advanced algorithms to forecast *"What is likely to happen?"*.

While these stages provide valuable knowledge, they remain incomplete without a final step: **prescriptive analysis**. Prescriptive data analysis goes beyond understanding and forecasting—it focuses on *"What should we do?"*. It integrates insights from descriptive, diagnostic, and predictive approaches to recommend optimal courses of action. In this way, prescriptive analysis establishes connectivity across the entire analytical spectrum, transforming raw data into actionable strategies.

Prescriptive methods are inherently decision-oriented. They rely on advanced techniques such as **optimization** to identify the best possible solutions under given constraints, **simulation** to model complex systems and evaluate outcomes under uncertainty, **game theory** to analyse competitive and cooperative scenarios, and **decision-analysis frameworks** to weigh trade-offs and risks systematically. Together, these tools empower organizations not only to anticipate the future but also to shape it deliberately, ensuring that data-driven decisions are both informed and strategically sound.

In essence, prescriptive analysis represents the culmination of the data analytics journey—where knowledge, prediction, and strategy converge to guide intelligent action.

4.1 EXPECTED LEARNING OUTCOMES

After completing this unit, you will be able to:

- Understand the concept and significance of prescriptive data analysis in the analytics continuum.
- Explain how prescriptive analysis integrates descriptive, diagnostic, and predictive insights.
- Identify techniques such as optimization, simulation, and decision analysis used in prescriptive analytics.
- Analyse how prescriptive models support decision-making under constraints and uncertainty.
- Understand the role of prescriptive analytics in recommending optimal actions and strategies.
- Develop the ability to apply prescriptive approaches for data-driven decision-making in real-world scenarios.

4.2 PRESCRIPTIVE DATA ANALYSIS

Prescriptive data analysis represents the pinnacle of the analytics continuum, where insights from descriptive, diagnostic, and predictive stages are transformed into actionable strategies. Unlike earlier stages that focus on understanding and forecasting, prescriptive analysis is inherently **decision-oriented**—it seeks to answer *"What should we do?"* by recommending the best possible course of action under given constraints and uncertainties.

To achieve this, prescriptive analytics leverages a set of powerful methodologies: **optimization, simulation, game theory, and decision-analysis methods**. Each plays a distinct role in shaping intelligent, data-driven decisions.

In the subsequent sections we are going to cover these methodologies one by one, but to begin with we need to understand them in brief, let's begin with Optimization and thereafter we will go ahead with other methodologies like Simulation, Game theory, and Decision analysis methods.

1. Optimization is the mathematical backbone of prescriptive analytics. It involves finding the most efficient solution to a problem given constraint such as cost, resources, or time.

- **Linear and Nonlinear Programming**: Used to allocate resources optimally across competing demands.

- **Integer Programming**: Essential for discrete decision-making, such as scheduling or logistics.

- **Multi-objective Optimization**: Balances trade-offs between conflicting goals, such as minimizing cost while maximizing quality.

In practice, optimization enables businesses to streamline supply chains, financial portfolios, and production schedules, ensuring that decisions are not just feasible but optimal.

2. Simulation complements optimization by modelling complex systems under uncertainty. It allows analysts to test scenarios and evaluate outcomes without real-world risks.

- **Monte Carlo Simulation**: Uses random sampling to estimate probabilities and outcomes in uncertain environments.
- **Discrete Event Simulation**: Models processes such as customer service queues or manufacturing lines.
- **System Dynamics**: Captures feedback loops and long-term behaviour in complex systems.

Simulation is particularly valuable when exact optimization is impractical due to uncertainty or complexity. For example, healthcare systems use simulation to evaluate patient flow, while finance applies it to stress-test investment strategies.

3. Game Theory introduces a strategic dimension to prescriptive analytics by analysing interactions among multiple decision-makers. It is crucial when outcomes depend not only on one's own choices but also on the actions of competitors, partners, or regulators.

- **Nash Equilibrium**: Identifies stable strategies where no player benefits from unilateral deviation.
- **Cooperative Game Theory**: Explores coalition-building and shared benefits.
- **Mechanism Design**: Structures incentives to achieve desired outcomes.

Applications range from pricing strategies in competitive markets to negotiations in supply chains and policymaking in regulatory environments. Game theory ensures that prescriptive analytics accounts for strategic interdependence.

4. Decision-Analysis Methods provides a structured framework for evaluating alternatives under uncertainty, incorporating both quantitative and qualitative factors.

- **Decision Trees**: Map out possible outcomes and associated probabilities.
- **Expected Utility Models**: Weigh risks and rewards to guide rational choices.
- **Multi-Criteria Decision Analysis (MCDA)**: Integrates diverse Objectives such as cost, sustainability, and customer satisfaction.

These methods are indispensable when decisions involve uncertainty, trade-offs, or subjective preferences. They help organizations make transparent, rational, and defensible choices.

Thus, we learned that Prescriptive data analysis is not merely the final stage of analytics—it is the **strategic engine** that transforms knowledge into action. By integrating optimization, simulation, game theory, and decision-analysis methods, organizations can move beyond prediction to deliberate influence, shaping outcomes that are efficient, competitive, and sustainable. In this way, prescriptive analytics completes the journey from *"What happened?"* to *"What should we do?"*, ensuring that data-driven decisions are both informed and impactful.

Now we take up the methodologies for Prescriptive data analysis and discuss them in detail, lets begin with Optimization

☞ **Check Your Progress 1**

Q1. Explain the role of Prescriptive Data Analysis in the data analytics continuum.

………………………………………………………………………………………

………………………………………………………………………………………

Q2. What is Prescriptive Data Analysis? How does it differ from Predictive Analysis?

………………………………………………………………………………………

………………………………………………………………………………………

Q3. List and briefly explain the key methodologies used in Prescriptive Data Analysis.

………………………………………………………………………………………

………………………………………………………………………………………

4.2.1 OPTIMIZATION

Optimization is a fundamental concept in data science, operations research, and prescriptive analytics. It provides a systematic approach to determine the best possible decision or solution from a set of alternatives while satisfying given constraints. In prescriptive analytics, optimization models analyse available data and recommend actions that maximize benefits or minimize costs under specific limitations such as resources, time, budget, or operational capacity.

Optimization techniques are widely used in areas such as supply chain management, production planning, financial portfolio management, transportation systems, and scheduling problems. Several mathematical methods support optimization in real-world decision-making. The most commonly used techniques include Linear and Nonlinear Programming, Integer Programming, and Multi-objective Optimization. Now, we will discuss each one of the techniques in brief.

1. Linear and Nonlinear Programming are important mathematical techniques used in optimization to determine the most efficient way to allocate limited resources while satisfying given constraints. These methods are widely applied in business analytics, operations management, engineering design, and financial planning.

- **Linear Programming (LP)** is a technique used when both the objective function and the constraints are linear relationships. In this method, the goal is to maximize or minimize a linear objective function, such as profit, cost, or time, subject to a set of linear constraints that represent limitations on resources. Linear programming assumes proportionality and additivity among variables, meaning that changes in decision variables affect the outcome in a directly proportional manner. The decision variables represent quantities that need to be determined in order to achieve the optimal solution. Linear programming models are typically solved using methods such as the Simplex method or graphical techniques for smaller problems. In practice, LP is commonly used in production planning, transportation and distribution problems, workforce allocation, and financial budgeting. For instance, a manufacturing firm may use linear programming to determine how many units of different products should be produced to maximize total profit while staying within limits of available labor, raw materials, and machine time.
- **Nonlinear Programming (NLP)** extends optimization to situations where the objective function or constraints involve nonlinear relationships. In many real-world systems, relationships between variables are not strictly linear and may involve quadratic, exponential, logarithmic, or other nonlinear functions. Nonlinear programming allows these complex relationships to be modeled more accurately. Because nonlinear problems may have multiple local optimum points rather than a single global optimum, they are often more computationally challenging to solve. Advanced algorithms such as gradient-based methods, interior-point methods, or heuristic approaches are often used to obtain optimal or near-optimal solutions. Nonlinear programming is widely used in engineering design optimization, financial portfolio management, machine learning parameter tuning, and energy system optimization. For example, in financial portfolio optimization, the relationship between expected return and investment risk is typically nonlinear, and nonlinear programming can help determine the optimal allocation of funds among different assets to achieve the best risk-return balance.

2. Integer Programming is an optimization technique used when some or all decision variables must take integer (whole-number) values. This type of programming is particularly important in situations where fractional solutions are not meaningful or feasible in practice. For example, in many operational decisions such as assigning employees, scheduling machines, or selecting projects, the outcomes must be expressed in discrete units rather than continuous values.

Integer programming can be classified into several forms depending on the nature of the decision variables.

- In pure integer programming, all decision variables are restricted to integer values.
- In mixed-integer programming (MIP), some variables are continuous while others must be integers.

- Another common form is binary integer programming, in which decision variables take only two possible values, typically 0 or 1, representing decisions such as yes/no, selected/not selected, or on/off.

Integer programming problems are generally more complex than linear programming problems because the requirement of integer solutions significantly increases the computational effort required to find the optimal solution. Techniques such as branch-and-bound, cutting plane methods, and branch-and-cut algorithms are commonly used to solve these problems.

Integer programming is widely applied in areas such as workforce scheduling, logistics planning, facility location selection, and project management. For instance, in airline crew scheduling, it is necessary to assign a whole number of crew members to each flight while ensuring compliance with labor regulations, working hours, and operational constraints. Similarly, in logistics management, integer programming can help determine the optimal number of delivery vehicles required and the routes they should follow in order to minimize transportation costs while meeting delivery deadlines.

3. Multi-objective Optimization is an advanced optimization technique used when a decision problem involves two or more Objectives that need to be optimized simultaneously. In many real-world situations, decision-makers must consider multiple goals that may conflict with each other. For example, an organization may aim to minimize operational costs while maximizing product quality or customer satisfaction. Achieving the best possible balance among such competing Objectives requires specialized optimization approaches.

Unlike single-objective optimization, where a single optimal solution can often be identified, multi-objective optimization typically produces a set of solutions known as Pareto optimal solutions. A solution is considered Pareto optimal if no objective can be improved without causing deterioration in at least one other objective. The collection of these optimal trade-off solutions forms the Pareto frontier, which helps decision-makers understand the trade-offs between different Objectives and select the most appropriate solution based on strategic priorities.

Various methods are used to solve multi-objective optimization problems, including weighted sum approaches, goal programming, evolutionary algorithms, and Pareto-based optimization techniques. These methods help transform multiple Objectives into a manageable form so that feasible solutions can be evaluated and compared.

Multi-objective optimization is widely used in fields such as sustainable manufacturing, transportation network planning, healthcare resource allocation, and environmental management. For example, in manufacturing systems, managers may attempt to reduce production costs while simultaneously improving product quality and minimizing environmental impact. Multi-objective optimization enables decision-makers to analyze these competing goals systematically and identify solutions that achieve an acceptable compromise among them. By supporting balanced and informed decision-making, this technique plays an important role in complex strategic planning and modern prescriptive analytics.

To understand the concept of Optimization, we are going to discuss Linear Programming as one of the optimization techniques

Linear Programming (LP) is a mathematical optimization technique used to determine the best possible outcome (maximum profit or minimum cost) in a model whose requirements are represented by linear relationships. It helps decision-makers allocate limited resources efficiently while satisfying certain restrictions.

Linear programming is widely used in production planning, transportation problems, scheduling, supply chain management, and financial planning. A linear programming model consists of several important components that define the decision problem.

Following are the Key Terms used in Linear Programming:

1. Decision Variables represent the quantities that the decision-maker needs to determine in order to achieve the optimal solution. These variables correspond to the possible actions or choices available in the problem.
For example, if a company produces two products, the number of units of each product to produce will be the decision variables.

Example:
Let
x = Number of units of Product A produced
y = Number of units of Product B produced

These variables will be determined through the optimization process.

2. Objective Function is a mathematical expression that represents the goal of the optimization problem. It specifies whether the aim is to maximize or minimize a particular quantity such as profit, cost, revenue, or production time.

Example:
If the profit from Product A is ₹40 per unit and from Product B is ₹30 per unit, the objective function becomes:

$$Z=40x+30y$$

Here

- Z = total profit
- The goal is to Maximize Z

3. Constraints are the limitations or restrictions placed on the decision variables. These usually arise due to limited availability of resources such as raw materials, labour hours, machine time, or budget.
Constraints are expressed as linear inequalities or equations.

Example: If product manufacturing requires limited labour and materials:
Labor constraint: $2x + y \leq 100$
Material constraint: $x + y \leq 80$
These inequalities restrict the feasible values of x and y.

4. Non-Negativity Condition says that the Decision variables cannot take negative values because production quantities cannot be negative i.e. $x \geq 0$, $y \geq 0$

Basic Steps in Linear Programming are:

1. Formulate the problem:
   - Define the decision variables,
   - Define the objective function, and
   - Define constraints.
2. Graph the feasible region: In two or three dimensions, the feasible region can be graphed.
3. Find the optimal solution: Use methods like the Simplex Method, Graphical Method (for two variables), or Interior-point methods to find the optimal solution.

Now let's understand linear programming by a Numerical Example

Example: A factory produces two products A and B.

- Profit from Product A = ₹40 per unit
- Profit from Product B = ₹30 per unit

Each unit requires resources as follows:

Resource	Product A	Product B	Available
Labor Hours	2	1	100
Raw Material	1	1	80

The company wants to determine how many units of each product should be produced to maximize profit.

Solution:

Step 1: Define Decision Variables

Let, x = number of units of Product A
y = number of units of Product B

Step 2: Objective Function

The objective is to maximize profit (Z).
$Z = 40x + 30y$
Maximize total profit Z.

Step 3: Formulate Constraints

Labor constraint: $2x + y \leq 100$
Raw material constraint: $x + y \leq 80$
Non-negativity constraints are $x \geq 0$, $y \geq 0$

Solution Using Graphical Method

The graphical method is applicable when there are only two decision variables.

Step 1: Convert inequalities to equations $2x + y = 100$ and $x + y = 80$

Step 2: Find intercepts

For 2x+y=100

Put x=0 in above equation we get y=100

Put y=0 in above equation we get x=50

Thus, the Points of intercept are: (0,100) and (50,0)

Similarly For x + y = 80

Put x=0 in above equation we get y=80

Put y=0 in above equation we get x=80

Thus, the Points of intercept are: (0,80) and (80,0)

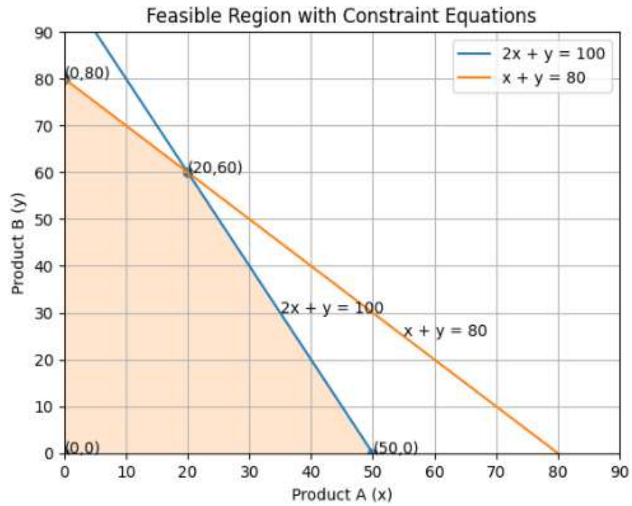

Step 3: Find Corner Points (Feasible Region)

The feasible region is formed by the intersection of constraints.

Corner points are:

1. A (0,0)
2. B (0,80)
3. C (20,60) (intersection of both constraints)
4. D (50,0)

Step 4: Evaluate Objective Function

Point	x	y	Z = 40x + 30y
A	0	0	0
B	0	80	2400
C	20	60	2600
D	50	0	2000

Step 5: Identify Optimal Solution

Maximum profit occurs at (x , y) = (20,60)

Maximum profit: Z= 40x + 30y = 40*20 + 30*60 = 2600

Final Optimal Solution: The company should produce 20 units of Product A and 60 units of Product B to Generate Maximum Profit of ₹2600. This solution satisfies all the constraints while giving the highest possible profit.

The graphical method shows that the optimal solution occurs at one of the corner points of the feasible region, which is a fundamental property of linear programming problems. We will learn about the graphical method also in our next example.

Example: Suppose a furniture company makes chairs and tables only. Each chair gives a profit of Rs. 20 whereas each table gives a profit of Rs. 30. Both products are processed by three machines M1 , M2 and M3. Each chair requires 3 hrs, 5 hrs and 2 hrs on M1 , M2 and M3 respectively, whereas the corresponding figures for each table are 3 hrs, 2 hrs, 6 hrs on M1 , M2 and M3. respectively. The machine M1 can work for 36 hrs per week, whereas M2 and M3 can work for 50 hrs and 60 hrs per week, respectively. How many chairs and tables should be manufactured per week to maximize the profit? "solve this problem by Linear programming

Step 1: Define Decision Variables

Let:

- x = number of **chairs** produced per week
- y = number of **tables** produced per week

These are the **decision variables** because the company must decide how many chairs and tables to manufacture.

Step 2: Formulate the Objective Function

Profit from:

- One chair = Rs. 20
- One table = Rs. 30

Therefore, total profit Z is:

$Z = 20x + 30y$

The objective is:

Maximize Z

Step 3: Formulate the Constraints

Machine M1

- Chair requires **3 hours**
- Table requires **3 hours**
- Available time = **36 hours**

$3x + 3y \leq 36$

or

$x+y\leq 12$

Machine M2

- Chair requires **5 hours**
- Table requires **2 hours**
- Available time = **50 hours**

$5x + 2y \leq 50$

Machine M3

- Chair requires **2 hours**
- Table requires **6 hours**
- Available time = **60 hours**

$2x + 6y \leq 60$

Non-Negativity Conditions

$x\geq 0,\ y\geq 0$

Step 4: Determine Corner Points (Graphical Solution)

Convert inequalities into equations.

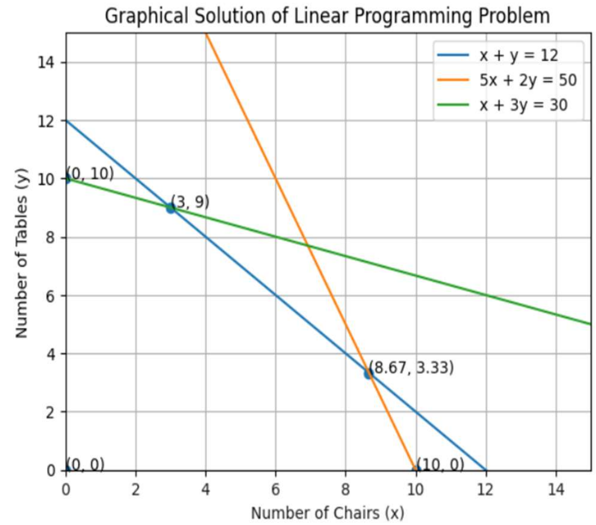

Constraint 1 : $x + y = 12$

Intercepts: x=12, y=0 and x=0, y=12

Constraint 2 : $5x + 2y = 50$

Intercepts: x=10, y=0 and x=0, y=25

Constraint 3 : $2x+6y=60$

Intercepts: x=30, y=0 and x=0, y=10

Intersection of Important Constraints

Intersection of $x + y = 12$ and $5x + 2y = 50$

Substitute $y = 12 – x$ in $5x + 2y = 50$

we get $5x+2(12−x) = 50$

⇨ $5x+24−2x=50$

$3x=26$

$x=8.67$ similarly we get $y=3.33$

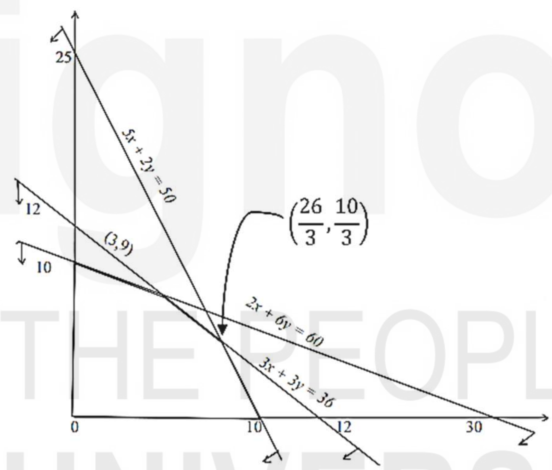

Feasible Corner Points

Approximate feasible points:

<u>Point</u>	<u>x</u>	<u>y</u>
A	0	10
B	3	9
C	8.67	3.33
D	10	0

Step 5: Evaluate Objective Function

$Z=20x+30y$

Point	x	y	Profit Z
A	0	10	300
B	3	9	330
C	8.67	3.33	273
D	10	0	200

Step 6: Optimal Solution says that the Maximum profit occurs at (x, y) = (3,9)

Therefore, to achieve maximum profit, the number of chairs to be produced is 3, and the number of tables to be produced is 9, leading to Maximum Profit: Z=330. Since tables provide higher profit per machine-hour under the given constraints, the optimal strategy is to produce only tables within the machine capacity limits.

We conclude this section with this example on Linear programming, the details related to Non-Linear programming, Integer programming and Multi Objective Optimization are out of scope of this course. However, this section had already given a brief discussion on these topics. Now, we will extend our discussion to the topic of Simulation in the next section.

☞ **Check Your Progress 2**

Q1. What is Optimization in Prescriptive Data Analysis? Explain its importance.

………………………………………………………………………………………………

………………………………………………………………………………………………

Q2. Explain the key components of a Linear Programming model.

………………………………………………………………………………………………

………………………………………………………………………………………………

Q3. Differentiate between Linear Programming and Integer Programming.

………………………………………………………………………………………………

………………………………………………………………………………………………

4.2.2 SIMULATION

Simulation is an important analytical technique used to model and analyse complex systems in order to understand their behaviour under different conditions. In many real-world situations, systems are affected by uncertainty, randomness, and dynamic interactions, making it difficult to predict outcomes using simple mathematical models. Simulation helps decision-makers evaluate the potential consequences of different decisions or scenarios without having to experiment directly in the real world.

The main idea behind simulation is to create a computational model that imitates the functioning of a real system. By running the model repeatedly with varying inputs, analysts can observe possible outcomes and measure system performance. This approach is particularly useful when analytical solutions are difficult or impossible to derive. Simulation is widely used in areas such as supply chain management, healthcare systems, financial risk analysis, transportation systems, and manufacturing operations.

Simulation models provide several advantages. They allow organizations to test alternative strategies, identify bottlenecks in processes, estimate risks, and understand the impact of uncertainty on decision-making. Among the most widely used simulation techniques are Monte Carlo Simulation, Discrete Event Simulation, and System Dynamics.

In the remaining part of this section, we are going to discuss these Simulation techniques, but the major emphasis will be on the Monte Carlo Simulation only. Later we will discuss the other two techniques Viz. Discrete Event Simulation, and System Dynamics in brief. Let's begin with the Monte Carlo Simulation.

Monte Carlo Simulation

Monte Carlo Simulation is a statistical technique that uses random sampling and probability distributions to estimate possible outcomes of a system. The method is particularly useful when the system involves uncertainty in input variables. Instead of using fixed values, Monte Carlo simulation treats inputs as random variables with defined probability distributions and repeatedly simulates the system to generate a range of possible results.

The process typically involves generating a large number of random samples for uncertain inputs, computing the output for each sample, and then analysing the resulting distribution of outcomes. This repeated experimentation allows analysts to estimate probabilities, expected values, and risk levels associated with different decisions.

Monte Carlo simulation is widely used in financial risk analysis, project management, investment portfolio evaluation, and engineering reliability studies. For example, in financial modelling, analysts may simulate thousands of possible market scenarios to estimate the probability that a portfolio will achieve a certain return or incur potential losses. Similarly, in project management, Monte Carlo simulation can be used to estimate the probability that a project will be completed within a specified time or budget by accounting for uncertainties in task durations.

The strength of Monte Carlo simulation lies in its ability to capture uncertainty and variability, providing decision-makers with a probabilistic understanding of possible outcomes rather than a single deterministic result.

Basic Idea of Monte Carlo Simulation involves following steps:

1. Define the problem and uncertain variables.
2. Determine the probability distribution of the variables.
3. Generate random numbers representing possible outcomes.
4. Map random numbers to outcomes based on probability ranges.
5. Repeat the process multiple times to simulate different scenarios.
6. Compute statistical measures such as average, variance, or probability of occurrence.

Mathematical Representation of Monte Carlo simulation often estimates the **expected value** of a random variable. The expected value can be expressed as: $E(X) = \Sigma_{(i=1 \text{ to } n)} X_i P(X_i)$

Where:

- Xi = possible outcome
- $P(Xi)$ = probability of that outcome
- $E(X)$ = expected value

In simulation, the estimated expected value after running N trials is: $\hat{E}(X) = (\Sigma_{(i=1 \text{ to } N)} X_i) / N$

Where:

- N = number of simulations runs
- Xi = outcome in the i^{th} simulation

As the number of simulations increases, the estimate becomes closer to the true expected value.

Numerical Example: A shop observes the following **daily demand distribution** for a product.

Demand (Units)	**Probability**
0	0.10
1	0.20
2	0.40
3	0.20
4	0.10

Simulate demand for **5 days** using Monte Carlo simulation.

Solution:

Step 1: Construct Cumulative Probability

Demand	*Probability*	*Cumulative Probability*	*Random Number Range*
0	0.10	0.10	00–09
1	0.20	0.30	10–29
2	0.40	0.70	30–69
3	0.20	0.90	70–89
4	0.10	1.00	90–99

Step 2: Generate Random Numbers

Assume the following random numbers:

Day	*Random Number*	*Demand*
1	15	1
2	62	2
3	85	3
4	08	0
5	47	2

Step 3: Compute Simulated Average Demand

Total demand: 1+2+3+0+2=8

Average demand: 8 / 5=1.6

The simulation estimates that the **average demand is approximately 1.6 units per day**. If the simulation were repeated with a larger number of trials (e.g., 100 or 1000 days), the estimate would become more accurate.

Monte Carlo Simulation offers several advantages when analysing complex systems that involve uncertainty and variability.

- One of its primary benefits is that it effectively handles uncertainty by incorporating random variables and probability distributions rather than relying on fixed input values. This enables decision-makers to understand the range of possible outcomes rather than a single deterministic result. Another important advantage is that Monte Carlo simulation can model complex systems and processes that are difficult or impossible to solve using traditional analytical methods. By simulating many possible scenarios, it allows analysts to explore how different factors interact and influence system performance.

- Monte Carlo simulation is also flexible and adaptable, meaning it can be applied to a wide variety of problems across disciplines such as finance, engineering, supply chain management, project management, and data science. It allows decision-makers to evaluate risks and probabilities by estimating the likelihood of different outcomes, which is particularly useful in situations involving uncertain market conditions, project timelines, or demand levels. Additionally, the method is easy to implement using modern computational tools such as spreadsheets, statistical software, and programming languages, making it accessible for practical use in academic and industry settings.

- Another advantage is that Monte Carlo simulation supports scenario analysis and experimentation without real-world consequences, enabling organizations to test alternative strategies and policies before implementing them in practice. It also helps in improving decision-making by providing insights into the expected value, variability, and potential risks associated with different options. Furthermore, as the number of simulation runs increases, the results become more accurate and reliable, making the method particularly valuable for long-term planning and forecasting. Overall, Monte Carlo simulation serves as a powerful tool for analyzing uncertain systems and supporting informed decision-making in complex environments.

In short, the Advantages of Monte Carlo Simulation are:

- Handles uncertainty and randomness effectively
- Useful when analytical solutions are difficult
- Provides probabilistic insights rather than single values
- Can model complex real-world systems

The details of Discrete Even Simulation and System Dynamics are as follows:

Discrete Event Simulation

Discrete Event Simulation (DES) is a modelling technique used to represent systems where changes occur at specific points in time as a result of events. In such systems, the state of the system changes only when certain events take place, such as the arrival of a customer, the completion of a task, or the breakdown of a machine.

In a discrete event simulation model, the system is represented as a sequence of events that occur over time. Each event triggers a change in the system's state, and the simulation tracks these changes to analyse system performance. DES is particularly useful for studying operational processes that involve queues, resource allocation, and workflow management.

This technique is widely applied in areas such as customer service centres, hospital operations, airport management, manufacturing systems, and transportation networks. For example, in a call centre, discrete event simulation can model the arrival of customers, the waiting time in queues, and the service time provided by agents. By simulating these events, managers can evaluate staffing strategies, estimate waiting times, and improve service efficiency.

Similarly, in manufacturing systems, discrete event simulation can help identify production bottlenecks, machine utilization rates, and optimal scheduling policies. Because it closely mirrors real operational processes, DES is a powerful tool for improving system performance and operational efficiency.

System Dynamics

System Dynamics is a simulation approach used to model and analyse complex systems that evolve over time due to interactions among their components. Unlike discrete event simulation, which focuses on individual events, system dynamics emphasizes continuous change, feedback loops, and long-term system behaviour.

System dynamics models represent systems using components such as stocks, flows, feedback loops, and time delays. Stocks represent accumulations (such as inventory or population), flows represent rates of change (such as production or growth rates), and feedback loops describe how different parts of the system influence each other.

This approach is particularly useful for understanding systems where decisions taken today influence future outcomes through complex feedback mechanisms. System dynamics is widely used in areas such as economic modelling, environmental sustainability studies, policy analysis, and organizational management.

For instance, in supply chain management, system dynamics can model how production rates, inventory levels, and customer demand interact over time. Such models help decision-makers understand phenomena like the bullwhip effect, where small changes in demand can lead to large fluctuations in supply chain operations. Similarly, in environmental studies, system dynamics models are used to analyse the long-term impact of policies on population growth, resource consumption, and climate change.

By capturing feedback structures and long-term interactions, system dynamics helps organizations and policymakers understand the dynamic complexity of large-scale systems and design strategies that lead to sustainable outcomes.

With this we conclude this section, and in this section, we learned that Simulation plays a vital role in modern data-driven decision-making and prescriptive analytics. By creating computational representations of real-world systems, simulation enables organizations to experiment with different strategies and analyze their potential impacts before implementing them. Techniques such as Monte Carlo Simulation, Discrete Event Simulation, and System Dynamics allow analysts to model uncertainty, operational processes, and long-term system behavior. These methods provide valuable insights into complex problems, helping decision-makers reduce risks, improve efficiency, and design more effective policies and strategies in uncertain and dynamic environments.

The next topic of discussion in this unit is game theory, an in distinguished part of Prescriptive Data analysis,

☞ **Check Your Progress 3**

Q1. What is Simulation? Explain its role in Prescriptive Data Analysis.

……………………………………………………………………………………………

……………………………………………………………………………………………

Q2. Explain the concept and steps involved in Monte Carlo Simulation.

……………………………………………………………………………………………

……………………………………………………………………………………………

Q3. Differentiate between Discrete Event Simulation and System Dynamics.

……………………………………………………………………………………………
……………………………………………………………………………………………

4.2.3 Game Theory

Game Theory is a mathematical framework used to analyse situations in which the outcome of a decision depends not only on one participant's actions but also on the actions of other participants. In many real-world environments such as markets, negotiations, and policy decisions, multiple decision-makers interact strategically, each trying to maximize their own benefit. Game theory provides tools to study such interactions and predict the possible outcomes of strategic behaviour.

In the context of prescriptive analytics, game theory introduces a strategic perspective to decision-making. While traditional optimization techniques focus on finding the best decision under fixed conditions, game theory considers situations where multiple agents make decisions simultaneously or sequentially, and the outcome depends on the choices made by all participants. These agents, often referred to as players, may represent firms, individuals, governments, or organizations. Each player has a set of possible actions called strategies, and the result of the interaction is represented by a payoff, which indicates the benefit or loss to each player.

Game theory is widely used in fields such as economics, business strategy, political science, computer science, and operations research. It helps organizations understand competitive behaviour, design effective strategies, and anticipate the reactions of competitors or partners. Important concepts within game theory include Nash Equilibrium, Cooperative Game Theory, and Mechanism Design, each of which provides insights into different aspects of strategic interactions.

Now, we are going to discuss the concept of Nash Equilibrium, Cooperative Game Theory, and Mechanism Design. Let's begin with *Nash Equilibrium.*

Nash Equilibrium is One of the most fundamental concepts in game theory is the Nash Equilibrium, proposed by the mathematician John Nash. A Nash equilibrium occurs when each player in a game chooses a strategy that is optimal given the strategies chosen by other players. In other words, no player can improve their payoff by unilaterally changing their strategy while the other players keep their strategies unchanged.

At the Nash equilibrium, all players are making the best possible decisions based on the current situation, and the system reaches a stable state where no participant has an incentive to deviate from their chosen strategy. This concept is particularly useful for analysing competitive environments where organizations must anticipate the reactions of their competitors.

For example, in a competitive market, two companies may independently decide whether to set high or low prices for their products. Each company's profit depends not only on its own pricing decision but also on the pricing strategy adopted by the competitor. The Nash equilibrium represents a stable pricing combination where neither company benefits by changing its price strategy alone.

Nash equilibrium is widely applied in market competition analysis, pricing strategies, auction design, and negotiation scenarios, helping decision-makers identify stable outcomes in strategic interactions.

Next, we are going to discuss about *Cooperative Game Theory*

Cooperative Game Theory examines situations in which players can benefit by forming alliances or coalitions, while traditional game theory often focuses on competition between independent players. In the cases of Cooperative Game Theory i.e. forming alliances or coalitions, participants may collaborate and coordinate their strategies to achieve outcomes that are more beneficial than acting independently.

Cooperative game theory studies how the total benefits generated by cooperation can be distributed fairly among the participating members of a coalition. It addresses questions such as how to allocate profits among partners, how to ensure fairness in collaboration, and how to maintain stable alliances among participants.

A common example of cooperative game theory is found in business partnerships or joint ventures, where multiple organizations collaborate to achieve shared goals such as developing new products or entering new markets. The combined effort may produce greater value than individual efforts, and cooperative game theory provides mathematical methods to determine how the resulting gains should be shared.

This concept is also widely used in supply chain coordination, international trade agreements, cost-sharing arrangements, and collaborative research projects. By understanding how cooperation influences outcomes, organizations can design strategies that promote mutually beneficial partnerships.

Next, we are going to discuss about *Mechanism Design*

Mechanism Design is a branch of game theory that focuses on designing systems, rules, or institutions that guide participants toward desired outcomes. While traditional game theory analyses how players behave within existing rules, mechanism design works in the opposite direction: it designs the rules of the game so that rational participants are motivated to act in ways that achieve socially or economically desirable results.

Mechanism design relies on the concept of incentives. By structuring incentives appropriately, policymakers or organizations can influence the behaviour of participants to align individual interests with collective Objectives. For example, governments may design tax policies, subsidies, or regulatory frameworks to encourage businesses to adopt environmentally sustainable practices.

One well-known application of mechanism design is auction design, where the rules of bidding are structured to ensure fairness, transparency, and efficient allocation of resources. Online advertising auctions, spectrum auctions for telecommunications, and procurement bidding processes often rely on mechanism design principles.

Mechanism design is widely used in public policy, market regulation, voting systems, contract theory, and digital platforms, ensuring that systems operate efficiently even when participants act in their own self-interest.

We learned that Game theory plays an important role in prescriptive analytics by incorporating strategic interdependence into decision-making models. in context of Prescriptive analytics, the **applications of game theory** include:

- Pricing Strategies in Competitive Markets: Businesses analyse competitors' actions to determine optimal pricing decisions.
- Negotiations in Supply Chains: Manufacturers, suppliers, and distributors use strategic models to negotiate contracts and profit-sharing arrangements.
- Policy and Regulatory Decision-Making: Governments design policies that influence the behaviour of industries and individuals.
- Auction and Bidding Systems: Strategic models help design efficient auction mechanisms.
- Resource Allocation and Conflict Resolution: Game theory helps identify fair and stable solutions when multiple stakeholders compete for limited resources.

In this section we understood that Game theory provides a powerful framework for understanding strategic interactions among multiple decision-makers. By analysing how individuals or organizations respond to each other's actions, it helps predict outcomes in competitive and cooperative environments. Concepts such as Nash Equilibrium, Cooperative Game Theory, and Mechanism Design enable analysts to study stable strategies, collaborative opportunities, and incentive structures. In prescriptive analytics, the integration of game theory ensures that decision models account not only for optimization and uncertainty but also for the strategic behaviour of interacting participants, leading to more realistic and effective decision-making solutions.

☞ **Check Your Progress 4**

Q1. What is Game Theory? Explain its role in Prescriptive Data Analysis.

…………………………………………………………………………………………

…………………………………………………………………………………………

Q2. What is Nash Equilibrium? Explain with an example.

…………………………………………………………………………………………

…………………………………………………………………………………………

Q3. Differentiate between Cooperative Game Theory and Mechanism Design.

……………………………………………………………………………………………
……………………………………………………………………………………………

4.2.4 DECISION-ANALYSIS METHODS

Decision-analysis methods provide systematic approaches for evaluating and selecting among different alternatives, particularly in situations involving uncertainty, multiple Objectives, and incomplete information. In real-world decision-making, organizations often face complex problems where several possible actions exist, and each action may lead to different outcomes with varying probabilities and consequences. Decision analysis helps decision-makers examine these alternatives in a structured and logical manner.

In the context of prescriptive analytics, decision-analysis methods integrate both quantitative information (such as probabilities, costs, and revenues) and qualitative considerations (such as preferences, risk tolerance, and strategic priorities). These methods help decision-makers understand possible consequences of different choices, evaluate trade-offs, and identify the most suitable option based on available information and organizational goals.

Decision analysis is widely used in areas such as business strategy, project management, healthcare planning, public policy, and financial decision-making. Some of the most commonly used techniques in decision analysis include Decision Trees, Expected Utility Models, and Multi-Criteria Decision Analysis (MCDA). Each of these methods supports structured decision-making in different types of uncertain environments.

Now, we are going to understand the concepts of Decision Trees, Expected Utility Models, and Multi-Criteria Decision Analysis (MCDA). Let's begin with *Decision Trees*.

A Decision Tree is a graphical representation of a decision-making process that shows possible choices, uncertain events, and their associated outcomes. It resembles a tree-like structure consisting of decision nodes, chance nodes, and outcome branches.

- Decision nodes represent points where a decision-maker must choose among alternatives.
- Chance nodes represent uncertain events with associated probabilities.
- Branches represent possible outcomes or consequences of decisions and events.

Decision trees are particularly useful when decisions must be made in stages, and when outcomes depend on uncertain future events. The method allows decision-makers to visualize the sequence of decisions and their potential consequences.

To evaluate a decision tree, analysts often calculate the Expected Monetary Value (EMV) of each alternative by multiplying possible outcomes by their probabilities and summing the results. The alternative with the highest expected value is typically selected as the optimal decision.

Decision trees are commonly used in investment planning, medical treatment decisions, product launch strategies, and project risk evaluation. By clearly mapping possible scenarios and their likelihoods, decision trees help organizations understand risks and make informed choices.

Next, we are going to discuss about *Expected Utility Models*

Expected Utility Models, while expected monetary value focuses only on financial outcomes, many decision-makers also consider risk preferences when choosing among alternatives. Expected Utility Models extend decision analysis by incorporating the decision-maker's attitude toward risk.

In this approach, outcomes are evaluated using a utility function, which represents the level of satisfaction or value associated with different results. Instead of maximizing expected monetary value, the goal is to maximize expected utility, which accounts for both the magnitude of outcomes and the decision-maker's risk tolerance.

For example, two investment options may have the same expected financial return, but one may involve greater uncertainty. A risk-averse decision-maker may prefer the option with more stable outcomes, even if the expected monetary value is slightly lower. Expected utility theory helps capture such preferences mathematically.

Expected utility models are widely used in financial portfolio selection, insurance decisions, risk management, and strategic planning. By incorporating psychological and behavioural aspects of decision-making, these models provide a more realistic framework for evaluating uncertain alternatives.

Finally, we are going to discuss about *Multi-Criteria Decision Analysis (MCDA)*

Multi-Criteria Decision Analysis (MCDA): In many real-world situations, decisions cannot be evaluated based on a single criterion such as cost or profit. Instead, multiple factors must be considered simultaneously. Multi-Criteria Decision Analysis (MCDA) is a decision-analysis method that helps evaluate alternatives when multiple, often conflicting Objectives are involved.

MCDA involves identifying relevant decision criteria, assigning weights to reflect their importance, and scoring each alternative based on how well it performs on each criterion. The weighted scores are then combined to determine the overall ranking of alternatives.

For example, when selecting a new supplier, an organization may need to consider several factors such as cost, product quality, delivery reliability, environmental sustainability, and customer satisfaction. MCDA provides a structured way to evaluate these diverse criteria and determine which option offers the best overall performance.

Various techniques fall under MCDA, including weighted scoring models, analytic hierarchy process (AHP), and technique for order preference by similarity to ideal solution (TOPSIS). These methods enable decision-makers to balance competing priorities and select solutions that best align with strategic Objectives.

MCDA is widely used in project selection, environmental policy evaluation, urban planning, and resource allocation.

Now it's time to conclude and understand the Importance of Decision-Analysis Methods, the Decision-analysis methods play a crucial role in modern decision-making environments because they help address challenges such as uncertainty, complexity, and competing Objectives. By structuring the decision process and explicitly considering probabilities, preferences, and multiple criteria, these methods promote transparency and rationality in decision-making.

They also encourage decision-makers to clearly define Objectives, identify relevant alternatives, and evaluate potential outcomes systematically. This structured approach reduces the likelihood of biased or intuitive decisions and supports evidence-based reasoning.

☞ **Check Your Progress 5**

Q1. What are Decision-Analysis Methods? Explain their role in Prescriptive Data Analysis.

………………………………………………………………………………………

………………………………………………………………………………………

Q2. Explain the concept of Decision Trees and Expected Utility Models.

………………………………………………………………………………………

………………………………………………………………………………………

Q3. What is Multi-Criteria Decision Analysis (MCDA)? Why is it important?

………………………………………………………………………………………

………………………………………………………………………………………

4.3 SUMMARY

The content presented in this unit "Prescriptive Data Analysis" represents the most advanced stage of data analytics, focusing on recommending optimal actions based on available data, models, and constraints. While descriptive analytics explains past events and predictive analytics forecasts future trends, prescriptive analytics determines what actions should be taken to achieve desired outcomes. It integrates techniques from operations research, statistics, artificial intelligence, and decision science to support effective decision-making. One of the core methods used in prescriptive analytics is optimization, which aims to find the best solution under given constraints such as time, cost, or resources. Techniques like linear programming, nonlinear programming, integer programming, and multi-objective optimization help organizations allocate resources efficiently, balance competing goals, and improve decision-making in areas such as production planning, logistics, and financial management.

Another important component of prescriptive analytics is simulation, which allows decision-makers to evaluate system behavior under uncertainty by testing different scenarios without real-world risk. Methods such as Monte Carlo simulation, discrete event simulation, and system dynamics help analyze complex processes and forecast possible outcomes. Game theory further enhances prescriptive analytics by studying strategic interactions among multiple decision-makers, using concepts such as Nash equilibrium, cooperative game theory, and mechanism design to understand competitive and collaborative behavior. In addition, decision-analysis methods provide structured frameworks for evaluating alternatives under uncertainty. Techniques such as decision trees, expected utility models, and multi-criteria decision analysis (MCDA), including methods like the Analytic Hierarchy Process (AHP), help decision-makers compare alternatives, consider multiple Objectives, and select the most appropriate course of action. Together, these approaches enable organizations to make informed, rational, and optimal decisions in complex environments.

4.4 ANSWERS

☞ **Check Your Progress 1**

Q1. Explain the role of Prescriptive Data Analysis in the data analytics continuum.
Solution: refer to section 4.0 & 4.2
Q2. What is Prescriptive Data Analysis? How does it differ from Predictive Analysis?
Solution: refer to section 4.0 & 4.2
Q3. List and briefly explain the key methodologies used in Prescriptive Data Analysis.
Solution: refer to section 4.0 & 4.2

☞ **Check Your Progress 2**

Q1. What is Optimization in Prescriptive Data Analysis? Explain its importance.
Solution: refer to section 4.2.1
Q2. Explain the key components of a Linear Programming model.
Solution: refer to section 4.2.1
Q3. Differentiate between Linear Programming and Integer Programming.
Solution: refer to section 4.2.1

☞ **Check Your Progress 3**

Q1. What is Simulation? Explain its role in Prescriptive Data Analysis.
Solution: refer to section 4.2.2
Q2. Explain the concept and steps involved in Monte Carlo Simulation.
Solution: refer to section 4.2.2
Q3. Differentiate between Discrete Event Simulation and System Dynamics.
Solution: refer to section 4.2.2

☞ **Check Your Progress 4**

Q1. What is Game Theory? Explain its role in Prescriptive Data Analysis.

Solution: refer to section 4.2.3

Q2. What is Nash Equilibrium? Explain with an example.

Solution: refer to section 4.2.3

Q3. Differentiate between Cooperative Game Theory and Mechanism Design.

Solution: refer to section 4.2.3

☞ **Check Your Progress 5**

Q1. What are Decision-Analysis Methods? Explain their role in Prescriptive Data Analysis.

Solution: refer to section 4.2.4

Q2. Explain the concept of Decision Trees and Expected Utility Models.

Solution: refer to section 4.2.4

Q3. What is Multi-Criteria Decision Analysis (MCDA)? Why is it important?

Solution: refer to section 4.2.4

4.5 REFERENCES/FURTHER READINGS

1. *Business Analytics: Data Analysis and Decision Making* by S. Christian Albright and Wayne L. Winston, published by Cengage Learning (7th Edition, 2020, ISBN: 978-0357131787).
2. *Introduction to Operations Research* by Frederick S. Hillier and Gerald J. Lieberman, published by McGraw-Hill Education (10th Edition, 2014, ISBN: 978-0073523453).
3. *Decision Analysis for Management Judgment* by Paul Goodwin and George Wright, published by Wiley (5th Edition, 2014, ISBN: 978-1119974017).
4. *Handbook of Statistical Analysis and Data Mining Applications* by Robert Nisbet, John Elder, and Gary Miner, published by Elsevier (2nd Edition, 2017, ISBN: 978-0124166455).
5. *Data Mining and Analysis: Fundamental Concepts and Algorithms* by Mohammed J. Zaki and Wagner Meira Jr., published by Cambridge University Press (1st Edition, 2014, ISBN: 978-0521766333).
6. *Core Concepts in Data Analysis: Summarization, Correlation and Visualization* by *Boris Mirkin, published by Springer (1st Edition, 2011, ISBN: 978-0857292872).*
7. *Applied Predictive Modeling, Max Kuhn and Kjell Johnson, 1st Edition, Springer, 2013.*
8. *An Introduction to Statistical Learning with Applications in R, Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani, 1st Edition, Springer, 2013.*
9. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Trevor Hastie, Robert Tibshirani and Jerome Friedman, 2nd Edition, Springer, 2009.*
10. *An Introduction to Statistical Learning with Applications in R by Gareth James, Daniela Witten, Trevor Hastie, and Robert Tibshirani, published by Springer (2nd Edition, 2021, ISBN: 978-1071614174).*
11. *Time Series Analysis: Forecasting and Control by George E. P. Box, Gwilym M. Jenkins, Gregory C. Reinsel, and Greta M. Ljung, published by Wiley (5th Edition, 2015, ISBN: 978-1118675021)*